BAN404 Statistical Learning - Exam Companion

Brede N. Espelid

Spring 2025

Contents

1 General R Setup and Tips

1.1 Commonly Used Libraries

- `ISLR` or `ISLR2`: For datasets from the book.

- `boot`: For bootstrap functions (`boot()`, `cv.glm()`).

- `glmnet`: For Ridge and LASSO (`glmnet()`, `cv.glmnet()`).

- `tree`: For classification and regression trees (`tree()`).

- `randomForest`: For Random Forests (`randomForest()`, `varImpPlot()`).

- `gbm`: For Gradient Boosting Machines (`gbm()`).

- `gam`: For Generalized Additive Models (`gam()`, `s()`, `lo()`).

- `e1071`: For Support Vector Machines (`svm()`).

- `MASS`: For LDA, QDA (`lda()`, `qda()`).

- `tidyverse` (or individual packages like `dplyr`, `ggplot2`): For data manipulation and plotting.

- `fastDummies`: For creating dummy variables (`dummy_cols()`).

- `insuranceData`: For `dataCar` dataset (Exam 2021).

- `Ecdat`: For `Computers` dataset (Exam 2022).

1.2 Key R Operations

- Setting seed: `set.seed(number)` for reproducibility.

- Splitting data (50/50 example):

```
n <- nrow(mydata)
set.seed(123)
train_indices <- sample(1:n, floor(n/2))
train_data <- mydata[train_indices, ]
test_data <- mydata[-train_indices, ]
```

- Converting to factors: `mydata$variable <- as.factor(mydata$variable)`

- Model fitting:

- Linear Regression: `lm(y ~ x1 + x2, data=train`$_d ata)LogisticRegression$:

- Evaluation:

  - MSE: `mean((predictions - actuals)^2)`

- Confusion Matrix: `table(predicted`$_c lass, actual_c lass)Accuracy$:

- LASSO/Ridge: Remember to create a model matrix (e.g., `x <- model.matrix(y   .   -1,` `data=train`$_d ata))andaresponsevector$

2 Lecture Summaries and Key Concepts (ISLR)

2.1 Lecture 1: Introduction to Statistical Learning

2.1.1 Defining Business Analytics

- *Analytics (Davenport & Harris, 2007)*: "the extensive use of data, statistical and quantitative analysis, explanatory and predictive models and fact based management to drive decisions and actions."

- *Business Analytics*: Applying analytics in the context of firms making the best possible decisions efficiently using available data.

2.1.2 Categories of Business Analytics

- **Descriptive Analytics**: Understanding past and present through data exploration (e.g., salary structures). Not model-based, not prescriptive.

- **Predictive Analytics**: Forecasting future events or unknown properties (e.g., stock prices, probability of tax evasion). This course primarily focuses here.

- **Prescriptive Analytics**: Recommending actions based on optimization or statistical models (e.g., resource allocation, best medication choice).

2.1.3 Essential Components for Business Analytics Problems

- Management Science / Operations Research: Formulating decision problems from business questions.

- Data Analysis / Statistics: Choosing data and methods to estimate/predict quantities of interest.

- Programming / Database Management: Implementing solutions.

2.1.4 Core Concepts in Statistical Learning

- Definition: A vast set of tools for understanding data.

- Example: Improving sales based on advertising budgets (TV, Radio, Newspaper).

  - Inputs (Predictors, Independent Variables, Features): $X_1, X_2, ..., X_p$.
  - Output (Response, Dependent Variable): Y.

- The Data Generating Process (DGP): $Y = f(X) + \epsilon$

  - f: Systematic information that X provides about Y. Unknown.
  - ϵ: Random error. Zero mean, independent of X. Represents irreducible error.

2.1.5 Goals of Estimating f: Prediction vs. Inference

- **Prediction Focus**:

  - Goal: Estimate Y using $\hat{Y} = \hat{f}(X)$.
  - Accuracy depends on reducible and irreducible error.
  - *Reducible error*: Error from $\hat{f}$ not being a perfect estimate of f. Can be reduced with better models.
  - *Irreducible error*: Due to ϵ. Inherent variability/noise.

- **Inference Focus**:

  - Goal: Understand the relationship between X and Y.
  - Key questions: Which predictors are associated? Strength of association? Functional form of f?

2.1.6 Types of Learning

- **Supervised Learning**: Both predictors X_i and response Y_i are observed.

  - Regression Problems: Y is quantitative (e.g., price, blood pressure).
  - Classification Problems: Y is qualitative/categorical (e.g., disease/no disease, purchase/no purchase).

- **Unsupervised Learning**: Only predictors X_i are observed.

  - Goal: Discover structure (e.g., clustering, PCA).

2.1.7 Assessing Model Accuracy

- **Measuring Quality of Fit - Regression**:

  - *Mean Squared Error (MSE)*: $MSE_{Train} = \text{avg}(y_i - \hat{f}(x_i))^2$.
  - Goal is to minimize $MSE_{Test} = E(Y_0 - \hat{f}(x_0))^2$ for new unseen data (x_0, Y_0).

- **Measuring Quality of Fit - Classification**:

  - *Training Error Rate*: $\frac{1}{n}\sum_{i=1}^{n} I(y_i \neq \hat{y}_i)$.
  - *Test Error Rate*: $\text{Avg}(I(y_0 \neq \hat{y}_0))$ for new unseen data.
  - *Bayes Classifier*: A theoretical ideal that minimizes test error rate. Assigns to class k for which $P(Y = k|X = x_0)$ is largest.
  - *K-Nearest Neighbors (KNN) Classifier*: Practical method, estimates $P(Y = k|X = x_0)$ based on fraction of K nearest neighbors belonging to class k.

2.1.8 The Bias-Variance Trade-Off

- For a given x_0, the expected test MSE: $E(y_0 - \hat{f}(x_0))^2 = \text{Var}(\hat{f}(x_0)) + [\text{Bias}(\hat{f}(x_0))]^2 + \text{Var}(\epsilon)$.

- **Bias**: Error from simplifying a complex real-world problem. More flexible models typically have lower bias.

- **Variance**: How much $\hat{f}$ would change with different training data. More flexible models typically have higher variance.

- **Overfitting vs. Underfitting**:

  - *Underfitting*: High bias, model too simple (e.g., linear model for non-linear data). Low training and test performance typically.

  - *Overfitting*: Low bias (on training data), high variance. Model fits training noise. Low training error, high test error. (Slides p. 17-22 illustrate with polynomial regression on limited training data, showing poor generalization to test data).

2.1.9 The Curse of Dimensionality

- Phenomenon: Test error tends to increase as dimensionality (p) grows, unless the additional predictors are truly associated with the response.

- Reason: In high dimensions, data becomes sparse. Nearest neighbors can be very far away, making local methods (like KNN) less effective.

- Volume of space increases exponentially: s^p. To maintain density, n must grow exponentially.

- Most points in a high-dimensional hypercube are close to the boundary. (R code 'cube.R', 'hypercube.R', 'edges.R').

Listing 1: R code: Fraction of points near boundary (edges.R concept)

```
GetRatio <- function(dim){
  1 - (0.9)^dim # (1 - 20.05)^dim is volume of inner hypercube
}
p_vals <- 1:10 # For demonstration, ISLR uses up to p=10 for nearest
    neighbor distance plots
ratios <- sapply(p_vals, GetRatio)
plot(p_vals, ratios, type="b", xlab="Number of Dimensions (p)",
     ylab="Fraction of Volume Near Boundary (within 10% of range)",
     main="Curse of Dimensionality: Volume Concentration")
abline(h=0.5, col="red", lty=2) # Shows when over 50% of volume is "near
    " boundary
text(5, 0.4, "p=5, ~41% near boundary", pos=4)
text(10, 0.6, "p=10, ~65% near boundary", pos=4)
```

2.1.10 Understanding Loss Functions

- Squared Error Loss: $L(Y, \hat{f}(X)) = (Y - \hat{f}(X))^2$. Minimizing expected squared error leads to prediction $\hat{f}(x) = E(Y|X = x)$.

  - Unconditional case (predict with constant θ): $\theta_{opt} = E(Y)$. Sample estimate $\bar{y}$.

- Absolute Error Loss: $L(Y, \hat{f}(X)) = |Y - \hat{f}(X)|$. Minimizing expected absolute error leads to prediction $\hat{f}(x) = \text{median}(Y|X = x)$.

  - Unconditional case: $\theta_{opt} = \text{median}(Y)$. Sample estimate is sample median.

- Implication: The "best" prediction can change depending on how we define "loss" or "error."

2.1.11   Lab: Introduction to R

- Basic commands, graphics, data indexing, loading data, summaries.

- Relevant files: 'ch2-lab.R' (from GitHub), 'ex2.8.Rmd' (Canvas).

2.1.12   Exercise: ISLR 2.8 (College Data)

- Application of descriptive statistics and basic plotting in R.

- See 'ex2.8.Rmd' (Canvas) or 'ch2-applied.R' (GitHub) for solution ideas.

[12pt,a4paper]article [utf8]inputenc amsmath amssymb graphicx [left=2cm, right=2cm, top=2cm, bottom=2cm]geometry hyperref booktabs longtable enumitem xcolor listings float
    BAN404 Statistical Learning - Exam Companion Brede Espelid Spring 2024

Contents

2.2 Lecture 2: Linear Regression and K-Nearest Neighbors (KNN)

2.2.1 Revisiting Overfitting with Polynomial Regression

- **Context**: Example using a small training dataset (10 observations, x, y) to fit polynomial models of increasing degrees ($k = 1, 2, 3, 4$). The true underlying relationship is quadratic ($k = 2$).

- **Training MSE Behavior** (Slides L2 p. 5):

    - Linear model ($k = 1$): `trMSE1 = 14.37876`

    - Quadratic model ($k = 2$): `trMSE2 = 0.4588453`

    - Cubic model ($k = 3$): `trMSE3 = 0.4320397`

    - Quartic model ($k = 4$): `trMSE4 = 0.4003628`

    - *Observation*: Training MSE decreases as model flexibility (polynomial degree) increases. The 4th order polynomial, being most flexible, has the lowest training MSE. This is typical: more flexible models can better fit the noise in the training data.

```
# Assuming y_train are the true training y values
# and reg1, reg2, reg3, reg4 are lm objects for k=1,2,3,4
# fitted ONLY on the first 10 training observations.
# y_pred_reg1 <- predict(reg1, newdata=training_data_first_10)
# trMSE1 <- mean((y_train_first_10 - y_pred_reg1)^2)
# ... similar for reg2, reg3, reg4
```

Listing 2: R code for Training MSE (Conceptual - from Slides L2 p.5 values)

- **Test MSE Behavior** (Slides L2 p. 6-8):

    - The full dataset actually has 20 observations. The first 10 were training, the next 10 are used as a hold-out test set.

    - Linear model ($k = 1$): `teMSE1 = 3522.564` (Significant underfitting)

    - Quadratic model ($k = 2$): `teMSE2 = 0.9843595` (Best test performance, matches true DGP)

    - Cubic model ($k = 3$): `teMSE3 = 173.0981` (Overfitting starts)

    - Quartic model ($k = 4$): `teMSE4 = 8582.551` (Severe overfitting)

- **Visualizing Over/Underfitting** (Slides L2 p. 7): The plot shows the linear model (black) clearly missing the curve. The quadratic model (blue) tracks the full dataset well. The 3rd/4th order polynomial (green, if shown for k=4 on slide 7, or specifically for k=4 on slide 22 of L1 slides) fits the initial 10 points well but deviates wildly for the test points.

- **Conclusion** (Slide L2 p. 9): Models that are too flexible (e.g., high-degree polynomials here) can achieve low training error by fitting noise, but generalize poorly to unseen test data (high test error). The goal is to find a model complexity that balances fitting the signal without fitting the noise, typically leading to the lowest test error.

2.2.2 Simple Linear Regression

- **Model**: $Y = \beta_0 + \beta_1 X + \epsilon$. Y is response, X is predictor, β_0 is intercept, β_1 is slope, ϵ is error term.

- **Estimating Coefficients using Least Squares**:

  - Objective: Minimize the Residual Sum of Squares (RSS).

  $$RSS = \sum_{i=1}^{n} e_i^2 = \sum_{i=1}^{n} (y_i - \hat{y}_i)^2 = \sum_{i=1}^{n} (y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_i))^2$$

  (Slide L2 p. 28 shows the minimization objective $\min_{\alpha, \beta} g(\alpha, \beta)$).

  - Least Squares Estimates:

  $$\hat{\beta}_1 = \frac{\sum_{i=1}^{n}(x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^{n}(x_i - \bar{x})^2} = \frac{S_{xy}}{S_{xx}}$$

  $$\hat{\beta}_0 = \bar{y} - \hat{\beta}_1 \bar{x}$$

  - R Implementation (`ISLR` `Auto` data, predict `mpg` from `horsepower`):

```
1  # Using Auto dataset from ISLR package
2  library(ISLR)
3  attach(Auto) # or use Auto$horsepower, Auto$mpg
4  lm.fit <- lm(mpg ~ horsepower) # shorthand for lm(mpg ~ horsepower,
        data=Auto)
5  summary(lm.fit)
6  # Output from summary(lm.fit) (similar to Slide L2 p.36 for TV~sales
        ):
7  # Coefficients:
8  #              Estimate Std. Error t value Pr(>|t|)
9  # (Intercept) 39.935861   0.717499    55.66   <2e-16
10 # horsepower  -0.157845   0.006446   -24.49   <2e-16
11
12 plot(horsepower, mpg)
13 abline(lm.fit, col="red", lwd=3)
14 # To manually calculate coefficients (Slide L2 p.37 concept for X, y
        ):
15 # X_design <- cbind(1, horsepower) # Design matrix
16 # y_response <- mpg
17 # beta_hat_manual <- solve(t(X_design) %% X_design) %% t(X_design)
        %% y_response
18 # print(beta_hat_manual) # Matches lm.fit coefficients
19 detach(Auto)
```

Listing 3: Simple Linear Regression in R (ISLR Ch 3 Lab, Slides L2 p. 29, L2 p.36)

2.2.3 Assessing Accuracy of Coefficient Estimates

- The true relationship is $Y = \beta_0 + \beta_1 X + \epsilon$. $\hat{\beta}_0, \hat{\beta}_1$ are estimates.

- **Standard Error (SE)** of estimates: Measures sampling variability. If we refit the model on different random samples, SE tells us how much $\hat{\beta}_j$ would typically vary.

  - $SE(\hat{\beta}_0)^2 = \sigma^2 \left[\frac{1}{n} + \frac{\bar{x}^2}{\sum_{i=1}^{n}(x_i - \bar{x})^2} \right]$

  - $SE(\hat{\beta}_1)^2 = \frac{\sigma^2}{\sum_{i=1}^{n}(x_i - \bar{x})^2}$, where $\sigma^2 = \text{Var}(\epsilon)$.

- σ^2 is usually unknown and estimated by $RSE^2 = RSS/(n-2)$ for simple linear regression.
- Distribution of $\hat{\beta}_1$ (Slide L2 p.30, assuming $\epsilon_i \sim N(0, \sigma^2)$ and are independent):

$$\hat{\beta}_1 \sim N\left(\beta_1, \frac{\sigma^2}{\sum(x_i - \bar{x})^2}\right)$$

- **Confidence Intervals (CI)**: A $(1-\alpha)\%$ CI for β_j is a range that contains the true β_j with $(1-\alpha)\%$ probability in repeated sampling.

  - For β_1: $\hat{\beta}_1 \pm t_{\alpha/2, n-2} \cdot SE(\hat{\beta}_1)$. (Slide L2 p.31 uses 1.96 from z-dist for large n).
  - Example from `summary(lm.fit)()`: The `confint(lm.fit)()` function gives CIs.

- **Hypothesis Testing** for β_1:

  - $H_0 : \beta_1 = 0$ (No linear relationship between X and Y).
  - $H_1 : \beta_1 \neq 0$ (There is a linear relationship).
  - Test statistic: $t = \frac{\hat{\beta}_1 - 0}{SE(\hat{\beta}_1)}$. This t-statistic follows a t-distribution with $n-2$ degrees of freedom under H_0.
  - p-value: The probability of observing a $|t|$-statistic as large or larger than the one computed, assuming H_0 is true. A small p-value (e.g., ¡ 0.05) provides evidence against H_0.
  - `summary(lm.fit)()` output provides t-values and p-values ('Pr(¿—t—)') for each coefficient.

2.2.4 Assessing Overall Model Accuracy

- **Residual Standard Error (RSE)**:

$$RSE = \hat{\sigma} = \sqrt{\frac{RSS}{n-p-1}}$$

(For simple LR, $p = 1$, so $n-2$ in denominator. Slide L2 p. 38).

  - An estimate of $\sigma = \text{std.dev.}(\epsilon)$.
  - Represents the average amount that the response will deviate from the true regression line.
  - Measured in units of Y.

- **R-squared (R^2)**:
$$R^2 = \frac{TSS - RSS}{TSS} = 1 - \frac{RSS}{TSS}$$
where $TSS = \sum(y_i - \bar{y})^2$ is Total Sum of Squares. (Slide L2 p. 38)

  - Proportion of variance in Y that is explained by $X(s)$.
  - $0 \leq R^2 \leq 1$. Closer to 1 indicates better fit.
  - For simple linear regression, $R^2 = (\text{Cor}(X, Y))^2$. (ISLR Ch3 Exercise 3.7 provides a proof for this, which involves algebraic manipulation of RSS, TSS and the formula for correlation).
  - R^2 always increases or stays the same when more predictors are added, even if they are irrelevant. This makes it less suitable for comparing models with different numbers of predictors.

2.2.5 Multiple Linear Regression

- **Model**: $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p + \epsilon$.

- **Estimating Coefficients**:

    - Minimize RSS: $\sum_{i=1}^{n}(y_i - (\hat{\beta}_0 + \hat{\beta}_1 x_{i1} + \cdots + \hat{\beta}_p x_{ip}))^2$.
    - **Matrix Notation** (Slides L2 p. 33-35):
        * Response vector $\mathbf{y}$ ($n \times 1$).
        * Design Matrix $\mathbf{X}$ ($n \times (p+1)$), first column usually all 1s for intercept.
        * Coefficient vector $\boldsymbol{\beta}$ ($(p+1) \times 1$).
        * Model: $\mathbf{y} = \mathbf{X}\boldsymbol{\beta} + \boldsymbol{\epsilon}$.
        * RSS: $g(\boldsymbol{\beta}) = (\mathbf{y} - \mathbf{X}\boldsymbol{\beta})'(\mathbf{y} - \mathbf{X}\boldsymbol{\beta})$.
        * OLS Estimator: $\hat{\boldsymbol{\beta}} = (\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'\mathbf{y}$.
        * Variance-Covariance Matrix of $\hat{\boldsymbol{\beta}}$: $\mathrm{Var}(\hat{\boldsymbol{\beta}}) = \sigma^2(\mathbf{X}'\mathbf{X})^{-1}$. Diagonal elements are $SE(\hat{\beta}_j)^2$.

- **Hypothesis Testing in Multiple Regression**:

    - **For individual coefficients** β_j: $H_0 : \beta_j = 0$ (predictor X_j has no effect, holding others constant).
        * t-statistic: $t_j = \frac{\hat{\beta}_j - 0}{SE(\hat{\beta}_j)}$. Follows t_{n-p-1} under H_0.
        * Provided in `summary(lm.fit)()`.
    - **Overall F-test (Model Significance)**: $H_0 : \beta_1 = \beta_2 = \cdots = \beta_p = 0$ (none of the predictors explain Y).
        * $F = \frac{(TSS-RSS)/p}{RSS/(n-p-1)} = \frac{MSR}{MSE}$.
        * Under H_0, $F \sim F_{p,n-p-1}$.
        * If F is large (small p-value), reject H_0, meaning at least one predictor is useful.
        * Provided in `summary(lm.fit)()`.

- **R Implementation** (Example from ISLR Ch3 Lab):

```
library(MASS) # For Boston dataset
lm.fit.multi <- lm(medv ~ lstat + age, data=Boston)
summary(lm.fit.multi)
lm.fit.all <- lm(medv ~ ., data=Boston) # Using all predictors
summary(lm.fit.all)
```

Listing 4: Multiple Linear Regression for Boston Housing Data

2.2.6 Prediction and Confidence/Prediction Intervals

- Prediction for a new set of predictor values $\mathbf{x}_0$: $\hat{y}_0 = \mathbf{x}_0'\hat{\boldsymbol{\beta}}$.

- **Confidence Interval for Mean Response** $E(Y|X = \mathbf{x}_0)$:

    - Quantifies uncertainty about the average value of Y for a specific $\mathbf{x}_0$.
    - $\hat{y}_0 \pm t_{\alpha/2, n-p-1} \cdot SE(\hat{E}(Y|X = \mathbf{x}_0))$.
    - For simple LR (Slide L2 p.32): $\hat{Y} \pm 1.96\hat{\sigma}\sqrt{\frac{1}{n} + \frac{(x_0 - \bar{x})^2}{\sum(x_i - \bar{x})^2}}$. The 1.96 is $z_{0.025}$.

– In R: `predict(lm.fit, newdata=..., interval="confidence")()`.

- **Prediction Interval for Individual Response Y_0 at $X = \mathbf{x}_0$:**

  – Quantifies uncertainty about a single future observation Y_0 for a specific $\mathbf{x}_0$.

  – Accounts for both uncertainty in $\hat{\boldsymbol{\beta}}$ and irreducible error ϵ. Always wider than CI.

  – $\hat{y}_0 \pm t_{\alpha/2, n-p-1} \cdot \sqrt{RSE^2 + SE(\hat{E}(Y|X = \mathbf{x}_0))^2}$.

  – For simple LR (Slide L2 p.32): $\hat{Y} \pm 1.96\hat{\sigma}\sqrt{1 + \frac{1}{n} + \frac{(x_0-\bar{x})^2}{\sum(x_i-\bar{x})^2}}$.

  – In R: `predict(lm.fit, newdata=..., interval="prediction")()`.

```
1  # Using simple linear model lm.fit <- lm(mpg ~ horsepower, data=Auto)
2  new_hp_values <- data.frame(horsepower=c(98, 150, 200))
3  # Confidence intervals for average mpg
4  predict(lm.fit, newdata=new_hp_values, interval="confidence")
5  # Prediction intervals for individual car's mpg
6  predict(lm.fit, newdata=new_hp_values, interval="prediction")
```

Listing 5: Confidence and Prediction Intervals in R

2.2.7 Model Extensions and Potential Problems

- **Qualitative Predictors:**

  – Create $k-1$ dummy variables for a k-level categorical predictor. One level is baseline.

  – Coefficients are interpreted relative to the baseline.

  – R's `lm()` handles factors automatically. `contrasts()` shows encoding.

- **Interaction Terms:** $X_1 \cdot X_2$. Allows effect of X_1 on Y to depend on the level of X_2.

  – Syntax in R: `Y ~ X1 + X2 + X1:X2` or `Y ~ X1X2`.

  – Hierarchical Principle: If an interaction term $X_1 : X_2$ is included, the main effects X_1 and X_2 should also be included, even if their p-values are not significant.

- **Non-linear Relationships:**

– **Polynomial Regression:** Include $X^2, X^3, \ldots$ as predictors. In R: `lm(Y ~ X + I(X^2))`, or `lm(Y ~ po`

- **Potential Problems in Linear Regression:**

  Non-linearity of Data: If true relationship is non-linear, linear model is a poor fit. Detect with residual plots (residuals vs. fitted values, or residuals vs. predictors). Solution: transformations (log, sqrt), polynomial terms, or non-linear models.

  Correlation of Error Terms (ϵ_i): Standard errors will be underestimated, CIs too narrow, p-values too small. Occurs often with time series data. Detect with Durbin-Watson test, plotting residuals vs. time, or ACF of residuals.

  Non-constant Variance of Errors (Heteroscedasticity): Variance of ϵ_i depends on X_i. Funnel shape in residual vs. fitted plot. Invalidates SEs, CIs, hypothesis tests. Solution: transform Y (e.g., $\log Y$, $\sqrt{Y}$), use weighted least squares.

  Outliers: Observations with y_i far from model prediction. Large residuals. Can unduly influence model. Detect with studentized residuals (residuals divided by estimated SE; values ¿ —3— are suspect). Solution: remove if data entry error, otherwise be cautious.

High-Leverage Points: Observations with unusual x_i values. Leverage statistic h_{ii} (diagonal of hat matrix $\mathbf{X}(\mathbf{X}'\mathbf{X})^{-1}\mathbf{X}'$). h_{ii} between $1/n$ and 1. Average $h_{ii} = (p+1)/n$. Points with h_{ii} much larger than average have high leverage. Can have large impact on $\hat{\beta}$.

Collinearity/Multicollinearity: Two or more predictors are highly correlated. Difficult to separate individual effects. SEs of $\hat{\beta}_j$ become large, t-stats small. Overall F-test might be significant while individual p-values are not. Detect with correlation matrix of predictors, Variance Inflation Factor (VIF). $VIF_j = 1/(1 - R^2_{X_j|X_{-j}})$. VIF ¿ 5 or 10 indicates problem. Solution: drop one correlated variable, combine variables (e.g., PCA), use ridge regression.

```r
# Assuming lm.fit.all <- lm(medv ~ ., data=Boston)
library(car) # for vif function
vif_values <- vif(lm.fit.all)
print(vif_values)
# Example: tax and rad often have high VIF in Boston dataset
```

Listing 6: Checking VIF in R

2.2.8 K-Nearest Neighbors (KNN) Regression

- **Concept**: A non-parametric method that predicts Y for a given x_0 by averaging the Y values of its K "closest" neighbors in the training data.

- **Algorithm for prediction at** x_0 (Slide L2 p.14):

  1. Identify the K training observations (x_i, y_i) that are closest to x_0. This set of neighbors is N_0.
  2. The KNN regression fit for x_0 is $\hat{f}(x_0) = \frac{1}{K} \sum_{x_i \in N_0} y_i$.

- **Distance Metric**:

  - For a single predictor x: $d(x, x_0) = |x - x_0|$.
  - For multiple predictors $\mathbf{x}$: Euclidean distance $\sqrt{\sum_{j=1}^{p} (x_j - x_{0j})^2}$. It's crucial to scale predictors to similar ranges if they are on different scales, otherwise predictors with larger scales will dominate the distance calculation.

- **Choice of K (The Tuning Parameter)**:

  - $K = 1$ (Slide L2 p.17): Lowest bias, highest variance. The prediction function is very flexible and interpolates the training data. Can be very noisy.
  - Large K (e.g., $K = 20$ on Slide L2 p.18, or $K = n$): Higher bias, lower variance. The prediction function becomes smoother. If $K = n$, it predicts the global mean $\bar{y}$ for all x_0.
  - Optimal K: Balances bias and variance to minimize Test MSE. Typically chosen using cross-validation.
  - Slides L2 p.21-24 illustrate this by plotting Squared Bias, Variance, and Test MSE vs. K for a simulated dataset, showing a U-shaped Test MSE curve with an optimal $K \approx 36$.

- **R Example for KNN Regression** (Slides L2 p.15, and extended concept from 'ex3.8.R'):

16

```
1  # Function from slide L2 p.15 (for single predictor)
2  knn_function_single_pred <- function(x0, x_train, y_train, K_val=20) {
3    distances <- abs(x_train - x0)
4    ordered_indices <- order(distances)
5    neighbor_indices <- ordered_indices[1:K_val]
6    predicted_y <- mean(y_train[neighbor_indices])
7    return(predicted_y)
8  }
9
10 # Example usage (conceptual, needs a loop or apply for multiple x0)
11 # library(ISLR) # For Auto dataset
12 # attach(Auto)
13 # x_values_sorted <- sort(horsepower)
14 # y_predictions_knn <- sapply(x_values_sorted, knn_function_single_pred,
15 #                       x_train=horsepower, y_train=mpg, K_val=5)
16 # plot(horsepower, mpg)
17 # lines(x_values_sorted, y_predictions_knn, col="blue", lwd=2)
18 # detach(Auto)
19
20 # For actual KNN, use packages like 'FNN' or 'class' (for classification
     )
21 # library(FNN)
22 # knn.reg(train = as.matrix(horsepower_train), test = as.matrix(
     horsepower_test),
23 #         y = mpg_train, k = 5)
```
Listing 7: Conceptual KNN Regression in R

- **Comparison to Linear Regression**:
  - *Assumptions*: Linear regression assumes a linear functional form. KNN is non-parametric and makes no such assumption.
  - *Performance*:
    * If the true relationship $f(X)$ is close to linear, linear regression usually performs better.
    * If $f(X)$ is highly non-linear, KNN can outperform linear regression, especially if n is large and p (number of predictors) is small.
  - *Curse of Dimensionality*: KNN's performance degrades rapidly as p increases because the "nearest" neighbors can be very far away in high-dimensional space, making the local average less meaningful.
  - *Interpretability*: Linear regression coefficients are easy to interpret. KNN is less interpretable (a black box).

2.2.9 ISLR Chapter 3 Lab Highlights

- **Simple Linear Regression** (using `Boston` data, `medv` vs `lstat`):
  - Fitting: `lm(medv ~ lstat, data=Boston)()`.
  - Output: `summary()`, `names()`, `coef()`, `confint()`.
  - Prediction: `predict()` with `interval="confidence"` and `interval="prediction"`.
  - Plotting: `plot(lstat, medv)()`, `abline(lm.fit)()`.
  - Diagnostics: `par(mfrow=c(2,2))()`, `plot(lm.fit)()` (Residuals vs Fitted, Normal Q-Q, Scale-Location, Residuals vs Leverage). `hatvalues()`.

17

- **Multiple Linear Regression**:

  - Fitting: `lm(medv ~ lstat + age)()`, `lm(medv ~ ., data=Boston)()`.
  - VIF: `vif(lm.fit)()` from `car` package.
  - Updating models: `update(lm.fit, . .-age)()`.

- **Interaction Terms**: `lm(medv ~ lstat*age)`.

- **Non-linear Transformations**: `lm(medv ~ lstat + I(lstat^2))`, `lm(medv ~ poly(lstat, 5))`, `lm(medv ~ log(rm))`.

  - Comparing models: `anova(lm.fit.linear, lm.fit.quadratic)()`.

- **Qualitative Predictors** (using `Carseats` data):

  - Factors are handled automatically. `contrasts(Carseats$ShelveLoc)()`.

2.2.10   Key Insights from ISLR Chapter 3 Exercises

- **ISLR 3.7 (Conceptual)**: Proves that for simple linear regression, $R^2 = (\text{Cor}(X, Y))^2$. The proof involves expanding the definitions of R^2, TSS, RSS, and $\text{Cor}(X, Y)$, and showing their algebraic equivalence using the formulas for $\hat{\beta}_0$ and $\hat{\beta}_1$.

- **ISLR 3.8 (Applied - Auto dataset)**:

  - Fit `lm(mpg ~ horsepower)`.
  - Interpret relationship: Significant negative relationship.
  - Predict `mpg` for `horsepower=98`, get confidence and prediction intervals.
  - Plot data and regression line.
  - Diagnostic plots: Reveal non-linearity (curved pattern in residuals vs. fitted).

- **ISLR 3.9 (Applied - Auto dataset)**:

  - Scatterplot matrix: `pairs(Auto)()`.
  - Correlation matrix: `cor(subset(Auto, select=-name))()`.
  - Multiple regression `lm(mpg ~ . -name, data=Auto)`.
  - Identify significant predictors (displacement, weight, year, origin).
  - Diagnostic plots for multiple regression: Similar non-linearity issues.
  - Explore interaction terms (e.g., `cylinders*displacement`).
  - Explore non-linear transformations (e.g., `log(weight)`, `sqrt(horsepower)`, `I(acceleration^2`
  - A better model might involve transforming the response, e.g., `lm(log(mpg) ~ ..., data=Auto)`, which can improve linearity and homoscedasticity of residuals.

- **ISLR 3.10 (Applied - Carseats dataset)**:

  - Fit `lm(Sales ~ Price + Urban + US)`.
  - Interpret coefficients: Price (negative effect), Urban (not significant), US (positive effect).
  - Model selection: Refit without Urban: `lm(Sales ~ Price + US)`.

- Compare model fit (RSE, R2).
- Confidence intervals for coefficients of the selected model.
- Check for outliers and high leverage points using diagnostic plots.

- **ISLR 3.13, 3.14 (Simulation Exercises)**:
  - Simulate data with known $\beta_0, \beta_1, \sigma^2$.
  - Fit models, compare $\hat{\beta}$ to true β.
  - Observe effect of noise (σ^2) on model fit and confidence intervals.
  - Demonstrate effect of collinearity on coefficient estimates and their significance.
  - Effect of outliers/leverage points.

2.3 Lecture 3: Classification - Logistic Regression

2.3.1 Introduction to Classification

- **Goal**: Predict a qualitative (categorical) response variable, Y.

- Unlike regression (quantitative Y), classification assigns an observation to a class.

- Y can take values in a set of K classes, e.g., $\{1, 2, \ldots, K\}$.

- The probability that Y belongs to class k is p_k, with $\sum_{k=1}^{K} p_k = 1$.

- Models estimate p_k as a function of predictors X_i: $p_k = P(Y_i = k | X_i = x_i) = f_k(x_i)$.

2.3.2 Measuring Prediction Quality in Classification

- **Error Rate**: Proportion of misclassified observations.

- Training Error Rate:
$$\text{trainER} = \frac{1}{n} \sum_{i=1}^{n} I(\hat{y}_i \neq y_i)$$

  where $I(\cdot)$ is the indicator function (1 if true, 0 if false).

- Test Error Rate:
$$\text{testER} = \frac{1}{m} \sum_{i=n+1}^{n+m} I(\hat{y}_i \neq y_i)$$

  Calculated on unseen test data of size m.

2.3.3 Why Not Linear Regression for Classification?

- If Y is binary (0/1), linear regression $\hat{Y} = \hat{\beta}_0 + \hat{\beta}_1 X$ can produce probabilities $\hat{p}(X)$ outside $[0, 1]$.

- The coding of classes (e.g., 0/1 vs. 1/2) affects the linear regression fit, which is undesirable for a classification method.

- Linear probability model (LPM): Using OLS on a 0/1 coded Y. (Slides L3 p. 7-8 illustrate this with `ISLR Default` data, showing predicted probabilities < 0 or > 1).

```
1  # Using Default dataset from ISLR
2  # y is 0/1 coded default status
3  # linprob <- lm(y ~ balance, data=Default)
4  # summary(linprob) # Shows significant relationship
5  # plot(balance, y)
6  # abline(linprob, col="red") # Line goes below 0 and above 1
```
Listing 8: Linear Probability Model Example (Slides L3 p.8)

2.3.4 Logistic Regression Model

- Models the probability $p(X) = P(Y = 1|X)$ using the logistic function:

$$p(X) = \frac{e^{\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p}}{1 + e^{\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p}} = \frac{1}{1 + e^{-(\beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p)}}$$

  This ensures $0 \le p(X) \le 1$. (Slide L3 p.10 shows $p(x)$ for one predictor).

- **Logit Transformation / Log-Odds**:

$$\log\left(\frac{p(X)}{1 - p(X)}\right) = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p$$

  The logit is linear in X. $\frac{p(X)}{1-p(X)}$ is called the odds.

- **Interpretation of Coefficients** (Slide L3 p.11):

  - A one-unit increase in X_j, holding other predictors constant, changes the log-odds by β_j.
  - Equivalently, it multiplies the odds by e^{β_j}.
  - Example (Slide L3 p.11, `Default` data, $Y =$ default, $X =$ studentYes): `glm(default ~ student, family="binomial", data=Default)` Coefficients: (Intercept) -3.5041, studentYes 0.4049. $e^{0.4049} \approx 1.50$. The odds of default for a student are 1.50 times the odds for a non-student.

2.3.5 Estimating Coefficients: Maximum Likelihood Estimation (MLE)

- Goal: Find $\hat{\beta}_0, \hat{\beta}_1, \ldots, \hat{\beta}_p$ that maximize the likelihood of observing the given data.

- **Likelihood Function** for binary $Y_i \in \{0, 1\}$:

$$L(\beta_0, \ldots, \beta_p) = \prod_{i:y_i=1} p(x_i) \prod_{i':y_{i'}=0} (1 - p(x_{i'}))$$

  where $p(x_i) = P(Y_i = 1|X_i = x_i; \beta_0, \ldots, \beta_p)$.

- Often easier to maximize the **log-likelihood function** (Slide L3 p.14-15):

$$\ell(\beta_0, \ldots, \beta_p) = \sum_{i=1}^{n} [y_i \log(p(x_i)) + (1 - y_i) \log(1 - p(x_i))]$$

$$\ell(\beta_0, \beta_1) = \sum_{i=1}^{n} \left[y_i(\beta_0 + \beta_1 x_i) - \log(1 + e^{\beta_0 + \beta_1 x_i})\right]$$

  (for single predictor, derived from the formula on Slide L3 p.15)

- Maximization is done numerically (e.g., Iteratively Reweighted Least Squares - IRLS).

- R: `glm()` with `family=binomial` uses MLE.

20

2.3.6 Making Predictions and Evaluating Accuracy

- Prediction is a two-step process (Slide L3 p.10):

  1. Estimate $P(Y = 1|X = x)$ using the fitted logistic model: $\hat{p}(x)$.
  2. Classify as 1 if $\hat{p}(x) >$ threshold (often 0.5), else 0. (Slide L3 p.9 example: predict default if $p > 0.5$, implies balance ¿ 2000).

- **Confusion Matrix** (Slides L3 p.16-17):

		Predicted Class	
		Negative (0)	Positive (1)
Actual	Negative (0)	(TN)	(FP)
Class	Positive (1)	(FN)	(TP)

- Common Metrics from Confusion Matrix:

  - *Overall Accuracy*: $(TN + TP)/(Total)$
  - *Sensitivity (True Positive Rate, Recall)*: $TP/(TP + FN)$
  - *Specificity (True Negative Rate)*: $TN/(TN + FP)$
  - *Precision (Positive Predictive Value)*: $TP/(TP + FP)$
  - *False Positive Rate*: $FP/(FP + TN) = 1 - \text{Specificity}$

- R Example (`Default` data, predicting `default`, Slides L3 p.17-18):

```r
# Fit model (can be on full data or training data)
# logprob <- glm(default ~ student + balance + income, data=Default,
    family="binomial")
# pred_probs <- predict(logprob, type="response")
# pred_class <- ifelse(pred_probs > 0.5, "Yes", "No") # Using 0.5
    threshold
# conf_matrix_full <- table(Default$default, pred_class)
# print(conf_matrix_full)

# Example with train/test split (Slide L3 p.18)
set.seed(123)
n <- nrow(Default)
train_indices <- sample(1:n, n/2)
train_data <- Default[train_indices,]
test_data <- Default[-train_indices,]

logprob_train <- glm(default ~ student + balance + income, data=train_
    data, family="binomial")
pred_probs_test <- predict(logprob_train, newdata=test_data, type="
    response")
pred_class_test <- ifelse(pred_probs_test > 0.5, "Yes", "No")
conf_matrix_test <- table(test_data$default, pred_class_test)
print(conf_matrix_test)
#    FALSE TRUE
# No   4809   20
# Yes   116   55
accuracy_test <- sum(diag(conf_matrix_test)) / sum(conf_matrix_test)
# accuracy_test is approx 0.9728
```

Listing 9: Logistic Regression and Confusion Matrix (Slides L3 p.17-18)

- **Issue with Imbalanced Classes** (Slide L3 p.19): If one class is much larger (e.g., "No default" is 96.7%), a naive classifier predicting the majority class for all observations can achieve high overall accuracy.

– Example: Predicting "Not default" for all in the test set on Slide L3 p.18 $(4809+20 = 4829$ "No" in test) would give $(4809+20)/(4809+20+116+55) \approx 4829/5000 \approx 0.9658$ if "No" was the only prediction for the No actuals. More accurately, if we only predict "No", accuracy $= (TN + 0) /$ Total $= (4809+116)/(5000)$ if predicting no for "yes" actuals also. The slide implies classifying all 4839 actual "No" correctly, which means a classifier always predicting "No" gets TN=4809, FP=0, FN=116+55=171, TP=0. Acc $= 4809 / 5000 = 0.9618$. The slide's calculation 4839/5000 seems to refer to the number of non-defaulters in the test set, implying a naive rule "predict No Default" would be correct for 4809 true "No"s, if we consider total test obs as 5000. If the test set had 4839 actual "No" and 161 actual "Yes", and we predict "No" always, accuracy $= 4839/5000 = 0.9678$.

– Point: High overall accuracy is misleading here. Sensitivity (correctly identifying "Yes" defaults) is often more important for rare events.

2.3.7 Multiple Logistic Regression

- Model: $\log\left(\frac{p(X)}{1-p(X)}\right) = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p$.

- Estimation and interpretation are analogous to simple logistic regression.

- **Confounding**: Relationship between a predictor and response can be distorted if other relevant predictors are omitted.

2.3.8 Logistic Regression for ¿2 Response Classes (Multinomial Logistic Regression)

- Not explicitly covered in slides, but ISLR discusses it.

- One class is chosen as baseline. Model $K - 1$ log-odds ratios relative to the baseline.

- E.g., for 3 classes (1, 2, 3), baseline class 3:

$$\log\left(\frac{P(Y = 1|X)}{P(Y = 3|X)}\right) = \beta_{01} + \beta_{11}X_1 + \ldots$$

$$\log\left(\frac{P(Y = 2|X)}{P(Y = 3|X)}\right) = \beta_{02} + \beta_{12}X_1 + \ldots$$

- Probabilities sum to 1.

2.3.9 Exercise: ISLR 4.13a-d (or old 4.10a-d)

(Corresponds to 'ex4.13.R' file, first part)

- (a) Explore `Weekly` data: `summary()`, `pairs()`, `cor()`. Look for patterns, e.g., Volume increases over Year. Lags don't show strong correlations with Today's return or Direction.

- (b) Logistic regression `Direction` ~ `Lag1+Lag2+Lag3+Lag4+Lag5+Volume`:

```
# library(ISLR)
# data(Weekly)
# glm.fit <- glm(Direction ~ Lag1 + Lag2 + Lag3 + Lag4 + Lag5 + Volume,
#                data = Weekly, family = binomial)
# summary(glm.fit)
# Output: Lag2 appears to be statistically significant (p-value ~0.03).
```
Listing 10: Logistic Regression on Weekly Data (ISLR 4.13b)

- (c) Confusion Matrix and Interpretation:

```
1  # glm.probs <- predict(glm.fit, type="response")
2  # glm.pred <- rep("Down", length(glm.probs))
3  # glm.pred[glm.probs > 0.5] <- "Up"
4  # conf_matrix <- table(Weekly$Direction, glm.pred)
5  # print(conf_matrix)
6  # #          glm.pred
7  # # Direction Down  Up
8  # #      Down   54  430
9  # #      Up     48  557
10 # accuracy <- mean(glm.pred == Weekly$Direction) # (54+557)/1089 = 0.561
11 # # The model correctly predicts "Up" 557/(48+557) = 92.1% of the time
     it goes Up.
12 # # It correctly predicts "Down" 54/(54+430) = 11.2% of the time it goes
     Down.
13 # # The model is biased towards predicting "Up".
```

Listing 11: Confusion Matrix for Weekly Data Logistic Model (ISLR 4.13c)

- (d) Train/Test Split, fit model with `Lag2` only, evaluate on test:

  - Training: Years ¡ 2009. Test: Years 2009-2010.
  - Fit `glm(Direction   Lag2, data=train`$_data, family = binomial).Predictontestdata, created$
  - *Result in 'ex4.13.R'*: Accuracy ≈ 0.625. This is better than chance (50%) and slightly better than always predicting "Up" for the test period (which would be $61/104 \approx 0.5865$).

2.4   Lecture 4: Classification - Discriminant Analysis & KNN

2.4.1   Linear Discriminant Analysis (LDA)

- Alternative to logistic regression, especially when classes are well-separated or n is small and predictors are approx. normal.

- **Bayes' Theorem for Classification** (Slides L4 p.3-4):

  - Assign observation $X = x$ to class k that maximizes posterior probability $p_k(x) = P(Y = k|X = x)$.
  - $p_k(x) = \frac{\pi_k f_k(x)}{\sum_{l=1}^{K} \pi_l f_l(x)}$
    * $\pi_k = P(Y = k)$: Prior probability of class k. Estimated as proportion of class k in training data ($\hat{\pi}_k = n_k/n$). (Slide L4 p.12)
    * $f_k(x) = P(X = x|Y = k)$: Density of X for an observation from class k.
  - Bayes Classifier: Ideal, but $f_k(x)$ is usually unknown.

- **LDA Assumptions (for one predictor $p = 1$, $K = 2$ classes)** (Slide L4 p.8):

  1. $f_k(x)$ is Gaussian (Normal): $X|Y = k \sim N(\mu_k, \sigma_k^2)$.
  2. LDA specifically assumes common variance: $\sigma_1^2 = \sigma_2^2 = \cdots = \sigma_K^2 = \sigma^2$.

- **LDA Discriminant Function ($\delta_k(x)$)** (Slides L4 p.10-11):

  - Maximize $p_k(x)$ is equivalent to maximizing $\log(p_k(x))$, and further equivalent to maximizing $\delta_k(x)$ after removing terms common to all classes.

- For $p = 1$:

$$\delta_k(x) = x\frac{\mu_k}{\sigma^2} - \frac{\mu_k^2}{2\sigma^2} + \log(\pi_k)$$

- Decision rule: Assign x to class k for which $\delta_k(x)$ is largest.
- The decision boundary between two classes k and l (where $\delta_k(x) = \delta_l(x)$) is linear in x.

- **Parameter Estimation for LDA ($p = 1$)** (Slide L4 p.12):

  - $\hat{\pi}_k = n_k/n$ (proportion of training obs in class k).
  - $\hat{\mu}_k = \frac{1}{n_k}\sum_{i:y_i=k} x_i$ (average of x_i for class k).
  - $\hat{\sigma}^2 = \frac{1}{n-K}\sum_{k=1}^{K}\sum_{i:y_i=k}(x_i - \hat{\mu}_k)^2$ (pooled variance). (Slide L4 p.12 has $\hat{\sigma}_k^2$ for QDA, then $\hat{\sigma}^2$ for LDA is essentially a weighted average if $p = 1$, or derived from pooled covariance for $p > 1$). The provided slide shows $\hat{\sigma}_k^2$ which is for QDA. For LDA with $p = 1$, the formula simplifies to a pooled estimate based on class-wise sums of squares. The formula in the slide for $\hat{\sigma}_k^2$ is for QDA, not LDA's common variance. ISLR p.141 gives common $\hat{\sigma}^2$.

- **LDA for $p > 1$ Predictors** (Slide L4 p.15):

  - Assume $X|Y = k \sim N_p(\boldsymbol{\mu}_k, \boldsymbol{\Sigma})$ (multivariate Gaussian with common covariance matrix $\boldsymbol{\Sigma}$).
  - Discriminant function:

$$\delta_k(\mathbf{x}) = \mathbf{x}'\boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_k - \frac{1}{2}\boldsymbol{\mu}_k'\boldsymbol{\Sigma}^{-1}\boldsymbol{\mu}_k + \log(\pi_k)$$

  - Decision boundaries are linear (hyperplanes).

- **R Implementation** (`MASS` package, Slides L4 p.13-14, 19):

```
# Conceptual R code from slides for LDA (single predictor 'balance' from
    Default data)
# x <- Default$balance
# cl <- Default$default # Factor with levels "No", "Yes"
# nk <- table(cl)
# n <- length(cl)
# pi_hat <- nk/n
# mu_hat <- as.matrix(by(x, cl, mean))
# # Pooled variance s2 (simplified for p=1, assuming K=2 classes)
# s2_no <- sum((x[cl=="No"] - mu_hat[1])^2)
# s2_yes <- sum((x[cl=="Yes"] - mu_hat[2])^2)
# s2_pooled <- (s2_no + s2_yes) / (n - 2) # This is (n-K)
#
# delta_k <- function(x_val, k_idx, mu_vec, s2_val, pi_vec) {
#    # k_idx would be 1 for "No", 2 for "Yes"
#    mu_k <- mu_vec[k_idx]
#    pi_k <- pi_vec[k_idx]
#    return(x_val  mu_k/s2_val - mu_k^2 / (2s2_val) + log(pi_k))
# }
# # To predict: calculate delta1 and delta2 for each x, assign to class
     with larger delta.
# delta1_vals <- delta_k(x, 1, mu_hat, s2_pooled, pi_hat)
# delta2_vals <- delta_k(x, 2, mu_hat, s2_pooled, pi_hat)
# pred_lda_manual <- ifelse(delta2_vals > delta1_vals, "Yes", "No")
# table(cl, pred_lda_manual)
```

```
24
25  # Using MASS::lda() (from ex4.13.R for Weekly data)
26  # library(MASS)
27  # train_indices <- (Weekly$Year < 2009)
28  # train_data <- Weekly[train_indices,]
29  # test_data <- Weekly[!train_indices,]
30  # lda.fit <- lda(Direction ~ Lag2, data=train_data)
31  # lda.pred_obj <- predict(lda.fit, newdata=test_data)
32  # lda.class <- lda.pred_obj$class
33  # conf_matrix_lda <- table(test_data$Direction, lda.class)
34  # print(conf_matrix_lda) # Test accuracy was 0.625
35  # head(lda.pred_obj$posterior) # Shows posterior probabilities
```

Listing 12: LDA on Default Data (Slides L4 p.13-14 conceptual, ex4.13.R actual)

2.4.2 Quadratic Discriminant Analysis (QDA)

- Similar to LDA, but assumes each class k has its own covariance matrix $\mathbf{\Sigma}_k$.

- Assumption: $X|Y = k \sim N_p(\boldsymbol{\mu}_k, \mathbf{\Sigma}_k)$.

- Discriminant function $\delta_k(x)$ becomes quadratic in x:

$$\delta_k(\mathbf{x}) = -\frac{1}{2}\log|\mathbf{\Sigma}_k| - \frac{1}{2}(\mathbf{x} - \boldsymbol{\mu}_k)'\mathbf{\Sigma}_k^{-1}(\mathbf{x} - \boldsymbol{\mu}_k) + \log(\pi_k)$$

- Decision boundaries are quadratic.

- **LDA vs. QDA Trade-off**:

  - LDA: Estimates p parameters for $\boldsymbol{\mu}_k$ (for each class), and $p(p+1)/2$ for common $\mathbf{\Sigma}$. Less flexible.

  - QDA: Estimates p parameters for $\boldsymbol{\mu}_k$ and $p(p+1)/2$ for $\mathbf{\Sigma}_k$ (for each of K classes). More flexible.

  - QDA has higher variance but lower bias (if true decision boundary is non-linear).

  - LDA better if n is small or common covariance assumption is reasonable. QDA better if n is large and common covariance is violated.

- R Implementation (`MASS` package):

```
1  # library(MASS)
2  # qda.fit <- qda(Direction ~ Lag2, data=train_data) # Using same train/
      test as LDA
3  # qda.pred_obj <- predict(qda.fit, newdata=test_data)
4  # qda.class <- qda.pred_obj$class
5  # conf_matrix_qda <- table(test_data$Direction, qda.class)
6  # print(conf_matrix_qda) # Test accuracy was 0.5865 (predicted "Up"
      always)
```

Listing 13: QDA on Weekly Data (from ex4.13.R)

2.4.3 Comparison of Classification Methods (Logistic, LDA, QDA, KNN)

- **Logistic Regression vs. LDA**:

  - Both produce linear decision boundaries (if $p > 1$).

- LDA assumes Gaussian $f_k(x)$ with common Σ. LR makes no such assumption about $f_k(x)$, models $P(Y = k|X)$ directly.
- If Gaussian assumption holds, LDA is more efficient (stable estimates with smaller n).
- LR is generally more robust if Gaussian assumption is violated.
- In practice, often similar performance.
- LDA parameter estimates can be unstable if classes are well-separated (LR too). (Slide L4 p.16)
- LDA more stable than LR for small n if X is approx. multinormal. (Slide L4 p.16)
- LDA is a natural approach for $K > 2$ classes. (Slide L4 p.16)

- **KNN**:

  - Completely non-parametric. Makes no assumptions about decision boundary shape.
  - Can outperform parametric methods if boundary is highly non-linear.
  - Requires large n for good performance, especially if p is large (curse of dimensionality).
  - Optimal K is crucial.

- **QDA**:

  - Compromise between non-parametric KNN and rigid LDA/LR.
  - Assumes Gaussian $f_k(x)$ but allows different Σ_k.
  - Good if true boundary is moderately non-linear and Gaussian assumption is reasonable.

2.4.4 Evaluating Classification Models (Revisited)

- **Threshold Choice** (Slide L4 p.18):

  - Default threshold for binary classification is 0.5 for $P(Y = 1|X)$.
  - Can be adjusted based on costs of misclassification or to balance sensitivity/specificity.
  - Example: For `Default` data, one might use threshold ¡ 0.5 (e.g., 0.2) to identify more potential defaulters, accepting more false positives.

- **ROC Curve (Receiver Operating Characteristics)** (Slides L4 p.19-21):

  - Plots True Positive Rate (Sensitivity) vs. False Positive Rate (1 - Specificity) for various threshold values.
  - A good classifier has ROC curve far from the 45-degree line (random guessing) towards top-left corner.
  - **AUC (Area Under Curve)**: Summary measure of ROC performance.
    * AUC = 1: Perfect classifier.
    * AUC = 0.5: Random guessing.
    * AUC ¿ 0.7 generally considered acceptable, ¿0.8 good, ¿0.9 excellent.

- R code for ROC (conceptual from Slides L4 p.21, using LDA predictions):

```
1  # Assume lda1 fitted on cl ~ x (Default data)
2  # pr <- predict(lda1)$posterior # Posterior probabilities P(Y=k|X)
3  # cl_numeric <- as.numeric(Default$default) - 1 # 0 for No, 1 for Yes
4  #
5  # thrange <- seq(0.01, 0.99, by=0.01) # Thresholds to test
6  # roc_data <- data.frame(FPrate=numeric(length(thrange)),
7  #                        TPrate=numeric(length(thrange)))
8  #
9  # for (i in 1:length(thrange)) {
10 #    th <- thrange[i]
11 #    pred_class <- ifelse(pr[, "Yes"] > th, 1, 0) # Predict "Yes" if P(
       Yes|X) > th
12 #
13 #    TP <- sum(pred_class == 1 & cl_numeric == 1)
14 #    FN <- sum(pred_class == 0 & cl_numeric == 1)
15 #    FP <- sum(pred_class == 1 & cl_numeric == 0)
16 #    TN <- sum(pred_class == 0 & cl_numeric == 0)
17 #
18 #    roc_data$TPrate[i] <- TP / (TP + FN) # Sensitivity
19 #    roc_data$FPrate[i] <- FP / (FP + TN) # 1 - Specificity
20 # }
21 # plot(roc_data$FPrate, roc_data$TPrate, type="l", col="red", lwd=2,
22 #      xlab="False Positive Rate", ylab="True Positive Rate", main="ROC
       Curve")
23 # abline(a=0, b=1, lty=2) # Line of no discrimination
24 # # Packages like 'pROC' or 'ROCR' can do this more easily and calculate
       AUC.
```

Listing 14: Generating ROC Curve Data (Conceptual, Slides L4 p.21)

2.4.5 K-Nearest Neighbors (KNN) Classifier

- To classify a test observation x_0:

    1. Find the K training points closest to x_0 (neighborhood N_0).
    2. Estimate conditional probability for class j: $P(Y = j|X = x_0) = \frac{1}{K}\sum_{i \in N_0} I(y_i = j)$.
    3. Assign x_0 to the class with the largest probability (majority vote among neighbors).

- Choice of K is critical, often selected by CV.

- Small K: Flexible boundary, low bias, high variance.

- Large K: Smoother, less flexible boundary, high bias, low variance.

- R Implementation (`class` package):

```
1  # library(class)
2  # train_X_knn <- as.matrix(train_data$Lag2)
3  # test_X_knn <- as.matrix(test_data$Lag2)
4  # train_Direction_knn <- train_data$Direction
5  #
6  # set.seed(1) # For reproducibility if there are ties
7  # knn.pred_k1 <- knn(train=train_X_knn, test=test_X_knn, cl=train_
       Direction_knn, k=1)
8  # conf_matrix_knn_k1 <- table(test_data$Direction, knn.pred_k1)
9  # print(conf_matrix_knn_k1) # Test accuracy 0.5
10 #
11 # # Trying different K values (e.g., K=10)
```

```
12  # # knn.pred_k10 <- knn(train=train_X_knn, test=test_X_knn, cl=train_
        Direction_knn, k=10)
13  # # accuracy_k10 <- mean(knn.pred_k10 == test_data$Direction)
```

Listing 15: KNN Classification (from ex4.13.R for Weekly data)

2.4.6 Exercises: ISLR 4.13e-i (or old 4.10e-i)

(Corresponds to 'ex4.13.R' file, second part)

- (e) LDA on `Direction`   `Lag2` (train/test split as before):

  - Fit LDA, predict on test set, calculate accuracy.
  - *Result from 'ex4.13.R'*: Accuracy ≈ 0.625. Same as logistic regression with only `Lag2`.

- (f) QDA on `Direction`   `Lag2`:

  - Fit QDA, predict, calculate accuracy.
  - *Result from 'ex4.13.R'*: Accuracy ≈ 0.5865. QDA predicts "Up" for all test observations, performing worse than LDA/logistic.

- (g) KNN with K=1 on `Direction`   `Lag2`:

  - Fit KNN (K=1), predict, calculate accuracy.
  - *Result from 'ex4.13.R'*: Accuracy ≈ 0.50. Performs poorly.

- (h) Compare Logistic, LDA, QDA, KNN(K=1):

  - Logistic regression (with `Lag2`) and LDA performed best and identically (accuracy 0.625).
  - QDA and KNN(K=1) performed worse.

- (i) Experiment with different combinations/transformations/K for KNN:

  - Logistic with `Lag2:Lag1` interaction: Accuracy ≈ 0.5865.
  - LDA with `Lag2:Lag1` interaction: Accuracy ≈ 0.5769.
  - QDA with `Lag2 + sqrt(abs(Lag2))`: Accuracy ≈ 0.5769.
  - KNN with K=10: Accuracy ≈ 0.5769.
  - KNN with K=100: Accuracy ≈ 0.5577.
  - *Conclusion*: Original simpler logistic regression and LDA with only `Lag2` were the best performers among these experiments.

2.5 Lecture 5: Resampling Methods - Cross-Validation

2.5.1 Introduction to Resampling Methods

- **Sampling Distribution** : The distribution of a statistic (e.g., sample mean $\bar{X}$) if we were to repeatedly draw samples from the population.

– Example: Estimating average height $\mu = E(X)$ of Norwegian adults. Each sample of $n = 100$ gives one estimate $\hat{\mu} = \bar{x}$. The collection of these $\bar{x}$ values from many hypothetical samples forms the sampling distribution. (Slides L5 p.3-5 show two such sample means: 183.678 and 180.678).

– **Central Limit Theorem (CLT)** (Slides L5 p.6, 9): For large n, the sampling distribution of $\bar{X}$ is approximately Normal: $\bar{X} \sim N(\mu, \sigma^2/n)$.

– We can estimate this sampling distribution using a single sample by plugging in estimates $\hat{\mu} = \bar{x}$ and $\hat{\sigma} = s$: $\bar{X} \sim N(\bar{x}, s^2/n)$. (Slides L5 p.10-11). This allows for constructing confidence intervals (Slide L5 p.12).

- **Why Resampling?** (Slide L5 p.13):

  – Large sample approximations (like CLT) may not hold for small n or for complex statistics where the asymptotic distribution is hard to derive.

  – Resampling methods mimic sampling from a population by instead repeatedly sampling from our *observed sample data.*

  – Often used to estimate test error $E((Y_i - \hat{f}(X_i))^2)$ or to assess variability of model parameters.

2.5.2 Cross-Validation (CV)

- **Core Idea**:

  1. Split data into a *training set* and a *test (or validation) set*.
  2. Fit (train/estimate) the model on the training set.
  3. Evaluate model performance (e.g., predict) on the test set.

- **Purpose**:

  – *Model Assessment*: Estimating the test error of a final chosen model.

  – *Model Selection*: Choosing the appropriate level of flexibility (e.g., degree of polynomial, value of K in KNN).

- **Types of CV**: Validation Set Approach, Leave-One-Out CV (LOOCV), k-Fold CV.

2.5.3 The Validation Set Approach

- **Procedure** (Slide L5 p.15):

  1. Randomly split the n observations into a training set (e.g., $m \approx n/2$ observations) and a validation/test set ($n - m$ observations).
  2. Fit the model using only the training data to get $\hat{f}$.
  3. Evaluate $\hat{f}$ on the validation set by calculating the test MSE (or other error metric):

  $$\text{Test MSE}_{\text{valid}} = \frac{1}{n - m} \sum_{i \in \text{valid\_set}} (y_i - \hat{f}(x_i))^2$$

- **R Example (`ISLR Auto` data)** (Slides L5 p.16-21):

  – Predict `mpg` from `horsepower`.

```r
library(ISLR)
# names(Auto) # mpg, cylinders, displacement, horsepower, ...
n <- nrow(Auto)
set.seed(1) # For reproducibility of the random split
draw <- sample(1:n, size=floor(n/2)) # Indices for training set
train_data <- Auto[draw, ]
test_data <- Auto[-draw, ]

# Model 1: Linear regression
mod1 <- lm(mpg ~ horsepower, data=train_data)
# summary(mod1) # (Slide L5 p.18 shows an example output)

# Model 2: Polynomial regression (degree 2)
mod2 <- lm(mpg ~ horsepower + I(horsepower^2), data=train_data)
# summary(mod2) # (Slide L5 p.20 shows an example output)

# Predicting on test data and calculating Test MSE
pred1_test <- predict(mod1, newdata=test_data)
mse1_test <- mean((test_data$mpg - pred1_test)^2)
# mse1_test -> 27.36073 (from slide L5 p.21, actual value depends on
    seed)

pred2_test <- predict(mod2, newdata=test_data)
mse2_test <- mean((test_data$mpg - pred2_test)^2)
# mse2_test -> 20.29991 (from slide L5 p.21, actual value depends on
    seed)
# Polynomial model has lower Test MSE in this specific split.
```

Listing 16: Validation Set Approach in R (Slides L5 p.16, 18, 21)

- **Drawbacks of Validation Set Approach** (ISLR Ch 5.1.1, Conceptual Ex 5.3b.i):

  1. *High Variability*: The estimated test MSE can be highly variable depending on which observations end up in the training vs. validation set. Different splits can lead to different conclusions about model performance. (The slide L5 p.21 asks "Did you get the exact same numbers?" highlighting this variability).

  2. *Overestimation of Test Error*: Only a subset of data ($m < n$) is used for training. Models fit on less data tend to perform worse. So, validation set MSE might overestimate the test error of a model fit to the full dataset.

2.5.4 Leave-One-Out Cross-Validation (LOOCV)

- **Procedure** (Slide L5 p.22):

  1. For each observation $i = 1, \ldots, n$:
     - Hold out observation (x_i, y_i) (this is the "test set" of size 1).
     - Fit the model $\hat{f}^{(-i)}$ using the remaining $n - 1$ observations (training set).
     - Predict y_i using $\hat{f}^{(-i)}(x_i)$ and calculate the squared error: $MSE_i = (y_i - \hat{f}^{(-i)}(x_i))^2$.

  2. The LOOCV estimate of test MSE is the average of these n errors:

  $$CV_{(n)} = \text{LOOCV MSE} = \frac{1}{n} \sum_{i=1}^{n} MSE_i$$

- **R Example (Auto data)** (Slide L5 p.23):

```
1  # MSE1_loo <- numeric(n) # For linear model
2  # MSE2_loo <- numeric(n) # For quadratic model
3  #
4  # for(i in 1:n) {
5  #   train_loo <- Auto[-i, ]
6  #   test_loo <- Auto[i, ]
7  #
8  #   mod1_loo <- lm(mpg ~ horsepower, data=train_loo)
9  #   pred1_loo <- predict(mod1_loo, newdata=test_loo)
10 #   MSE1_loo[i] <- (test_loo$mpg - pred1_loo)^2
11 #
12 #   mod2_loo <- lm(mpg ~ horsepower + I(horsepower^2), data=train_loo)
13 #   pred2_loo <- predict(mod2_loo, newdata=test_loo)
14 #   MSE2_loo[i] <- (test_loo$mpg - pred2_loo)^2
15 # }
16 #
17 # mean_MSE1_loo <- mean(MSE1_loo) # Slide says 24.23
18 # mean_MSE2_loo <- mean(MSE2_loo) # Slide says 19.25 (Quadratic is
      better)
19
20 # Using boot::cv.glm() for LOOCV with linear models (more efficient)
21 # library(boot)
22 # glm.fit.linear <- glm(mpg ~ horsepower, data=Auto)
23 # cv.err.linear <- cv.glm(Auto, glm.fit.linear) # Default is LOOCV
24 # print(cv.err.linear$delta[1]) # First element is raw LOOCV MSE
25 #
26 # glm.fit.quad <- glm(mpg ~ poly(horsepower, 2), data=Auto) # poly() is
      preferred
27 # cv.err.quad <- cv.glm(Auto, glm.fit.quad)
28 # print(cv.err.quad$delta[1])
```

Listing 17: LOOCV for Auto Data (Conceptual from Slide L5 p.23)

- **Advantages of LOOCV**:

  - *Less Bias*: Uses $n-1$ observations for training in each fold, so $\hat{f}^{(-i)}$ is very similar to $\hat{f}$ fit on all n observations. LOOCV MSE tends to be an almost unbiased estimate of test error for a model trained on n observations.

  - *No Randomness*: Result is always the same, no variability due to random splits.

- **Disadvantages of LOOCV** (ISLR Ch 5.1.2, Conceptual Ex 5.3b.ii):

  - *Computationally Expensive*: Model must be fit n times. Can be very slow for complex models or large n. (For OLS linear regression and polynomial regression, a computational shortcut exists: $CV_{(n)} = \frac{1}{n} \sum (\frac{y_i - \hat{y}_i}{1 - h_{ii}})^2$, where h_{ii} is leverage of i-th obs. So, only need to fit full model once).

  - *High Variance*: The n training sets are highly overlapping (share $n-2$ observations). The MSE_i values are highly correlated, so their average ($CV_{(n)}$) can have high variance. It can be a poor estimator of the true test MSE.

2.5.5   k-Fold Cross-Validation

- **Procedure**:

  1. Randomly divide the n observations into k non-overlapping groups (folds) of approximately equal size (n/k).

2. For each fold $j = 1, \ldots, k$:

    − Hold out fold j as the test set.

    − Fit the model $\hat{f}^{(-j)}$ using the other $k - 1$ folds as the training set.

    − Calculate $MSE_j = \frac{1}{n_j} \sum_{i \in \text{fold } j} (y_i - \hat{f}^{(-j)}(x_i))^2$, where n_j is size of fold j.

3. The k-fold CV estimate of test MSE is:

$$CV_{(k)} = \frac{1}{k} \sum_{j=1}^{k} MSE_j$$

- Common choices for k: 5 or 10. (Slide L5 p.24 notes $k = n$ is LOOCV).

- **Advantages over LOOCV**:

  - *Computationally Cheaper*: Model fit only k times.

  - *Lower Variance*: Training sets for each fold are less overlapping than in LOOCV, leading to less correlated MSE_j values. $CV_{(k)}$ often has lower variance than $CV_{(n)}$.

- **Bias-Variance Trade-off for Choice of k**:

  - *Bias*: $CV_{(k)}$ uses training sets of size $n(k-1)/k$, which is smaller than $n - 1$ (for LOOCV). So $CV_{(k)}$ might have slightly more bias as an estimator of test error for a model fit on full data, compared to LOOCV.

  - *Variance*: $CV_{(k)}$ generally has lower variance than LOOCV.

  - $k = 5$ or $k = 10$ often strike a good balance.

- **CV for Classification Problems**: Use error rate instead of MSE.

$$CV_{(k)} = \frac{1}{k} \sum_{j=1}^{k} \text{Err}_j = \frac{1}{k} \sum_{j=1}^{k} \left(\frac{1}{n_j} \sum_{i \in \text{fold } j} I(y_i \neq \hat{y}_i) \right)$$

- R Implementation (using `boot::cv.glm()`):

```r
# library(boot)
# For k-fold CV, specify K argument in cv.glm
# Example: 10-fold CV for linear model on Auto data
# glm.fit.auto <- glm(mpg ~ horsepower, data=Auto)
# cv.err.10fold <- cv.glm(Auto, glm.fit.auto, K=10)
# print(cv.err.10fold$delta[1]) # Raw 10-fold CV MSE

# Example: 10-fold CV for logistic regression on Default data (ISLR Ch5
    Ex5)
# glm.fit.default <- glm(default ~ income + balance, data=Default,
    family=binomial)
# cv.err.default.10fold <- cv.glm(Default, glm.fit.default, K=10) # Cost
    function needed for error rate
# To get error rate for classification, must define a cost function for
    cv.glm
# Or implement k-fold CV manually:
# k <- 10
# folds <- sample(cut(seq(1,nrow(Default)),breaks=k,labels=FALSE))
# cv_errors <- numeric(k)
# for(j in 1:k){
#    test_indices <- which(folds==j, arr.ind=TRUE)
#    test_data_fold <- Default[test_indices, ]
```

```
19  #    train_data_fold <- Default[-test_indices, ]
20  #    fit_fold <- glm(default ~ income + balance, data=train_data_fold,
    family=binomial)
21  #    probs_fold <- predict(fit_fold, newdata=test_data_fold, type="
    response")
22  #    preds_fold <- ifelse(probs_fold > 0.5, "Yes", "No")
23  #    cv_errors[j] <- mean(preds_fold != test_data_fold$default)
24  # }
25  # mean_cv_error_10fold <- mean(cv_errors)
```

Listing 18: k-Fold CV using boot::cv.glm

2.5.6 Pros and Cons of Different CV Approaches

Table 1: Comparison of Cross-Validation Methods (Slide L5 p.25, ISLR Ch 5.1.4)

Property	Validation Set	LOOCV	k-Fold CV
Bias of Test Error Estimate	Higher	Lower (Approx. Unbiased)	Intermediate
Variance of Test Error Estimate	Higher	Higher	Lower
Computational Time	Best (1 fit)	Worst (n fits)	Intermediate (k fits)

- The table summarizes that LOOCV is good for low bias in estimating test error but can have high variance and is computationally expensive. k-Fold CV (e.g., $k = 5$ or 10) often provides a better bias-variance trade-off for the estimate of test error itself and is computationally more feasible. The validation set approach is simple but can be unreliable.

2.5.7 The Right and Wrong Way to do Cross-Validation

- **Crucial Point**: Any variable selection, feature engineering, or parameter tuning step that relies on the outcomes (y_i) must be performed *inside* the CV loop, using only the training data of that specific fold.

- **Wrong Way**: Perform variable selection on the entire dataset first, then use CV to estimate the error of the chosen model. This "leaks" information from the test folds into the model selection process, leading to an overly optimistic (too low) estimate of test error.

- **Right Way**: In each CV fold j:

  1. Perform variable selection (or other tuning) using *only the training part* of fold j.
  2. Fit the model selected in step 1 using *only the training part* of fold j.
  3. Evaluate this model on the *test part* (hold-out portion) of fold j.

2.5.8 Exercise: ISLR 5.5 (Cross-validation for Logistic Regression)

- (a) Fit logistic regression `default   income + balance` on full dataset.

- (b) Validation Set Approach:

  - Split data 50/50.
  - Fit model on training.

- Predict on validation, calculate test error rate (misclassification).
- *Result from 'ex5.5.R'*: Test error rate $\approx 2.36\%$ to 2.86% depending on split.

- (c) Repeat (b) multiple times: Shows variability of validation set error.

- (d) Validation Set with `student` dummy: Add `student` predictor.

  - Fit `default` `income + balance + student` on training.
  - Test error rate $\approx 2.64\%$. Does not seem to improve much over model without `student` using this single validation split. (More robust comparison would require repeated CV or k-fold CV).

- (Extra Task from 'ex5.5.R': LOOCV on smaller subset)

  - Illustrates manual LOOCV loop for logistic regression on 500 observations.
  - For logistic regression, `cv.glm()` with default settings (or `K=nrow(data)`) performs LOOCV but requires a custom cost function to output error rate instead of deviance for binomial models.

```
# n_subset <- 500
# subset_data <- Default[sample(1:nrow(Default), n_subset), ]
# loocv_errors <- numeric(n_subset)
# for(i in 1:n_subset) {
#    train_fold <- subset_data[-i, ]
#    test_fold <- subset_data[i, ]
#    fit_fold <- glm(default ~ income + balance + student,
#                    data=train_fold, family=binomial)
#    prob_fold <- predict(fit_fold, newdata=test_fold, type="response")
#    pred_fold <- ifelse(prob_fold > 0.5, "Yes", "No")
#    loocv_errors[i] <- (pred_fold != test_fold$default)
# }
# mean_loocv_error <- mean(loocv_errors)
```

Listing 19: Manual LOOCV for Classification Error (Conceptual)

2.6 Lecture 6: Resampling Methods - The Bootstrap

2.6.1 Introduction to the Bootstrap

  * **Purpose**: A powerful and widely applicable tool for quantifying uncertainty associated with a given estimator or statistical learning method.
    · Primarily used to estimate the standard error (SE) of an estimator.
    · Can also be used to construct confidence intervals.
    · Useful when the true sampling distribution of a statistic is unknown or difficult to derive analytically (e.g., for medians, quantiles, complex model parameters).
  * **Core Idea**: Mimic the process of obtaining new sample sets from the population by repeatedly sampling *with replacement* from the original observed dataset.
  * Each "bootstrap sample" has the same size n as the original dataset.
  * Some observations from the original dataset may appear multiple times in a bootstrap sample, while others may not appear at all.

2.6.2 The Bootstrap Procedure for Estimating Standard Error

1. Let the original dataset be $Z = \{z_1, z_2, \ldots, z_n\}$.
2. Generate B independent bootstrap samples $Z^1, Z^2, \ldots, Z^B$. Each Z^b is obtained by drawing n observations from Z with replacement.
3. For each bootstrap sample Z^b, compute the statistic of interest, $\hat{\alpha}^b$. This could be a sample mean, median, regression coefficient, etc.
4. The bootstrap estimate of the standard error of $\hat{\alpha}$ (the statistic computed on the original data) is the standard deviation of the B bootstrap estimates:

$$
SE_B(\hat{\alpha}) = \sqrt{\frac{1}{B-1} \sum_{b=1}^{B} \left(\hat{\alpha}^b - \frac{1}{B} \sum_{r=1}^{B} \hat{\alpha}^r \right)^2}
$$

5. A large number of bootstrap samples B (e.g., $B = 1000$ or more) is typically used.

2.6.3 Example: Estimating SE of Investment Allocation

* Task: Minimize variance of portfolio $\alpha X + (1 - \alpha)Y$.
* Optimal $\alpha = \frac{\sigma_Y^2 - \sigma_{XY}}{\sigma_X^2 + \sigma_Y^2 - 2\sigma_{XY}}$.
* Suppose we have $n = 100$ observations of (X, Y). We can estimate $\hat{\sigma}_X^2, \hat{\sigma}_Y^2, \hat{\sigma}_{XY}$ from this sample, and then $\hat{\alpha}$.
* To find $SE(\hat{\alpha})$ using bootstrap:
    1. Draw B bootstrap samples (X^b, Y^b) of size $n = 100$ from the original 100 observations.
    2. For each bootstrap sample b, calculate $\hat{\alpha}^b$ using the sample variances and covariance from that bootstrap sample.
    3. Compute $SE_B(\hat{\alpha})$ as the standard deviation of these B values of $\hat{\alpha}^b$.
* This provides a measure of how much $\hat{\alpha}$ would vary if we repeatedly sampled 100 observations from the true population.

2.6.4 R Implementation of Bootstrap

* The `boot` package provides the `boot()` function for easy bootstrap implementation.
* **Key arguments for `boot()`** (ISLR Ch 5 Ex 5.6):
    * `data`: The original dataset.
    * `statistic`: A function that takes two arguments: `data` and `index`. The `index` argument will contain the indices of the observations selected for a particular bootstrap sample. The function should return the statistic(s) of interest computed on `data[index,]`.
    * `R`: The number of bootstrap replicates B.
* Example: Bootstrap SE for coefficients of logistic regression on `Default` data (ISLR Ch 5 Ex 5.6).

```
1  library(ISLR) # For Default dataset
2  library(boot)   # For boot() function
3
4  # (a) Fit the logistic regression model once on original data
```

```
5  glm.fit.full <- glm(default ~ income + balance, data = Default,
       family = binomial)
6  # summary(glm.fit.full)$coef # Shows original estimates and SEs
       from asymptotic theory
7
8  # (b) Define the statistic function for bootstrap
9  # This function will be called by boot() for each bootstrap
       sample.
10 # It needs to return the coefficients of the logistic regression
        model.
11 boot.fn.coeffs <- function(data, index) {
12   # Fit model on the bootstrap sample (data[index,])
13   fit <- glm(default ~ income + balance, data = data[index, ],
         family = binomial)
14   return(coef(fit)) # Return the vector of estimated
         coefficients
15 }
16
17 # (c) Run the bootstrap
18 set.seed(1) # For reproducibility
19 boot_results <- boot(data = Default, statistic = boot.fn.coeffs,
        R = 1000) # B=1000 replicates
20
21 # Print the bootstrap results
22 print(boot_results)
23 # Output includes:
24 # original: Coefficients from the model fit on the original data
25 # bias: Difference between mean of bootstrap estimates and
       original estimate
26 # std. error: Bootstrap estimate of SE for each coefficient
27
28 # (d) Comparison: Bootstrap SEs are usually close to SEs from
       summary(glm.fit.full),
29 # but bootstrap is more general and doesn't rely on asymptotic
       theory.
30 # E.g., from Ex 5.6 solution:
31 # Original glm SE for income: 4.99e-06, balance: 2.27e-04
32 # Bootstrap SE for income:  ~4.5e-06, balance: ~2.3e-04 (approx
        from R=50 in book example)
```

Listing 20: Bootstrap for Logistic Regression Coefficients (ISLR Ch 5 Ex 5.6)

* Example: Bootstrap SE for the sample median (ISLR Ch 5 Ex 5.9f, Boston data).

```
1  library(MASS) # For Boston dataset
2  library(boot)
3  # attach(Boston) # or use Boston$medv
4
5  # (e) Calculate median on original data
6  medv.median.original <- median(Boston$medv) # 21.2
7
8  # (f) Define statistic function and run bootstrap
9  boot.fn.median <- function(data_vector, index) {
10   return(median(data_vector[index]))
11 }
12 set.seed(1)
13 boot_median_results <- boot(data = Boston$medv, statistic = boot
       .fn.median, R = 1000)
14 print(boot_median_results)
15 # Bootstrap Statistics :
```

```
16  #      original    bias     std. error
17  # t1      21.2  -0.0098        0.3874 (from ISLR solution)
18  # The bootstrap SE for the median is ~0.387.
19  # detach(Boston)
```
Listing 21: Bootstrap SE for Median (ISLR Ch 5 Ex 5.9f)

2.6.5   Bootstrap Confidence Intervals

∗ Standard $(1 - \alpha)\%$ Normal-based CI: $\hat{\alpha} \pm z_{1-\alpha/2} \cdot SE_B(\hat{\alpha})$. Relies on $\hat{\alpha}$ being approx. Normal and $SE_B(\hat{\alpha})$ being a good estimate.

∗ **Percentile Confidence Interval**: A more direct method.
  1. Obtain B bootstrap estimates $\hat{\alpha}^1, \dots, \hat{\alpha}^B$.
  2. Sort them: $\hat{\alpha}^{(1)} \le \hat{\alpha}^{(2)} \le \cdots \le \hat{\alpha}^{(B)}$.
  3. A $(1-\alpha)\%$ percentile CI is $[\hat{\alpha}^{(\lfloor B\alpha/2 \rfloor)}, \hat{\alpha}^{(\lceil B(1-\alpha/2) \rceil)}]$.
  4. E.g., for a 95% CI ($\alpha = 0.05$) with $B = 1000$, use the 25th and 975th sorted bootstrap estimates.

∗ Other types: Bias-Corrected and Accelerated (BCa) intervals (often better but more complex).

∗ R: `boot.ci()` function from `boot` package can compute various CIs.

```
1  # Assuming boot_results from logistic regression example above
2  # conf_intervals <- boot.ci(boot_results, type = c("norm", "perc
     ", "bca"), index = 2) # For 'income' coeff
3  # print(conf_intervals)
4  # index=1 for intercept, index=2 for income, index=3 for balance
5
6  # For median example:
7  # boot.ci(boot_median_results, type=c("norm", "perc", "bca"))
```
Listing 22: Bootstrap Confidence Intervals using boot.ci

2.6.6   When is Bootstrap Useful?

∗ When it's hard to derive SEs or CIs analytically (e.g., for medians, quantiles, ratios of parameters, complex model outputs).

∗ For small sample sizes where asymptotic theory for SEs might not hold well.

∗ To check assumptions of simpler SE formulas.

2.6.7   Limitations of Bootstrap

∗ **Assumption**: The empirical distribution (from the original sample) is a good approximation of the true population distribution. This may not hold if the sample is very small or not representative.

∗ **Extreme Quantiles/Parameters**: Bootstrap may perform poorly for estimating properties related to the tails of a distribution if the original sample doesn't capture those tails well.

∗ **Computational Cost**: Can be intensive if B is large and the statistic computation is complex.

∗ **Not a substitute for more data**: Bootstrap quantifies uncertainty based on the current sample; it doesn't improve the point estimate itself in the way more data would.

* **Dependence**: Standard bootstrap assumes i.i.d. observations. Modifications exist for time series (e.g., block bootstrap) or dependent data, but are more complex.

2.6.8 Conceptual Exercise Insights

* **(a,b)** Probability a specific observation j is not in a bootstrap sample: $(1-1/n)$.
* **(c)** Probability observation j is not in a bootstrap sample of size n: $(1 - 1/n)^n$.
* **(d,e,f,g)** As $n \to \infty$, $(1 - 1/n)^n \to e^{-1} \approx 0.368$. So, probability observation j is in the bootstrap sample approaches $1 - e^{-1} \approx 0.632$.
    · This means, on average, about 63.2% of original observations are present in any given bootstrap sample (these are the "in-bag" samples). The remaining 36.8

```
pr_in_bootstrap <- function(n) {
  return(1 - (1 - 1/n)^n)
}
n_values <- c(5, 100, 10000)
sapply(n_values, pr_in_bootstrap)
# For n=5:   0.67232
# For n=100: 0.63397
# For n=10000:0.63214
# Converges to 1 - exp(-1)
# 1 - exp(-1) # approx 0.63212
```

Listing 23: Probability an Observation is in a Bootstrap Sample

2.6.9 Exercise: ISLR 5.9 (Bootstrap for `Boston` data)

(Corresponds to 'ch5-applied.R' Exercise 9)

* **(a)** Estimate $\hat{\mu} = \text{mean}(\texttt{medv})$.
* **(b)** Estimate $SE(\hat{\mu})$ using formula $s/\sqrt{n}$.

```
# library(MASS)
# data(Boston)
# mu_hat_medv <- mean(Boston$medv) # 22.53281
# se_mu_hat_formula <- sd(Boston$medv) / sqrt(nrow(Boston)) #
    0.4088611
```

Listing 24: Mean and SE of Mean for Boston *medv*

* **(c)** Estimate $SE(\hat{\mu})$ using bootstrap.

```
# library(boot)
# boot.fn_mean <- function(data_vector, index) {
#    return(mean(data_vector[index]))
# }
# set.seed(1)
# boot_mean_results <- boot(Boston$medv, boot.fn_mean, R=1000)
# print(boot_mean_results) # Bootstrap SE is ~0.4119 (close to
    formula)
```

Listing 25: Bootstrap SE of Mean for Boston *medv*

* **(d)** Compare SEs and construct 95% CI for μ.
    · Bootstrap SE is similar to formula-based SE.

· CI from `t.test(Boston$medv)()`: e.g., [21.73, 23.34].

· Bootstrap CI (e.g., percentile or BCa from `boot.ci()`) should be similar.

∗ (e) Estimate $\hat{\mu}_{med} = $ median(`medv`). (Result: 21.2)

∗ (f) Estimate $SE(\hat{\mu}_{med})$ using bootstrap. (Result SE ≈ 0.38, shown in earlier R example).

∗ (g) Estimate $\hat{\mu}_{0.1} = $ 10th percentile of `medv`. (Result: 12.75)

∗ (h) Estimate $SE(\hat{\mu}_{0.1})$ using bootstrap.

```
# boot.fn_q10 <- function(data_vector, index) {
#     return(quantile(data_vector[index], probs=0.1))
# }
# set.seed(1)
# boot_q10_results <- boot(Boston$medv, boot.fn_q10, R=1000)
# print(boot_q10_results) # Bootstrap SE is ~0.5113
```

Listing 26: Bootstrap SE for 10th Percentile of Boston*medv*

2.7 Lecture 7: Linear Model Selection Methods

2.7.1 Introduction to Model Selection

∗ Model: $Y = \beta_0 + \beta_1 X_1 + \cdots + \beta_p X_p + \epsilon$.

∗ **Conflicting Goals**:

· *Prediction Accuracy*: We want a model that predicts well on new, unseen data.

· *Model Interpretability*: We want a simpler model that is easier to understand and explain.

∗ **Challenge with many predictors (p large relative to n)**:

· Potentially more information about $Y \implies$ better predictions.

· Estimating many parameters $\beta_j \implies$ high estimation uncertainty (variance) $\implies$ poor predictions.

· This leads to a trade-off:

· *Prediction Bias*: Due to an incorrectly specified model (e.g., omitting important variables).

· *Prediction Variance*: Due to estimating too many parameters from limited data.

∗ **Goal of Model Selection**: Find a model that optimally balances this bias-variance trade-off to achieve good prediction accuracy and/or interpretability.

2.7.2 Three Types of Model Selection Techniques

1. **Subset Selection**: Identify a subset of the p predictors that we believe are most relevant to the response. Then fit a model using OLS on this reduced set.

   ∗ Examples: Best Subset Selection, Forward Stepwise, Backward Stepwise.

2. **Shrinkage (Regularization)**: Fit a model involving all p predictors, but the estimated coefficients are constrained or "shrunk" towards zero relative to the OLS estimates. This reduces variance.

   ∗ Examples: Ridge Regression, Lasso. (Covered in Lecture 8)

3. **Dimension Reduction**: Project the p predictors into an M-dimensional sub-space, where $M < p$. Then fit a linear regression model using these M projections as predictors.

   * Examples: Principal Components Regression (PCR), Partial Least Squares (PLS).

2.7.3   Best Subset Selection

* **Algorithm** (Slide L6 p.4):

   1. Let $\mathcal{M}_0$ be the null model (intercept only, no predictors): $Y = \beta_0 + \epsilon$.
   2. For $k = 1, 2, \ldots, p$:
      · Fit all $\binom{p}{k}$ models that contain exactly k predictors.
      · Pick the best among these $\binom{p}{k}$ models based on some criterion (e.g., highest R^2 or lowest RSS for a fixed k). Call this model $\mathcal{M}_k$.
   3. Select the single best model from $\mathcal{M}_0, \mathcal{M}_1, \ldots, \mathcal{M}_p$ using a criterion that accounts for model complexity/overfitting (e.g., Adjusted R^2, C_p, BIC, AIC, or cross-validated test MSE). Cannot use raw R^2 or RSS for this final step as they always improve with more predictors.

* **Computational Cost**: Involves fitting 2^p models. If $p = 50$, this is $> 10^{15}$ models, which is computationally infeasible. Generally feasible for $p \leq 30 - 40$.

* **R Example** (`ISLR Auto data`) (Slides L6 p.5-11, using $p = 6$ predictors):
   · Predictors: cylinders, displacement, horsepower, weight, acceleration, age (derived from year).
   · *Full Model* (Slide L6 p.6): `lm(mpg ., data=Auto_modified)`. $AdjustedR^2 = 0.8063$.
   · `Step 1: Fit` $\mathcal{M}_0$ (Slide L6 p.7): `M0=lm(y 1)`. RSS for $\mathcal{M}_0$ is TSS.
   · `Step 2: Find best model` $\mathcal{M}_k$ `for each` k (Slide L6 p.8 for $k = 2$): For $k = 2$, the model with weight and age has the lowest RSS (4568.95) among all 2-predictor models.
   · `Step 3: Select overall best model` (Slide L6 p.9): Using Adjusted R^2, the model with weight and age ($\mathcal{M}_2$) has the highest Adj $R^2 = 0.8072$.

```
1  # require(ISLR)
2  # Auto$age <- 83 - Auto$year # Create age variable
3  # Auto_subset <- Auto[, !(names(Auto) %in% c("name", "origin
      ", "year"))] # Select relevant columns
4
5  # require(leaps)
6  # regfit.full <- regsubsets(mpg ~ ., data=Auto_subset, nvmax
      = 6) # nvmax = p
7  # reg.summary <- summary(regfit.full)
8  # print(reg.summary)
9  # names(reg.summary) # Shows "which", "rsq", "rss", "adjr2",
      "cp", "bic"
10
11 # Results from summary (Slide L6 p.10)
12 # cbind(reg.summary$which[,-1], adjR2=round(reg.summary$
      adjr2,4))
13 # Shows which variables are in the best model of each size,
      and its AdjR2.
```

40

```
14  # ##    cylinders displacement horsepower weight acceleration
        age   adjR2
15  # ## 1     FALSE        FALSE      FALSE    TRUE        FALSE
        FALSE 0.6918 (weight)
16  # ## 2     FALSE        FALSE      FALSE    TRUE        FALSE
        TRUE  0.8072 (weight, age)
17  # ## 3     FALSE        FALSE      FALSE    TRUE        TRUE
        TRUE  0.8071 (weight, acc, age)
18  # ## ...
19
20  # plot(regfit.full, scale="adjr2", col=gray.colors(10)) #
        Visual (Slide L6 p.11)
21  # The plot shows that AdjR2 peaks for the model with 'weight
        ' and 'age'.
```

Listing 27: Best Subset Selection with leaps::regsubsets (Slides L6 p.10-11)

* **Example with 3 predictors** (X_1, X_2, X_3) (Slides L6 p.18-19):
  · $\mathcal{M}_0 : Y = \beta_0 + \epsilon$
  · $\mathcal{M}_1$: Best of $\{X_1\}, \{X_2\}, \{X_3\}$ (based on R^2)
  · $\mathcal{M}_2$: Best of $\{X_1, X_2\}, \{X_1, X_3\}, \{X_2, X_3\}$ (based on R^2)
  · $\mathcal{M}_3 : Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \beta_3 X_3 + \epsilon$
  · Choose among $\mathcal{M}_0, \mathcal{M}_1, \mathcal{M}_2, \mathcal{M}_3$ using Adj R^2, C_p, BIC, or CV test MSE.
  · Total models estimated: $2^3 = 8$. (Slide L6 p.20)

2.7.4 Stepwise Selection Methods

* Computationally cheaper alternatives to best subset selection, especially for large p.

* **Forward Stepwise Selection** (Slides L6 p.12):
  1. Start with null model $\mathcal{M}_0$ (intercept only).
  2. For $k = 0, \dots, p-1$:
     · Consider all $p - k$ models that add one additional predictor to $\mathcal{M}_k$.
     · Choose the best among these $p - k$ models (e.g., one with lowest RSS or highest R^2). Call this $\mathcal{M}_{k+1}$.
  3. Select the best model among $\mathcal{M}_0, \dots, \mathcal{M}_p$ using C_p, BIC, Adj R^2, or CV.
  4. Total models: $1 + \sum_{k=0}^{p-1}(p-k) = 1 + p(p+1)/2$. For $p = 50$, this is $1 + 1275 = 1276$ models.

* **R Example (Forward Stepwise for `Auto` data)** (Slides L6 p.13-14):

```
1  # regfit.fwd <- regsubsets(mpg ~ ., data=Auto_subset, nvmax =
       6, method="forward")
2  # summary.fwd <- summary(regfit.fwd)
3  # cbind(summary.fwd$which[,-1], adjR2=round(summary.fwd$adjr2
       ,4))
4  # # Results are identical to Best Subset for this Auto example
5  # plot(regfit.fwd, scale="adjr2", col=gray.colors(10))
6  # Number of models estimated: 1 (null) + 6+5+4+3+2+1 = 22.
       Slide says 21, likely not counting M0 in sum.
```

Listing 28: Forward Stepwise with leaps::regsubsets (Slides L6 p.13)

* **Backward Stepwise Selection** (Slides L6 p.15):

1. Start with full model $\mathcal{M}_p$ (all p predictors).
2. For $k = p, p-1, \ldots, 1$:
   - Consider all k models that remove one predictor from $\mathcal{M}_k$.
   - Choose the best among these k models (e.g., lowest RSS or highest R^2). Call this $\mathcal{M}_{k-1}$.
3. Select the best model among $\mathcal{M}_0, \ldots, \mathcal{M}_p$ using C_p, BIC, Adj R^2, or CV.
4. Total models: $1 + \sum_{k=1}^{p} k = 1 + p(p+1)/2$. Requires $n > p$ for initial full model.

* **R Example (Backward Stepwise for `Auto` data)** (Slides L6 p.16-17):

```
# regfit.bwd <- regsubsets(mpg ~ ., data=Auto_subset, nvmax =
    6, method="backward")
# summary.bwd <- summary(regfit.bwd)
# cbind(summary.bwd$which[,-1], adjR2=round(summary.bwd$adjr2
    ,4))
# # Results are identical to Best Subset and Forward for this
    Auto example
# plot(regfit.bwd, scale="adjr2", col=gray.colors(10))
```

Listing 29: Backward Stepwise with leaps::regsubsets (Slides L6 p.16)

* **Hybrid Approaches**: Combine forward and backward steps (e.g., add predictors then remove some if they become non-significant).

* Stepwise methods are not guaranteed to find the true best model out of all 2^p possibilities but are often good approximations and computationally efficient.

2.7.5 Choosing the Optimal Model (Criteria)

* After generating a sequence of models $\mathcal{M}_0, \ldots, \mathcal{M}_p$ (either from best subset or stepwise), we need to select one.

* Training error (RSS, R^2) is not suitable as it always improves with more variables.

* We need to estimate test error. Two approaches:

1. **Adjusting Training Error for Model Size**: Add a penalty for number of predictors (d).
   - **Mallow's C_p** (Slide L6 p.21): For OLS with estimated $\hat{\sigma}^2$ from full model.
     $$C_p = \frac{1}{n}(RSS_d + 2d\hat{\sigma}^2)$$
     Choose model with smallest C_p. If $C_p \approx d$, model has low bias.
   - **Akaike Information Criterion (AIC)** (Slide L6 p.21): Proportional to
     $$AIC \propto \frac{1}{n\hat{\sigma}^2}(RSS_d + 2d\hat{\sigma}^2)$$
     (Derived from maximizing log-likelihood minus a penalty $2d$). Choose model with smallest AIC. For linear models with Gaussian errors, C_p and AIC are proportional.
   - **Bayesian Information Criterion (BIC)** (Slide L6 p.21):
     $$BIC = \frac{1}{n}(RSS_d + \log(n)d\hat{\sigma}^2)$$

(For linear models with Gaussian errors, assuming $\hat{\sigma}^2$ is error variance).
BIC penalizes model size more heavily than AIC/Cp for $n > 7$. Tends
to select smaller models. Choose model with smallest BIC.

· **Adjusted R^2** (Slide L6 p.21):

$$R^2_{adj} = 1 - \frac{RSS_d/(n - d - 1)}{TSS/(n - 1)}$$

Penalizes for adding useless predictors. Choose model with largest Adj
R^2.

2. **Direct Estimation of Test Error using Cross-Validation**:
   · For each model size $k = 0, \ldots, p$, get the best k-variable model $\mathcal{M}_k$.
   · Compute its CV test error (e.g., 5-fold or 10-fold CV).
   · Select k that minimizes CV test error. This is often considered the
     most direct and reliable approach if computationally feasible.
   · "One-standard-error rule": Choose the simplest model whose CV error
     is within one standard error of the minimum CV error.

2.7.6 Exercise: ISLR 6.8

(Refers to 'ex6.8.R')

* (a) Generate data: $X \sim N(0, 1)$, $\epsilon \sim N(0, 1)$, $Y = \beta_0 + \beta_1 X + \cdots + \beta_p X^p + \epsilon$.
  For this exercise, data is simulated for $p = 100$ predictors, $n = 1000$ obs, where
  some β_j are zero. $Y = X\beta + \epsilon$.

* (b) Split into train (100 obs) and test (900 obs).

* (c) Best subset selection on training data, plot training MSE vs. subset size.
  Training MSE will decrease monotonically.

* (d) Plot test MSE vs. subset size. Will show a U-shape, identifying optimal
  number of predictors for test performance.

* (e) Find best model size based on test MSE.

* (f) Compare coefficients of best model to true βs.

* (g) Plot error $\sqrt{\sum(\hat{\beta}_j - \beta_j)^2}$ vs. subset size. This measures how well coeffi-
  cients are estimated. The minimum may not align with minimum test MSE.

* *Key R functions for ISLR 6.8*: `rnorm()`, `matrix()`, `sample()`, `leaps::regsubsets()`,
  `coef()`, `predict.regsubsets()` (custom function usually needed for 'regsub-
  sets' as it doesn't have a standard 'predict' method for new data directly).

```
# predict.regsubsets <- function(object, newdata, id, ...) {
#    form <- as.formula(object$call[[2]]) # Get formula from
     regsubsets object
#    mat <- model.matrix(form, newdata)   # Create model matrix
     for new data
#    coefi <- coef(object, id=id)         # Get coefficients for
     model of size 'id'
#    xvars <- names(coefi)
#    pred <- mat[, xvars] %% coefi        # Matrix
     multiplication
#    return(pred)
# }
```

Listing 30: Conceptual predict function for regsubsets (ISLR Ch6 Lab)

2.8 Lecture 8: Regularization (Shrinkage) Methods and PCR

2.8.1 Introduction to Shrinkage Methods

* Recall: Subset selection methods select a subset of predictors and fit OLS.
* **Shrinkage Methods**: Use all p predictors, but constrain or "shrink" the estimated coefficients $\hat{\beta}_j$ towards zero (or each other).
* **Goal**: Reduce variance at the cost of a small increase in bias, potentially leading to better prediction accuracy (lower test MSE).
* **OLS objective**: Minimize $RSS(\boldsymbol{\beta}) = \sum_{i=1}^{n}(y_i - \mathbf{x}_i'\boldsymbol{\beta})^2$.
* **Shrinkage methods add a penalty term** to the RSS:
    · Ridge Regression: Minimize $RSS(\boldsymbol{\beta}) + \lambda \sum_{j=1}^{p} \beta_j^2$.
    · LASSO: Minimize $RSS(\boldsymbol{\beta}) + \lambda \sum_{j=1}^{p} |\beta_j|$.
* The penalty term discourages large coefficient values.

2.8.2 Ridge Regression

* **Objective Function** (Slide L7 p.3):

$$
\min_{\boldsymbol{\beta}} \left\{ \sum_{i=1}^{n} \left(y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^{p} \beta_j^2 \right\}
$$

The term $\lambda \sum_{j=1}^{p} \beta_j^2$ is the L_2 penalty or shrinkage penalty.
* **Tuning Parameter** $\lambda \geq 0$ (Slide L7 p.3):
    · Controls the amount of shrinkage.
    · $\lambda = 0$: Penalty has no effect. Ridge estimates = OLS estimates.
    · $\lambda \to \infty$: Shrinkage penalty dominates. Coefficients $\hat{\beta}_j^{\text{ridge}} \to 0$ (for $j = 1, \ldots, p$). Model approaches intercept-only model.
    · β_0 (intercept) is typically not penalized.
* **Equivalent Formulation** (Constrained optimization, Slide L7 p.4):

$$
\min_{\boldsymbol{\beta}} \left\{ \sum_{i=1}^{n} \left(y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^{p} \beta_j^2 \leq s
$$

There is a one-to-one correspondence between λ and s.
* **Standardization of Predictors** (Slides L7 p.4-5):
    · OLS estimates are scale equivariant (if X_j is multiplied by c, $\hat{\beta}_j^{OLS}$ is divided by c).
    · Ridge regression estimates are *not* scale equivariant because the penalty $\lambda \sum \beta_j^2$ treats all β_j equally, regardless of the scale of X_j.
    · **Crucial**: Standardize predictors to have mean 0 and standard deviation 1 before fitting Ridge:

$$
\tilde{x}_{ij} = \frac{x_{ij} - \bar{x}_j}{s_j}
$$

(Slide L7 p.5 shows $\tilde{x}_{ij} = x_{ij}/\sqrt{\frac{1}{n}\sum(x_{ij} - \bar{x})^2}$, which is scaling to have RMS of 1, effectively standardizing if mean is 0. Standard practice is to scale to unit variance.)

* **Effect on Bias and Variance** (Slide L7 p.3):
  · Ridge regression reduces variance compared to OLS, especially when predictors are collinear.
  · It introduces some bias (coefficients are shrunk from OLS values).
  · Can lead to lower Test MSE if the reduction in variance outweighs the increase in squared bias.

* **R Implementation (glmnet package)** (Slides L7 p.9-12):
  · `glmnet()` function. Set `alpha=0` for Ridge.
  · Requires predictors `X` as a matrix and response `y` as a vector.
  · `cv.glmnet()` performs cross-validation to select optimal λ.

```r
# Simulated data from slides L7 p.6
# set.seed(1234567)
# n <- 100
# p_vars <- 10
# X_sim <- matrix(0, n, p_vars)
# for(j in 1:p_vars) X_sim[,j] <- rnorm(n, sd=j) #
    Predictors with increasing variance
# beta_sim <- 2:(p_vars+1)
# e_sim <- rnorm(n, sd=200)
# y_sim <- 1 + X_sim %% beta_sim + e_sim

# OLS fit (Slide L7 p.8)
# ols.fit <- lm(y_sim ~ X_sim)
# summary(ols.fit) # Shows some coeffs not significant, high
    RSE

library(glmnet)
# Ridge with small lambda (Slide L7 p.9)
# ridge.fit_small_lambda <- glmnet(X_sim, y_sim, alpha=0,
    lambda=0.01)
# cbind(coef(ols.fit), coef(ridge.fit_small_lambda)) #
    Coeffs very similar

# Ridge with large lambda (Slide L7 p.10)
# ridge.fit_large_lambda <- glmnet(X_sim, y_sim, alpha=0,
    lambda=1000000)
# cbind(coef(ols.fit), coef(ridge.fit_large_lambda)) # Ridge
     coeffs shrunk close to 0

# Choosing lambda by LOOCV (Slide L7 p.11)
# cv.ridge <- cv.glmnet(X_sim, y_sim, alpha=0, nfolds=n) #
    nfolds=n for LOOCV
# lambda_min_ridge <- cv.ridge$lambda.min # e.g., 28.27952
    from slide
# ridgemin.fit <- glmnet(X_sim, y_sim, alpha=0, lambda=
    lambda_min_ridge)
# cbind(coef(ols.fit), coef(ridgemin.fit)) # Shows shrunk
    coefficients

# Test predictions (Slide L7 p.12)
# # Assume xy_df created, split into train_df, test_df,
    Xtrain, ytrain, Xtest, ytest
# olstrain.fit <- lm(y ~ ., data=train_df) # Using original
    formula style with data.frame
# predols_test <- predict(olstrain.fit, newdata=test_df)
# ols_MSE_test <- mean((ytest - predols_test)^2) # e.g.,
    53559.83
```

```
35
36   # ridgetrain.fit <- glmnet(Xtrain, ytrain, alpha=0, lambda=
         lambda_min_ridge)
37   # predridge_test <- predict(ridgetrain.fit, newx=Xtest)
38   # ridge_MSE_test <- mean((ytest - predridge_test)^2) # e.g.,
         52444.16 (better than OLS)
```

Listing 31: Ridge Regression Example (Simulated Data, Slides L7 p.6-12)

2.8.3 The Lasso (Least Absolute Shrinkage and Selection Operator)

* **Objective Function** (Slide L7 p.13):

$$\min_{\boldsymbol{\beta}} \left\{ \sum_{i=1}^{n} \left(y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij} \right)^2 + \lambda \sum_{j=1}^{p} |\beta_j| \right\}$$

The term $\lambda \sum_{j=1}^{p} |\beta_j|$ is the L_1 penalty.

* **Key Property**: Unlike Ridge, the L_1 penalty can force some coefficient estimates to be *exactly zero* if λ is large enough.

* This means Lasso performs **variable selection**.

* **Tuning Parameter** $\lambda \geq 0$ (Slide L7 p.13):
  · $\lambda = 0$: Lasso estimates = OLS estimates.
  · $\lambda \to \infty$: All coefficients $\hat{\beta}_j^{\text{lasso}} \to 0$.
  · As λ increases, more coefficients are set to zero.

* **Equivalent Formulation** (Constrained optimization):

$$\min_{\boldsymbol{\beta}} \left\{ \sum_{i=1}^{n} \left(y_i - \beta_0 - \sum_{j=1}^{p} \beta_j x_{ij} \right)^2 \right\} \quad \text{subject to} \quad \sum_{j=1}^{p} |\beta_j| \leq s$$

* Like Ridge, predictors should be standardized before applying Lasso.

* **R Implementation (glmnet package)** (Slides L7 p.14-15):
  · Use `glmnet()` with `alpha=1`.
  · `cv.glmnet()` for choosing λ.

```
1    # Using X_sim, y_sim from Ridge example
2    # cv.lasso <- cv.glmnet(X_sim, y_sim, alpha=1, nfolds=n) #
         nfolds=n for LOOCV
3    # lambda_min_lasso <- cv.lasso$lambda.min # e.g., 13.12619
         from slide
4    # lassomin.fit <- glmnet(X_sim, y_sim, alpha=1, lambda=
         lambda_min_lasso)
5    # cbind(coef(ols.fit), coef(ridgemin.fit), coef(lassomin.fit
         ))
6    # # Lasso sets X1, X3, X4 to zero (coefficients are '.') in
         slide example
7
8    # Test predictions (Slide L7 p.15)
9    # # Using Xtrain, ytrain, Xtest, ytest from Ridge example
10   # lassotrain.fit <- glmnet(Xtrain, ytrain, alpha=1, lambda=
         lambda_min_lasso)
11   # predlasso_test <- predict(lassotrain.fit, newx=Xtest)
```

```
12  # lasso_MSE_test <- mean((ytest - predlasso_test)^2) # e.g.,
        51258.12 (best of OLS, Ridge, Lasso)
```

Listing 32: Lasso Example (Simulated Data, Slides L7 p.14-15)

2.8.4 Comparing Lasso and Ridge

* **Variable Selection**: Lasso can produce sparse models (some $\beta_j = 0$), Ridge cannot (all β_j are non-zero unless $\lambda = \infty$).
* **Performance**:
  · If many predictors have small/moderate effects, Ridge might perform better (keeps all variables).
  · If a smaller number of predictors have substantial effects and others are negligible, Lasso might perform better (by setting negligible ones to zero).
* **Geometric Interpretation** (ISLR Fig 6.7):
  · Ridge constraint $\sum \beta_j^2 \leq s$ is a circle (2D) or hypersphere. OLS solution contours (ellipses) are less likely to hit an axis exactly.
  · Lasso constraint $\sum |\beta_j| \leq s$ is a diamond (2D) or rhomboid. OLS solution contours are more likely to hit a corner, making a coefficient zero.
* **Elastic Net**: A compromise between Ridge and Lasso, penalty $\alpha \sum |\beta_j| + (1 - \alpha) \sum \beta_j^2$. (`glmnet()` can fit this by setting $0 < \alpha < 1$).

2.8.5 Selecting the Tuning Parameter λ

* Cross-validation (typically k-fold, e.g., 10-fold) is the standard method.
* For a grid of λ values:
  1. For each λ, perform k-fold CV.
  2. Calculate average test MSE (or other error metric) across folds for that λ.
  3. Choose λ that results in the lowest average CV test error (`lambda.min` from `cv.glmnet()`).
  4. Often, `lambda.1se` is also considered: largest λ such that error is within one standard error of the minimum. This gives a more parsimonious model with similar performance.

2.8.6 Dimension Reduction Methods

* Transform predictors $X_1, \ldots, X_p$ into a new set of $M < p$ predictors $Z_1, \ldots, Z_M$. Then fit OLS using Z_m.
* $Z_m = \sum_{j=1}^{p} \phi_{jm} X_j$ (linear combinations). (Slide L7 p.16)
* Two main methods: Principal Components Regression (PCR) and Partial Least Squares (PLS).

2.8.7 Principal Components Analysis (PCA) and Regression (PCR)

* **Principal Component Analysis (PCA)** (Slides L7 p.17-18):
  · A technique for dimension reduction that finds directions (principal components, PCs) in the data that capture the most variance.

- Z_1 (first PC): Linear combination of X_js that has maximal variance.
- Z_2 (second PC): Linear combination of X_js, uncorrelated with Z_1, that has maximal variance among all such combinations.
- And so on, up to p PCs.
- The ϕ_{jm} (loadings) define the components.
- Predictors should be standardized before PCA if on different scales.

* **Principal Components Regression (PCR)**:
  1. Perform PCA on the p predictors to get $Z_1, \ldots, Z_p$.
  2. Select the first $M < p$ principal components. M is a tuning parameter chosen by CV.
  3. Fit OLS regression of Y on $Z_1, \ldots, Z_M$.

* PCA is an unsupervised method (does not use Y to find components). PCR becomes supervised when Y is regressed on Z_m.

* PCR can reduce variance by using fewer components, especially if early PCs capture most of the signal related to Y.

* If most variation in predictors can be seen in a few PCs, we reduce overfitting. (Slide L7 p.18)

* **R Implementation (`pls` package)** (Slides L7 p.19-21):

```r
library(pls)
library(ISLR)
# Auto2 <- Auto[,!(names(Auto) %in% c("name","origin","year"))]
    # From slide L6 p.6
# Auto2$age <- 83 - Auto$year # From slide L6 p.6
# set.seed(1) # Assuming a seed for train/test split
# n_auto <- nrow(Auto2)
# train_idx_auto <- sample(1:n_auto, floor(n_auto/2))
# train_auto <- Auto2[train_idx_auto,]
# test_auto <- Auto2[-train_idx_auto,]

# Fit PCR on training data, choose M using CV (built-in)
# pcr.fit <- pcr(mpg ~ ., data=train_auto, scale=TRUE,
    validation="CV")
# summary(pcr.fit) # Shows % variance explained by X and mpg
    for different M
# validationplot(pcr.fit, val.type="MSEP") # Helps choose M

# Example output from summary(pcr.fit) on full Auto2 (Slide L7
    p.19):
# TRAINING: % variance explained
#       1 comps 2 comps 3 comps 4 comps 5 comps 6 comps
# X       99.76   99.96  100.00  100.00  100.00  100.00
# mpg     69.35   70.09   70.75   80.79   80.88   80.93
# First PC explains 99.76% of variance in X's, and 69.35% of
    variance in mpg.

# Training/test evaluation (Slide L7 p.21 - assuming M=4 chosen
    from CV)
# regall.fit <- lm(mpg ~ ., data=train_auto)
# pcr.fit_M4 <- pcr(mpg ~ ., ncomp=4, data=train_auto, scale=
    TRUE)
#
# predall_test <- predict(regall.fit, newdata=test_auto)
# MSEall_test <- mean((test_auto$mpg - predall_test)^2) # e.g.,
    11.47
```

```
29  #
30  # predpcr_test <- predict(pcr.fit_M4, newdata=test_auto)
31  # MSEpcr_test <- mean((test_auto$mpg - predpcr_test)^2) # e.g.,
        15.60
32  # In this specific split on slide L7, OLS performed better than
        PCR with M=4.
33  # Choice of M from CV on training data is crucial.
```

Listing 33: PCR Example (Auto Data, Slides L7 p.19, 21)

2.8.8 Partial Least Squares (PLS)

* Similar to PCR, but finds components Z_m in a supervised way.

* $Z_1 = \sum \phi_{j1} X_j$ where ϕ_{j1} are coefficients of simple linear regression of Y on each X_j. Z_1 is the linear combination most correlated with Y.

* Z_2 is found by first regressing Y on Z_1 and taking residuals. Then Z_2 is the linear combination of X_js (after orthogonalizing them w.r.t Z_1) that best explains these residuals.

* Can sometimes outperform PCR, especially if predictors are highly collinear or if the response is strongly related to directions with lower variance in X.

* Number of components M chosen by CV.

* R: `plsr()` in `pls` package.

2.8.9 Exercise: ISLR 6.11

(Refers to 'ex6.11.R')

* Goal: Compare Best Subset, Lasso, Ridge, PCR for predicting per capita crime rate (`crim`).

* Setup: Split `Boston` data into training/test. Prepare `X` matrix and `y` vector.

* (a) Best Subset Selection:
  · Use k-fold CV (e.g., 10-fold) to select optimal number of predictors.
  · Need a 'predict' method for 'regsubsets' (often custom written).
  · Fit best model of chosen size on full training set, evaluate on test set.
  · *From 'ex6.11.R'*: 10-fold CV suggests 9 predictors. Test RMSE ≈ 6.59.

* (a) Lasso:
  · Use `cv.glmnet()` on training data to find optimal λ.
  · Fit Lasso model with optimal λ on training data.
  · Predict on test data, calculate test RMSE.
  · *From 'ex6.11.R'*: Test RMSE (using λ_{1se}) ≈ 7.405.

* (a) Ridge Regression:
  · Use `cv.glmnet()` (`alpha=0`) on training data for λ.
  · Fit Ridge model with optimal λ on training data.
  · Predict on test data, calculate test RMSE.
  · *From 'ex6.11.R'*: Test RMSE (using λ_{1se}) ≈ 7.457.

* (a) PCR:

· Use `pcr()` with `validation="CV"` on training data to choose number of components M.
· Fit PCR model with chosen M on training data.
· Predict on test data, calculate test RMSE.
· *From 'ex6.11.R'*: CV suggests $M = 13$ components (full model). Test RMSE ≈ 6.546.

* (b) Conclusion: Compare test RMSEs. Best subset and PCR (with many components) performed best in this example.

* (c) Are models parsimonious?: Discuss which methods yield simpler models (Lasso, Best Subset if few variables chosen).

2.9 Lecture 9: Non-linear Models - Polynomials, Splines, Local Regression

2.9.1 Introduction to Non-linear Models

* Linear models assume a linear relationship between predictors and response, which is often too simplistic.

* This lecture explores methods to model non-linear relationships:
    · Polynomial Regression
    · Step Functions (Piecewise Constant Regression)
    · Regression Splines (Piecewise Polynomials)
    · Local Regression
    · Smoothing Splines (covered briefly, more detail in GAMs)

* Generalized Additive Models (GAMs) provide a framework to use these for multiple predictors (covered in Lecture 10).

2.9.2 Polynomial Regression

* Extends linear model by adding polynomial terms of predictors:

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \cdots + \beta_d x_i^d + \epsilon_i$$

* This is still a linear model in terms of coefficients β_j, so it can be fit using OLS.

* Choice of degree d is critical. Too low $d \implies$ underfitting. Too high $d \implies$ overfitting (wiggly fit).

* d often chosen by cross-validation or hypothesis tests (e.g., ANOVA to compare nested models).

* R: `lm(y ~ poly(x, degree=d))()` or `lm(y ~ x + I(x^2) + ...)()`. Using `poly()` creates orthogonal polynomials, which can be numerically more stable.

2.9.3 Step Functions (Piecewise Constant Regression)

* Divide the range of X into $K + 1$ disjoint regions using K cutpoints (knots) $c_1 < c_2 < \cdots < c_K$.

* Fit a constant (mean of Y) in each region:

$$y_i = \beta_0 + \beta_1 I(x_i < c_1) + \beta_2 I(c_1 \leq x_i < c_2) + \cdots + \beta_{K+1} I(x_i \geq c_K) + \epsilon_i$$

(Slide L8 p.3 shows a slightly different formulation with 5 coefficients for 4 knots, creating 5 regions).

* This is achieved by creating dummy variables for each interval.
* Choice of number and location of knots is important. Often chosen by CV or domain knowledge.
* R Example (Simulated data, Slides L8 p.3-7):

```
# Simulating data (Slides L8 p.3)
# set.seed(1); n=100; x=runif(n)
# c_knots=seq(0.10,0.90,length=4) # c1, c2, c3, c4
# b_coeffs=c(1,2,3,-1,1) # Beta for intercept and each interval
# y_step = b_coeffs[1] +
#          b_coeffs[2](x < c_knots[1]) +
#          b_coeffs[3]((x >= c_knots[1]) & x < c_knots[2]) +
#          b_coeffs[4]((x >= c_knots[2]) & x < c_knots[3]) + #
    Error in slide formula for b[4]
#          b_coeffs[5]((x >= c_knots[3]) & x < c_knots[4]) + #
    Error in slide formula for b[5]
#          0.1rnorm(n) # Should be 5 intervals for 5 beta
    coeffs after intercept if b[1] is overall mean
                    # Or one more interval for x >= c_knots[4]
                          if b[1] is first interval
# A more direct way in R is using cut()
# fit_step <- lm(y ~ cut(x, breaks=c(min(x)-.1, c_knots, max(x)
    +.1)))
# summary(fit_step)
# The coefficients will be differences from the intercept (mean
     of first interval).

# Example with Wage data (Slides L8 p.8)
# library(ISLR)
# attach(Wage)
# table(cut(age, 4)) # Divides age into 4 equally spaced
    intervals
# fit_wage_step <- lm(wage ~ cut(age, 4), data=Wage)
# coef(summary(fit_wage_step))
# # Shows average wage for first age group, and differences for
     others.
# detach(Wage)
```

Listing 34: Piecewise Constant Regression (Step Function) in R

* Drawback: Fit is not continuous, can miss trends at boundaries.

2.9.4 Basis Functions

* General idea: Transform X using a set of known basis functions $b_1(X), b_2(X), \ldots, b_K(X)$.
* Then fit linear model: $y_i = \beta_0 + \beta_1 b_1(x_i) + \cdots + \beta_K b_K(x_i) + \epsilon_i$.
* Examples:
  · Polynomial regression: $b_j(x) = x^j$.
  · Step functions: $b_j(x) = I(c_j \leq x < c_{j+1})$.
  · Periodic functions: $b_j(x) = \sin(\alpha_j x)$ or $\cos(\alpha_j x)$.

2.9.5 Regression Splines (Piecewise Polynomials)

* Fit separate low-degree polynomials in different regions defined by knots $\xi_1, \ldots, \xi_K$.

* **Constraints for Smoothness** (Slide L8 p.11): To avoid discontinuities, constraints are imposed at the knots.

  · Function must be continuous at knots.
  · First derivative continuous (smooth curve).
  · Second derivative continuous (smoother curve).

* A degree-d spline with K knots has $d+1+K$ degrees of freedom (parameters).

* **Cubic Spline** (degree 3) is common: Continuous, with continuous 1st and 2nd derivatives. Smooth and flexible.

* **Cubic Spline Basis Representation** (Slide L8 p.12): A cubic spline with K knots $\xi_1, \ldots, \xi_K$ can be modeled using a basis:

$$y_i = \beta_0 + \beta_1 x_i + \beta_2 x_i^2 + \beta_3 x_i^3 + \sum_{k=1}^{K} \beta_{k+3} h(x_i, \xi_k) + \epsilon_i$$

where $h(x, \xi) = (x - \xi)_+^3 = \begin{cases} (x - \xi)^3 & \text{if } x > \xi \\ 0 & \text{otherwise} \end{cases}$. This model has $K + 4$ parameters (degrees of freedom $= K + 4$).

* **Choosing Knots and Number of Knots (or Degrees of Freedom)**:

  · More knots $\implies$ more flexible, can overfit.
  · Knots often placed at uniform quantiles of X or based on domain knowledge.
  · Number of knots (or equivalently, degrees of freedom 'df') chosen by cross-validation.

* **Natural Splines** (ISLR Ch 7.4.3): Cubic splines constrained to be linear beyond boundary knots. Reduces variance at boundaries. Uses K df for K knots.

* **R Implementation (`splines` package)** (Slide L8 p.13, using `bs()`):

```
1  # library(ISLR)
2  # library(splines) # For bs() function
3  # attach(Wage)
4  # # Fit a cubic spline with knots at ages 25, 40, 60
5  # fit.spline <- lm(wage ~ bs(age, knots=c(25,40,60)), data=Wage
       )
6  # summary(fit.spline) # Shows coefficients for basis functions
7  #
8  # # Plotting the fit
9  # agelims <- range(age)
10 # age.grid <- seq(from=agelims[1], to=agelims[2])
11 # pred.spline <- predict(fit.spline, newdata=list(age=age.grid)
       )
12 # plot(age, wage, col="lightgrey")
13 # lines(age.grid, pred.spline, col="black", lwd=2)
14 # # Add vertical lines for knots
15 # abline(v=c(25,40,60), lty=2, col="red")
16 # detach(Wage)
17 # # Alternatively, specify degrees of freedom (df) instead of
       knots
```

```
18  # # fit.spline_df <- lm(wage ~ bs(age, df=6), data=Wage) # df=6
        for 3 knots cubic spline (df=K+3+1-1=K+3 usually, or K+
        degree)
19  # # bs() uses df = K + degree (for cubic, df=K+3). Intercept is
        separate.
20  # # So, 3 knots with bs() is df=3+3=6.
```
Listing 35: Regression Spline with Wage Data (Slide L8 p.13)

2.9.6 Smoothing Splines

* Goal: Find a function $g(x)$ that fits the data well ($RSS = \sum(y_i - g(x_i))^2$ is small) but is also smooth.

* **Objective Function** (Slide L9 p.3):

$$\min_g \left\{ \sum_{i=1}^{n}(y_i - g(x_i))^2 + \lambda \int (g''(t))^2 dt \right\}$$

  · First term: Measures closeness to data (RSS).
  · Second term: Roughness penalty. $\int(g''(t))^2 dt$ is large if $g(x)$ is "wiggly".
  · $\lambda \geq 0$ is a tuning parameter:
  · $\lambda = 0$: $g(x)$ interpolates data (can be very rough, zero RSS if all x_i unique). (Slide L9 p.2 figure)
  · $\lambda \to \infty$: $g''(t)$ must be zero, meaning $g(x)$ is a linear function (OLS line).

* **Solution**: The function $g(x)$ that minimizes this is a *natural cubic spline* with knots at every unique value of x_i. (Slide L9 p.3)

* Seems complex (many knots), but λ controls effective degrees of freedom (df_λ).

* λ is chosen by (generalized) cross-validation, often LOOCV.

* R: `smooth.spline(x, y, cv=TRUE)()` (Slide L9 p.5). `cv=TRUE` chooses λ by LOOCV.

```
1   # library(ISLR)
2   # attach(Wage)
3   # fit.smooth <- smooth.spline(age, wage, cv=TRUE) # Chooses
        lambda by LOOCV
4   # print(fit.smooth)
5   # # fit.smooth$df gives effective degrees of freedom, e.g.,
        6.79
6   #
7   # plot(age, wage, col="lightgrey")
8   # lines(fit.smooth, col="red", lwd=2) # Red line is CV-chosen
        smoothing spline
9   # # Can also specify df to get a specific lambda
10  # lines(smooth.spline(age, wage, df=16), col="blue", lwd=2) #
        More wiggly (higher df)
11  # detach(Wage)
```
Listing 36: Smoothing Spline with Wage Data (Slide L9 p.5)

2.9.7 Local Regression (LOESS/LOWESS)

* Fits separate simple models (e.g., linear or quadratic polynomials) in local neighborhoods of x_0.

53

* **Algorithm at a target point** x_0:

  1. Gather $s \cdot n$ training points closest to x_0 (the neighborhood). s is the span (e.g., 0.2 to 0.5).

  2. Assign weights to points in this neighborhood, with points closer to x_0 getting higher weights (e.g., tricube weight function). Points outside neighborhood get weight 0.

  3. Fit a weighted least squares regression of y on x using only points in the neighborhood with these weights.

  4. The fitted value $\hat{y}_0$ at x_0 is the prediction.

* This is repeated for all x_0 where a prediction is desired.

* Span s is the main tuning parameter, controls smoothness. Small $s \implies$ local, wiggly fit. Large $s \implies$ global, smoother fit.

* R: `loess(y   x, span=..., data=...)()`.

2.9.8   Kernel Estimators (Nadaraya-Watson)

* Model $y = f(x) + \epsilon$. (Slides L9 p.6)

* **Nadaraya-Watson Estimator**: A type of local averaging.

$$\hat{f}(x_0) = \sum_{i=1}^{n} w_i(x_0) y_i$$

where weights $w_i(x_0) = \dfrac{K\left(\frac{x_i - x_0}{h}\right)}{\sum_{j=1}^{n} K\left(\frac{x_j - x_0}{h}\right)}$ sum to 1.

* $K(\cdot)$ is a **kernel function**, a non-negative function that integrates to one and is symmetric about zero (e.g., Gaussian kernel, Epanechnikov kernel, boxcar kernel).

  · Gaussian kernel (using standard normal density ϕ): $K(u) = \phi(u)$. (Slide L9 p.6 shows weights with ϕ).

  · The plot on Slide L9 p.7 shows how $K((x - x_0)/h)$ gives higher weights to x_i close to x_0.

* h is the **bandwidth**, a tuning parameter controlling smoothness. Small $h \implies$ local, wiggly. Large $h \implies$ global, smooth. Chosen by CV.

* R Example (Slides L8 p.19, implementing N-W with Gaussian kernel):

```
# K_gauss <- function(x0_val, x_vec, h_val=0.2) {
#    u <- (x_vec - x0_val) / h_val
#    weights_numerator <- dnorm(u) # Using standard normal
      density
#    return(weights_numerator / sum(weights_numerator))
# }
#
# nw_estimator <- function(x0_val, x_train, y_train, h_val=0.2)
      {
#    weights <- K_gauss(x0_val, x_train, h_val)
```

```
 9 #    return(sum(weights   y_train))
10 # }
11 #
12 # # Example data from slides L8 p.16
13 # # set.seed(1); x_sim_nw <- rnorm(100, sd=20); y_sim_nw <- x_
      sim_nw2 + rnorm(100, sd=50)
14 # # plot(x_sim_nw, y_sim_nw)
15 # # ypred_knn_nw <- sapply(as.matrix(x_sim_nw), knn, x=x_sim_nw
      , y=y_sim_nw) # Using knn from slide L8 p.17
16 # # lines(sort(x_sim_nw), ypred_knn_nw[order(x_sim_nw)], col="
      blue")
17 # # ypred_nw_h5 <- sapply(as.matrix(x_sim_nw), nw_estimator,
18 # #                       x_train=x_sim_nw, y_train=y_sim_nw, h
      _val=5)
19 # # lines(sort(x_sim_nw), ypred_nw_h5[order(x_sim_nw)], col="
      red")
```

Listing 37: Nadaraya-Watson with Gaussian Kernel (Slides L8 p.19)

* **Local Linear Regression** (Slides L8 p.28): Can be seen as an extension of Nadaraya-Watson. Instead of local constant (average), fit a local linear model using weighted least squares, with weights from a kernel.

$$\min_{\beta_0,\beta_1} \sum_{i=1}^{n} K\left(\frac{x_i - x_0}{h}\right)(y_i - \beta_0 - \beta_1(x_i - x_0))^2$$

Then $\hat{f}(x_0) = \hat{\beta}_0$. (Slide formula has $\beta_1 x_i$, often $(x_i - x_0)$ is used for stability). LOESS is a form of local polynomial regression.

2.9.9   Exercises: ISLR 7.6, 7.9, (7.10, 7.11 for GAMs)

* **ISLR 7.6 (`Wage` data)**: (Corresponds to 'ex7.6.R')

  · (a) Polynomial regression for `wage    age`: Use CV to choose degree. ANOVA for comparison. Plot results. *Result from 'ex7.6.R'*: CV suggests degree 3 or 4 (varies with seed/method). ANOVA often suggests degree 3 or 4 as significant improvements over lower degrees.

  · (b) Step function for `wage    age`: Use CV to choose number of cuts. Plot. *Result from 'ex7.6.R'*: CV often suggests around 7-9 cuts.

* **ISLR 7.9 (`Boston` data)**: (Corresponds to 'ex7.9.R')

  · (a) Polynomial regression for `nox    dis`: Plot polynomial fits of various degrees.

  · (b) Report training RSS for different degrees. (Decreases monotonically).

  · (c) Use CV to select best degree. Plot CV error. (Often degree 3 or 4 is chosen).

  · (d) Regression spline with `bs()` for `nox    dis`: Specify knots (e.g., at quartiles of 'dis'). Report fit. Plot.

  · (e) Regression spline with varying df: Fit for df from 3 to 16. Report training RSS. (Decreases).

  · (f) Use CV to select best df for spline. Plot CV error.

2.10 Lecture 10: Generalized Additive Models (GAMs)

2.10.1 Introduction to GAMs

* Extends linear models to allow non-linear functions for *each* predictor, while maintaining an additive structure.

* Model form:

$$y_i = \beta_0 + f_1(x_{i1}) + f_2(x_{i2}) + \cdots + f_p(x_{ip}) + \epsilon_i$$

or, for expected value: $E(Y|X_1, \ldots, X_p) = \beta_0 + \sum_{j=1}^{p} f_j(X_j)$.

* Each $f_j(x_{ij})$ can be a smooth non-linear function (e.g., spline, local polynomial) or a linear term.

* **Additive**: Allows examining the effect of each X_j on Y individually, holding other predictors constant, by plotting $\hat{f}_j(x_j)$.

* Avoids the "curse of dimensionality" faced by fitting a general non-linear function $f(x_1, \ldots, x_p)$ directly in high dimensions.

2.10.2 Fitting GAMs: Backfitting Algorithm

* An iterative procedure to fit GAMs.

* **Example with two predictors** ($y = f_1(x_1) + f_2(x_2) + \epsilon$) (Slide L9 p.9-10):

  1. Initialize: e.g., $\hat{\beta}_0 = \bar{y}$, $\hat{f}_2(x_{i2}) = 0$ for all i.
  2. Iterate until convergence:
     · Estimate $\hat{f}_1$ by fitting a model to $y_i - \hat{\beta}_0 - \hat{f}_2(x_{i2})$ using x_{i1} (e.g., smoothing spline on x_1).
     · Estimate $\hat{f}_2$ by fitting a model to $y_i - \hat{\beta}_0 - \hat{f}_1(x_{i1})$ using x_{i2} (e.g., smoothing spline on x_2).
     · (Often, residuals are centered).
  3. This means iteratively fitting each f_j on "partial residuals" $y_i - \hat{\beta}_0 - \sum_{k \neq j} \hat{f}_k(x_{ik})$.

* R: The `gam` package (or `mgcv`) handles this.

2.10.3 GAMs for Regression

* R Example (`ISLR Wage` data, using `gam` package):

```
# library(ISLR)
# library(gam) # Older gam package. mgcv is more modern.
# attach(Wage)

# GAM with smoothing splines for year and age (df specified)
# gam.fit1 <- gam(wage ~ s(year, df=3) + s(age, df=3), data=
    Wage) # Slide uses df=3 for both
# plot.Gam(gam.fit1, terms="s(age, 3)", col="red", se=TRUE) #
    Slide L9 p.11, example with one term
# par(mfrow=c(1,3)) # To plot multiple terms (Slide L9 p.12)
```

```
10  # GAM including a linear term for education and smoothing
        splines for year & age
11  gam.fit2 <- gam(wage ~ s(year, df=4) + s(age, df=5) + education
        , data=Wage)
12  # Note: df might be different in slide examples, often chosen
        by CV or set.
13  # plot.Gam(gam.fit2, se=TRUE, col="red") # Plots all terms
14  # par(mfrow=c(1,3)); plot(gam.fit2, se=TRUE, col="red") #
        Example from ISLR book for Fig 7.12
15  # (Slide L9 p.12 shows plots for s(year), s(age), education. s
        () implies smoothing spline)
16  # The plots show the estimated non-linear effect of year and
        age, and linear effect of education.
17  # detach(Wage)
```

Listing 38: GAM for Regression with Wage Data (Slide L9 p.11-12)

* For factor predictors like 'education', `gam()` fits a separate constant for each level (similar to linear regression with dummy variables).

2.10.4 GAMs for Classification

* Extends logistic regression by allowing non-linear functions for predictors.

* Model for binary response $Y \in \{0, 1\}$ (Slide L9 p.13):

$$\log\left(\frac{P(Y = 1|X)}{1 - P(Y = 1|X)}\right) = \beta_0 + f_1(x_1) + \cdots + f_p(x_p)$$

* Each f_j can be a smooth function.

* Fitted using backfitting with weighted least squares for each step (as logistic regression is fit with IRLS).

* R Example (`Wage` data, predict if `wage > 250`, Slide L9 p.14):
```
1  # library(ISLR)
2  # library(gam)
3  # attach(Wage)
4  # gam.log.fit <- gam(I(wage > 250) ~ s(year, df=4) + s(age, df
        =5) + education,
5  #                          data=Wage, family=binomial)
6  # # par(mfrow=c(1,3))
7  # # plot.Gam(gam.log.fit, se=TRUE, col="red")
8  # # Shows effect of each predictor on the log-odds of wage >
        250.
9  # detach(Wage)
```

Listing 39: GAM for Classification (Slide L9 p.14)

2.10.5 Pros and Cons of GAMs

* **Pros**:

    · Automatically model non-linear relationships without manual specification of transformations.

57

- Additive nature allows examining effect of each predictor individually.
- More flexible than linear models, but more interpretable than fully non-parametric methods like KNN in high dimensions.
- Can improve predictive accuracy if true relationships are non-linear.

  * **Cons**:

    - Primarily additive, so important interactions can be missed unless explicitly included (e.g., 's(x1,x2)' or 'te(x1,x2)' in `mgcv`).
    - With many predictors, interpretation of many individual non-linear plots can still be complex.

2.10.6 Exercises: ISLR 7.10, 7.11

* **ISLR 7.10 (`College` data)**: (Corresponds to 'ex7.10.R')

  - (a) Forward stepwise selection for `Outstate` .. Identify best model by BIC/Cp/AdjR2. (E.g., 6 variables: Private, Room.Board, Terminal, perc.alumni, Expend, Grad.Rate).
  - (b) Fit GAM with selected variables, using smoothing splines (`s()`) for continuous predictors. Plot terms.
  - (c) Evaluate GAM on test set. Compare MSE with linear model using same predictors.
  - (d) Check for non-linear evidence using Anova for GAM terms ('summary(gam.fit)' in `gam` provides this for non-parametric effects).
  - *Result from 'ex7.10.R'*: GAM Test MSE can be slightly better than OLS. Anova often shows evidence of non-linearity for some predictors like 'Expend'.

* **ISLR 7.11 (Backfitting Algorithm)**: (Corresponds to 'ex7.11.R')

  - (a) Simulate data: $Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \epsilon$.
  - (b, c, d, e) Implement backfitting algorithm manually: Initialize $\hat{\beta}_1$. Iterate: 1. $a_i = y_i - \hat{\beta}_1 x_{i1}$. Regress a_i on X_2 to get $\hat{\beta}_{2,\text{iter}}$. 2. $a_i = y_i - \hat{\beta}_{2,\text{iter}} x_{i2}$. Regress a_i on X_1 to get $\hat{\beta}_{1,\text{iter}+1}$. Repeat until convergence. Plot $\hat{\beta}_0, \hat{\beta}_1, \hat{\beta}_2$ vs. iteration.
  - (f) Compare with OLS coefficients from `lm(Y  X1 + X2)()`. They should match.
  - (g) How many iterations? Typically very few for linear case.
  - *'ex7.11.R' shows code for this manual backfitting.*

2.11 Lecture 11: Tree-Based Methods

2.11.1 Introduction to Decision Trees

* **Core Idea**: Segment the predictor space into a number of simple, non-overlapping regions. For every observation that falls into a region R_m, make the same prediction.

∗ **Regression Trees**: Prediction is the mean of response values for training observations in R_m.

∗ **Classification Trees**: Prediction is the most common class (mode) for training observations in R_m.

∗ Example (Slides L10 p.2-3, ISLR `Hitters` data, predict `logSalary` from `Years` and `Hits`):

· First split: `Years < 4.5`.

· If `Years < 4.5`, predict mean logSalary for this group (e.g., 5.11).

· If `Years >= 4.5`, then split further by `Hits < 117.5`.

· If `Hits < 117.5`, predict mean (e.g., 6.00).

· If `Hits >= 117.5`, predict mean (e.g., 6.74).

· These splits define regions R_1, R_2, R_3.

2.11.2 Building Regression Trees

∗ **Stratifying Predictor Space**:

1. Divide predictor space $(X_1, \ldots, X_p)$ into J non-overlapping rectangular regions $R_1, \ldots, R_J$.

2. For any observation in region R_j, predict $\hat{y}_{R_j} = \text{mean}(y_i | x_i \in R_j)$.

3. Goal: Find regions that minimize RSS:

$$RSS = \sum_{j=1}^{J} \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2$$

∗ **Recursive Binary Splitting** (Greedy Algorithm, Slides L10 p.7-8, 14):

· Computationally infeasible to consider all possible partitions.

· Start with all data in one region.

· Iteratively split existing regions. At each step, choose the predictor X_j and cutpoint s that lead to the greatest possible reduction in RSS when splitting a region into $\{X | X_j < s\}$ and $\{X | X_j \geq s\}$.

· This is a top-down, greedy approach (best split at current step, may not lead to global optimum).

· The two new regions are then split further using the same principle.

· Continue until a stopping criterion is met (e.g., minimum number of observations per region/node, say 5 or 10 (Slide L10 p.14 `mincut=10`)).

∗ R Example (Manual splitting logic for `Hitters` data, Slides L10 p.9-12):

· Demonstrates finding the best split for `Hits` by minimizing RSS over possible cutpoints s.

· Then for `Years`.

· Compares RSS reduction from best `Hits` split vs. best `Years` split. Chooses `Years < 4.5` as the first split because it gives lower total RSS.

∗ R Implementation with `tree` package (Slides L10 p.15-16):

```
1  # library(ISLR)
2  # data(Hitters)
3  # Hitters <- Hitters[complete.cases(Hitters$Salary),] # Remove
       NAs for Salary
4  # Hitters$logSalary <- log(Hitters$Salary)
5
6  # library(tree)
7  # Fit tree to predict logSalary using Years and Hits
8  # tree.fit1 <- tree(logSalary ~ Years + Hits, data=Hitters,
9  #                   control = tree.control(nobs=nrow(Hitters),
       mincut=10)) # mincut is like minleaf
10 # summary(tree.fit1)
11 # plot(tree.fit1)
12 # text(tree.fit1, pretty=0, digits=3) # digits for precision of
        node labels
13
14 # Fit tree with all variables (Slide L10 p.16, more complex
       tree)
15 # tree.fit.allvars <- tree(logSalary ~ . - Salary, data=Hitters
       ,
16 #                          control = tree.control(nobs=nrow(
       Hitters), mincut=10))
17 # plot(tree.fit.allvars)
18 # text(tree.fit.allvars, pretty=0)
```

Listing 40: Fitting a Regression Tree with `tree` (Slides L10 p.15)

2.11.3   Tree Pruning

∗ **Problem**: Large trees grown by recursive binary splitting often overfit the training data, leading to poor test performance.

∗ **Strategy**:

1. Grow a large initial tree T_0 (e.g., by having a small 'mincut' or 'minsplit').

2. Prune this tree back to get a sequence of smaller subtrees.

∗ **Cost Complexity Pruning (Weakest Link Pruning)** (Slide L10 p.17):

· For each possible subtree $T \subseteq T_0$, consider a penalized RSS:

$$\sum_{m=1}^{|T|} \sum_{i:x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha|T|$$

where $|T|$ is the number of terminal nodes in subtree T.

· $\alpha \geq 0$ is a non-negative tuning parameter.

· $\alpha = 0$: T_0 is chosen (no penalty for complexity).

60

· As α increases, the penalty for more terminal nodes increases, favoring smaller trees.

· Goal: For a given α, find the subtree T_α that minimizes this criterion.

* **Selecting Optimal α (and thus best subtree)**:

    · Use k-fold cross-validation.

    · For each α in a range, get T_α.

    · Compute CV error for each T_α.

    · Choose α (and corresponding T_α) that minimizes CV error.

* R Implementation (`tree` package, Slides L10 p.18-19):

```
# Assuming tree.fit1 from before (logSalary ~ Years + Hits)
# cv.results <- cv.tree(tree.fit1, FUN=prune.tree) # FUN=prune.
    tree is for regression
# # For classification, FUN=prune.misclass can be used.
# names(cv.results) # "size", "dev", "k", "method"
# # 'size' is number of terminal nodes.
# # 'dev' is cross-validated deviance (RSS for regression, or
    deviance for classification).
# # 'k' is effectively alpha (cost-complexity parameter).
#
# plot(cv.results$size, cv.results$dev, type="b",
#      xlab="Number of Terminal Nodes (Tree Size)", ylab="CV
    Deviance (RSS)")
# # Identify the size that minimizes CV deviance.
# best.size <- cv.results$size[which.min(cv.results$dev)] # e.g
    ., 4 from slide
#
# # Prune the original tree to this best size
# pruned.tree <- prune.tree(tree.fit1, best=best.size)
# plot(pruned.tree)
# text(pruned.tree, pretty=0)
```

Listing 41: Pruning a Regression Tree (Slides L10 p.18-19)

2.11.4 Classification Trees

* Similar to regression trees, but Y is qualitative.

* **Prediction in a region R_m**: Predict the most common class among training observations in R_m.

* **Splitting Criteria** (Instead of RSS, Slide L10 p.21):

    · Let $\hat{p}_{mk}$ be the proportion of training observations in region R_m that belong to class k.

    · **Classification Error Rate**: $E_m = 1 - \max_k(\hat{p}_{mk})$. Proportion of misclassified obs in region R_m if we predict majority class.

    · Not sensitive enough for tree-growing (can lead to same error for different splits).

    · Often used for pruning.

    · **Gini Index**: $G_m = \sum_{k=1}^{K} \hat{p}_{mk}(1 - \hat{p}_{mk})$.

- · Measure of total variance across K classes, or node purity.

- · Small Gini $\implies$ node contains mostly observations from one class.

- · **Cross-Entropy (Deviance)**: $D_m = -\sum_{k=1}^{K} \hat{p}_{mk} \log(\hat{p}_{mk})$.

- · Similar to Gini index. Small value $\implies$ node is pure.

- · Gini or Cross-Entropy are typically used for growing the tree because they are more sensitive to changes in node probabilities than classification error rate.

* Recursive binary splitting proceeds by choosing splits that achieve the largest reduction in Gini index or cross-entropy.

* Pruning is done similarly to regression trees, using CV with error rate, Gini, or deviance.

* R Example (`ISLR` `Carseats` data, Slides L10 p.22-25):

```
# library(ISLR)
# data(Carseats)
# Carseats$High <- as.factor(ifelse(Carseats$Sales <= 8, "No",
    "Yes"))
#
# library(tree)
# tree.carseats <- tree(High ~ . - Sales, data=Carseats) #
    Exclude original Sales
# summary(tree.carseats)
# plot(tree.carseats)
# text(tree.carseats, pretty=0) # Shows classes and proportions
    at nodes
#
# # Pruning for classification tree
# cv.carseats <- cv.tree(tree.carseats, FUN=prune.misclass) #
    Use misclass error for pruning
# plot(cv.carseats$size, cv.carseats$k, type="b", # Note: Slide
    plots size vs deviance (k=-alpha)
#       xlab="Tree Size", ylab="Alpha (Cost-Complexity)")
# # Or plot deviance: plot(cv.carseats$size, cv.carseats$dev,
    type="b")
#
# # From slide L10 p.24, deviance plot suggests size 4 might be
    optimal
# best.size.class <- 4 # Example from slide
# pruned.carseats <- prune.misclass(tree.carseats, best=best.
    size.class)
# plot(pruned.carseats)
# text(pruned.carseats, pretty=0)
```

Listing 42: Classification Tree for Carseats Data (Slides L10 p.22-25)

2.11.5 Pros and Cons of Decision Trees

* **Pros (+)**:

- · Simple to understand and interpret.

- · Easy to visualize (graphical representation).

- · Can handle categorical predictors without dummy variables (natively in some packages like `tree`).

· Can capture non-linear relationships and interactions.

· Considered to mimic human decision-making.

* **Cons (-):**

· Predictive accuracy is often not as good as other methods.

· High variance: Small changes in training data can lead to very different tree structures.

· Greedy approach of recursive binary splitting is not guaranteed to find the globally optimal tree.

2.11.6   Exercise: ISLR 8.8a-c

(Corresponds to 'ex8.8.R' file, first part)

* (a) Split data into training and test sets.
* (b) Fit a regression tree to training data. Plot. Interpret. Compute test MSE.

```
# library(ISLR); library(tree)
# data(Carseats)
# set.seed(123) # Or seed from problem
# train_idx <- sample(1:nrow(Carseats), nrow(Carseats)/2)
# train_data <- Carseats[train_idx,]
# test_data <- Carseats[-train_idx,]
#
# tree.sales <- tree(Sales ~ ., data=train_data)
# summary(tree.sales) # Variables used, number of nodes,
    deviance
# plot(tree.sales); text(tree.sales, pretty=0)
# # Interpretation: e.g., ShelveLoc=Good and Price < 109.5 ->
    high sales
#
# pred.sales.test <- predict(tree.sales, newdata=test_data)
# test_mse_tree <- mean((test_data$Sales - pred.sales.test)^2)
# # test_mse_tree around 4-5 usually
```

Listing 43: ISLR 8.8b - Basic Regression Tree

* (c) Use cross-validation to determine optimal tree complexity (pruning). Does pruning improve test MSE?

```
# cv.sales <- cv.tree(tree.sales, FUN=prune.tree)
# plot(cv.sales$size, cv.sales$dev, type="b", xlab="Tree Size",
    ylab="CV Deviance")
# # Find optimal size (e.g., cv.sales$size[which.min(cv.sales$
    dev)])
# best_size_cv <- cv.sales$size[which.min(cv.sales$dev)] # e.g
    ., 15 in ex8.8.R
#
# pruned.sales <- prune.tree(tree.sales, best=best_size_cv)
# plot(pruned.sales); text(pruned.sales, pretty=0)
#
# pred.pruned.test <- predict(pruned.sales, newdata=test_data)
# test_mse_pruned <- mean((test_data$Sales - pred.pruned.test)
    ^2)
# # Compare test_mse_pruned with test_mse_tree.
# # In ex8.8.R, pruning didn't help or slightly worsened MSE
    for that particular seed.
```

Listing 44: ISLR 8.8c - Pruning Regression Tree

2.12 Lecture 12: Bagging, Random Forests, Boosting

2.12.1 Bagging (Bootstrap Aggregation)

* **Motivation** (Slide L11 p.3): Decision trees suffer from high variance. Averaging a set of observations reduces variance: $\text{Var}(\bar{X}) = \sigma^2/n < \sigma^2 = \text{Var}(X)$.

* **Idea** (Slide L11 p.4): Reduce prediction error variance by fitting many trees on different training sets and averaging their predictions. Since we only have one training set, we use bootstrap.

* **Bagging Algorithm for Regression Trees** (Slides L11 p.5, 12-13):

  1. Generate B bootstrap samples from the original training data (sample with replacement).

  2. For each bootstrap sample $b = 1, \ldots, B$:

     · Grow a regression tree $\hat{f}^b(x)$ on this bootstrap sample. Trees are typically grown deep (unpruned).

  3. To predict for a new point x_0: Average the predictions from all B trees:

  $$\hat{f}_{\text{bag}}(x_0) = \frac{1}{B} \sum_{b=1}^{B} \hat{f}^b(x_0)$$

* For classification, the prediction is the majority vote among the B trees. (Slide L11 p.13)

* Bagging for KNN (Slides L11 p.5, 8-10): Illustrates the concept. KNN is relatively stable, so bagging offers less improvement compared to unstable methods like trees. Test MSE improved slightly from 0.0607 to 0.0564 in the example.

* **Out-of-Bag (OOB) Error Estimation** (ISLR Ch 8.2.2, Slide L11 p.20):

  · Computationally efficient way to estimate test error without explicit CV.

  · Each bootstrap sample uses 2/3 of original observations. The remaining 1/3 are "out-of-bag" (OOB) for that tree.

  · For each observation i:

  · Predict y_i using only the trees for which observation i was OOB.

  · Average these predictions (or take majority vote).

  · Calculate overall OOB MSE or error rate. Provides a valid estimate of test error.

* R Example (Bagging for `Carseats` Sales, Slides L11 p.14-15):

```r
# library(ISLR); library(randomForest)
# data(Carseats)
# set.seed(1) # Assuming a seed from lecture example for
    Carseats train/test
# train_idx_cs <- sample(1:nrow(Carseats), nrow(Carseats)/2)
# train_cs <- Carseats[train_idx_cs,]
# test_cs <- Carseats[-train_idx_cs,]

# For bagging, set mtry = number of predictors (e.g., 10 if
    Sales is excluded)
# ncol(train_cs)-1 if Sales is the first column and no other
    exclusions.
```

```
10  # For Carseats , Sales is col 1. Other 10 vars are predictors.
11  # bagging.carseats <- randomForest(Sales ~ ., data=train_cs,
      mtry=10, importance=TRUE)
12  # # Default ntree=500
13  # pred.bag <- predict(bagging.carseats , newdata=test_cs)
14  # mse.bag <- mean((test_cs$Sales - pred.bag)^2)
15  # # Slide L11 p.15 shows mse.bag = 3.113909
16  # # Compare to single tree mse1 = 4.881557 (Slide L11 p.14) -
      Bagging improves.
```

Listing 45: Bagging with randomForest Package (Slide L11 p.15)

2.12.2 Random Forests

∗ **Improvement over Bagging**: Bagged trees can be highly correlated if there are a few very strong predictors that are always selected at top splits. Averaging correlated quantities doesn't reduce variance as much.

∗ **Random Forest Algorithm** (Slide L11 p.17):

1. Generate B bootstrap samples.

2. For each bootstrap sample, grow a tree:

   · At *each split* in the tree, randomly select a subset of m predictors out of the total p predictors.

   · Choose the best split among these m predictors only.

3. Average predictions (regression) or take majority vote (classification).

∗ **Choice of m**:

   · Typically $m \approx \sqrt{p}$ for classification.

   · Typically $m \approx p/3$ for regression.

   · If $m = p$, then Random Forest = Bagging.

∗ By restricting choice of predictors at each split, Random Forests decorrelate the trees, leading to greater variance reduction when averaged.

∗ R Example (`Carseats` Sales, Slides L11 p.18):

```
1  # Assuming train_cs , test_cs from Bagging example
2  # For Carseats , p=10 predictors. mtry=p/3 approx 3.
3  # rf.carseats <- randomForest(Sales ~ ., data=train_cs, mtry=3,
      importance=TRUE)
4  # pred.rf <- predict(rf.carseats , newdata=test_cs)
5  # mse.rf <- mean((test_cs$Sales - pred.rf)^2)
6  # # Slide L11 p.18 shows mserf = 3.568976. This is worse than
      bagging in this specific example.
7  # # Often RF is better or similar. Results can vary with seed
      and mtry choice.
```

Listing 46: Random Forest with randomForest Package (Slide L11 p.18)

2.12.3 Variable Importance Measures (for Bagging/RF)

∗ Trees make it hard to see overall variable importance when many trees are averaged.

* **Mean Decrease in Impurity (e.g., RSS or Gini)** (Slide L11 p.22):

  · For each tree, record total reduction in RSS (regression) or Gini index (classification) due to splits on each predictor.

  · Average this reduction over all B trees for each predictor.

  · Larger value $\implies$ more important.

  · R: `importance(rf.object)()` gives this (`IncNodePurity` column), `varImpPlot(rf.c`

* **Permutation Importance (Mean Decrease in Accuracy/MSE)**:

  · For each tree, calculate OOB error.

  · Then, for each predictor X_j:

  · Randomly permute the values of X_j in the OOB data for that tree.

  · Recalculate OOB error with permuted X_j.

  · The average increase in OOB error across all trees due to permuting X_j is its importance.

  · R: `importance(rf.object, type=1)()` or `importance=TRUE` in `randomForest()` and check `%IncMSE` (regression) or `MeanDecreaseAccuracy` (classification).

* R Example (`Carseats` Bagging, Slides L11 p.23):

```
1  # Assuming bagging.carseats from before
2  # varImpPlot(bagging.carseats)
3  # # Shows ShelveLoc and Price as most important by
      IncNodePurity (Gini/RSS decrease)
```

Listing 47: Variable Importance Plot (Slide L11 p.23)

2.12.4 Boosting

* **Core Idea**: Sequentially fit trees. Each new tree is fit to the *residuals* of the previous model, thus focusing on observations the model currently mispredicts.

* This is different from Bagging/RF where trees are grown independently.

* **Boosting Algorithm for Regression Trees** (Slide L11 p.25):

  1. Initialize $\hat{f}(x) = 0$ and residuals $r_i = y_i$ for all i.

  2. For $b = 1, \ldots, B$ (number of trees):

     (a) Fit a tree $\hat{f}^b$ with d splits (often small, e.g., $d = 1$ for stumps) to the training data (X, r) (i.e., predict current residuals).

     (b) Update the overall prediction: $\hat{f}(x) \leftarrow \hat{f}(x) + \lambda \hat{f}^b(x)$.

     (c) Update residuals: $r_i \leftarrow r_i - \lambda \hat{f}^b(x_i)$.

  3. The final boosted model is $\hat{f}(x) = \sum_{b=1}^{B} \lambda \hat{f}^b(x)$. (Slide L11 p.25 has λ multiplying sum, which is equivalent if λ is constant).

* **Tuning Parameters for Boosting**:

  · **Number of Trees** B: Unlike Bagging/RF, boosting can overfit if B is too large. Chosen by CV.

- **Shrinkage Parameter** λ (learning rate): Small positive number (e.g., 0.01, 0.1). Controls rate at which boosting learns. Smaller λ requires larger B.

- **Interaction Depth** d (or tree depth): Number of splits in each tree. $d = 1$ gives an additive model. $d > 1$ allows for interactions.

* R Example (Simulated data, Slides L11 p.26-31, `Hitters` data, Slides L11 p.32-37):

```
1  # library(ISLR); library(gbm)
2  # data(Hitters)
3  # Hitters <- Hitters[complete.cases(Hitters$Salary),]
4  # Hitters$logSalary <- log(Hitters$Salary)
5  # set.seed(1234) # From slide
6  # itrain_hit <- sample(1:nrow(Hitters), floor(nrow(Hitters)/2))
7  # traindata_hit <- Hitters[itrain_hit,]
8  # testdata_hit <- Hitters[-itrain_hit,]
9  # ytest_hit <- testdata_hit$logSalary
10
11 # Define formula (slide uses a subset of variables)
12 # formula1_hit <- logSalary ~ CRuns+CHits+HmRun+CWalks+CRBI+
     Hits+RBI
13
14 # Fit boosted model
15 # boost.hitters <- gbm(formula1_hit, data=traindata_hit,
     distribution="gaussian",
16 #                      n.trees=1000, interaction.depth=4,
     shrinkage=0.01)
17 # summary(boost.hitters) # Shows variable importance
18
19 # Predict and calculate Test MSE
20 # pred.gbm <- predict(boost.hitters, newdata=testdata_hit, n.
     trees=1000)
21 # mse.gbm <- mean((ytest_hit - pred.gbm)^2)
22 # # mse.gbm = 0.3591336 from slide L11 p.36
23 # # mse.simple_tree = 0.3851602 from slide L11 p.35 - Boosting
     improves.
```

Listing 48: Boosting with gbm Package (Hitters data, Slides L11 p.36)

2.12.5 Exercises: ISLR 8.8d-e, 8.11

* **ISLR 8.8d-e (`Carseats` data)**: (Corresponds to 'ex8.8.R' file, latter part)

- (d) Bagging: Fit bagged trees for `Sales`. Compute test MSE. Variable importance. *From 'ex8.8.R'*: Bagging with `mtry=10` (all predictors) gives Test MSE $\approx 2.5 - 2.6$. Price, ShelveLoc, Age are important.

- (e) Random Forest: Fit RF for `Sales`. Vary `mtry`. Compute test MSE. Variable importance. *From 'ex8.8.R'*: RF with `mtry=5` (approx $p/2$) gives Test MSE $\approx 2.8 - 2.9$. Sometimes bagging performed slightly better for this dataset/seed.

* **ISLR 8.11 (`Caravan` data)**: (Corresponds to 'ch8-applied.R' Exercise 11)

- Predict `Purchase`. Highly imbalanced data.
- (b) Boosting: Fit boosted trees. Variable importance.

- · (c) Evaluate Boosting: Predict on test set. Use threshold (e.g., 0.2). Compare with logistic regression. Focus on fraction of correctly predicted "Yes" among those predicted "Yes".
- · *From 'ch8-applied.R' Ex 11*: Boosting (threshold 0.2) correctly identifies 20

2.13 Lecture 13: Support Vector Machines (SVMs)

2.13.1 Introduction to Support Vector Machines

- ∗ SVMs are a powerful class of supervised learning algorithms used for classification and regression (though primarily classification is covered).
- ∗ We focus on binary classification (two classes).
- ∗ Progression of concepts:

  1. **Maximal Margin Classifier**: Assumes classes are perfectly separable by a linear boundary.
  2. **Support Vector Classifier (SVC)**: Extends maximal margin classifier to handle non-separable classes by allowing some misclassifications (soft margin). Still linear boundaries.
  3. **Support Vector Machine (SVM)**: Extends SVC to accommodate non-linear boundaries using kernels.

2.13.2 Hyperplanes

- ∗ A hyperplane in a p-dimensional space is a flat affine subspace of dimension $p-1$.
- ∗ Equation: $\beta_0 + \beta_1 X_1 + \beta_2 X_2 + \cdots + \beta_p X_p = 0$.
  - · If $p = 2$: Line, $X_2 = -\frac{\beta_0}{\beta_2} - \frac{\beta_1}{\beta_2} X_1$.
  - · If $p = 3$: Plane.
- ∗ A hyperplane divides the p-dimensional space into two half-spaces.

2.13.3 Classification by Separating Hyperplanes

- ∗ Data: n training observations $(\mathbf{x}_i, y_i)$, where $\mathbf{x}_i \in \mathbb{R}^p$ and $y_i \in \{-1, 1\}$.
- ∗ Decision rule for a test observation $\mathbf{x}$:
  - · Classify as $+1$ if $f(\mathbf{x}) = \beta_0 + \beta_1 x_1^+ \cdots + \beta_p x_p^> 0$.
  - · Classify as -1 if $f(\mathbf{x}) < 0$.

  (Slide L12 p.6-7 illustrates this for $p = 2$).
- ∗ For a separating hyperplane, all training observations satisfy $y_i f(\mathbf{x}_i) > 0$.
- ∗ The magnitude $|f(\mathbf{x})|$ indicates confidence in the classification (larger magnitude $\implies$ further from hyperplane).

2.13.4 The Maximal Margin Classifier (MMC)

* If classes are linearly separable, there can be infinitely many separating hyperplanes.
* **MMC**: The separating hyperplane that is farthest from the closest training observations from either class. (Slide L12 p.8).
* **Margin (M)**: The perpendicular distance from the hyperplane to the nearest training observation. MMC maximizes M.
* **Support Vectors**: Training observations that lie exactly on the margin (equidistant from the hyperplane). They "support" the hyperplane; if moved, the hyperplane would change.
* **Optimization Problem** (Slide L12 p.9):

$$\max_{\beta_0,\ldots,\beta_p,M} M$$

Subject to:

$$\sum_{j=1}^{p} \beta_j^2 = 1 \quad \text{(Normalization constraint)} \tag{1}$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip}) \geq M \quad \forall i = 1,\ldots,n \tag{2}$$

The term $y_i f(\mathbf{x}_i)$ is the functional margin. With $\sum \beta_j^2 = 1$, M becomes the geometric margin.

* Slide L12 p.10 shows the maximal margin hyperplane (black) and the margin boundaries (red/blue).
* **Limitation**: Only exists if classes are perfectly linearly separable. Sensitive to outliers. (Slide L12 p.11 shows a non-separable case).

2.13.5 Support Vector Classifier (SVC) / Soft Margin Classifier

* Handles non-linearly separable data or cases where a wider margin with a few errors is preferred.
* Allows some observations to be on the wrong side of the margin, or even on the wrong side of the hyperplane.
* **Optimization Problem** (Slide L12 p.12):

$$\max_{\beta_0,\ldots,\beta_p,\epsilon_1,\ldots,\epsilon_n,M} M$$

Subject to:

$$\sum_{j=1}^{p} \beta_j^2 = 1 \tag{3}$$

$$y_i(\beta_0 + \beta_1 x_{i1} + \cdots + \beta_p x_{ip}) \geq M(1 - \epsilon_i) \quad \forall i \tag{4}$$

$$\epsilon_i \geq 0, \quad \sum_{i=1}^{n} \epsilon_i \leq C \tag{5}$$

* **Slack Variables** ϵ_i (Slide L12 p.13):

$\cdot$ $\epsilon_i = 0$: Observation i is on the correct side of its margin.

$\cdot$ $0 < \epsilon_i \le 1$: Observation i is on the correct side of the hyperplane but on the wrong side of its margin (margin violation).

$\cdot$ $\epsilon_i > 1$: Observation i is on the wrong side of the hyperplane (misclassified).

$*$ **Tuning Parameter** C (Cost/Budget for violations, Slide L12 p.14):

$\cdot$ Controls the trade-off between margin width and number/severity of margin violations.

$\cdot$ Small C: Wide margin, more violations allowed (high bias, low variance). Similar to larger s in $\sum \beta_j^2 \le s$.

$\cdot$ Large C: Narrow margin, fewer violations allowed (low bias, high variance). Approaches MMC if data is separable.

$\cdot$ C is chosen by cross-validation.

$*$ Only observations on the margin or violating the margin (support vectors) affect the solution. (Slide L12 p.13)

$*$ Still produces a linear decision boundary.

2.13.6 Support Vector Machines (SVMs) - Non-linear Boundaries

$*$ Extends SVC to handle non-linear decision boundaries by using **kernels**.

$*$ **The Kernel Trick**: Instead of explicitly mapping data to a higher-dimensional space to find a linear separator, kernels compute inner products in that higher-dimensional space directly using original data.

$*$ The decision function becomes (Slide L12 p.19, simplified from dual form):

$$f(\mathbf{x}) = \beta_0 + \sum_{i \in \mathcal{S}} \alpha_i K(\mathbf{x}, \mathbf{x}_i)$$

where $\mathcal{S}$ is the set of support vectors, α_i are parameters, and $K(\cdot, \cdot)$ is a kernel function.

$*$ **Common Kernels** (Slide L12 p.19):

$\cdot$ **Linear Kernel**: $K(\mathbf{x}_i, \mathbf{x}_{i'}) = \mathbf{x}_i' \mathbf{x}_{i'}$. Recovers the SVC.

$\cdot$ **Polynomial Kernel** of degree d: $K(\mathbf{x}_i, \mathbf{x}_{i'}) = (1 + \gamma \mathbf{x}_i' \mathbf{x}_{i'} + r)^d$. (Slide has $1 + \sum x_{ij} x_{i'j}$). Tuning parameters: d, γ, cost C.

$\cdot$ **Radial Basis Function (RBF) Kernel / Gaussian Kernel**:

$$K(\mathbf{x}_i, \mathbf{x}_{i'}) = \exp\left(-\gamma \sum_{j=1}^{p} (x_{ij} - x_{i'j})^2\right) = \exp(-\gamma \|\mathbf{x}_i - \mathbf{x}_{i'}\|^2)$$

Tuning parameters: $\gamma > 0$, cost C.

$\cdot$ If γ is very large, fit is very local (wiggly, high variance).

$\cdot$ If γ is very small, fit is very smooth (like linear, high bias).

2.13.7 R Implementation of SVMs

* Main function: `svm()` from `e1071`.

* Key arguments:

  · `formula` or `x, y` input.

  · `data`: Training data.

  · `kernel`: "linear", "polynomial", "radial" (default), "sigmoid".

  · `cost`: Parameter C (SVC/SVM).

  · `degree`: For polynomial kernel.

  · `gamma`: For polynomial and radial kernels.

  · `scale = TRUE/FALSE`: Whether to standardize predictors. (Default TRUE). (Slide L12 p.17 shows FALSE for manual data).

* **Tuning Parameters** using `tune()` from `e1071`:

```
1  # library(e1071)
2  # set.seed(1) # For reproducibility
3  # # Example data from ISLR Ch9 Lab / Slides L12 p.15
4  # x_sim_svm <- matrix(rnorm(202), ncol=2)
5  # y_sim_svm <- c(rep(-1,10), rep(1,10))
6  # x_sim_svm[y_sim_svm==1,] <- x_sim_svm[y_sim_svm==1,] + 1 #
       Create separable classes
7  # dat_sim_svm <- data.frame(x=x_sim_svm, y=as.factor(y_sim_svm)
       )
8
9  # # Tune linear SVM for cost
10 # tune.out.linear <- tune(svm, y ~ ., data=dat_sim_svm, kernel
       ="linear",
11 #                          ranges=list(cost=c(0.001, 0.01, 0.1,
       1, 5, 10, 100)))
12 # summary(tune.out.linear) # Shows best cost and CV error
13 # bestmod.linear <- tune.out.linear$best.model
14 # summary(bestmod.linear) # Shows parameters of best model,
       number of SVs
15
16 # # Plotting the SVM decision boundary (Slides L12 p.17)
17 # # svmfit.linear <- svm(y ~ ., data=dat_sim_svm, kernel="
       linear", cost=10, scale=FALSE)
18 # # plot(svmfit.linear, dat_sim_svm) # For 2D data
19
20 # # Example with radial kernel (Slides L12 p.21)
21 # # svmfit.radial <- svm(y ~ ., data=dat_sim_svm, kernel="
       radial", cost=10, gamma=1, scale=FALSE)
22 # # plot(svmfit.radial, dat_sim_svm)
```

Listing 49: Tuning SVM with e1071::tune (ISLR Ch9 Lab, Ex 9.7)

* **Support Vectors**: Can be accessed via `svmfit$index`.

2.13.8 SVMs with More than Two Classes

* **One-Versus-One (OVO)**: Construct $\binom{K}{2}$ SVMs, each comparing a pair of classes. Classify a new observation by majority vote among these classifiers.

* **One-Versus-All (OVA)**: Fit K SVMs, each separating one class from the rest. Classify to the class whose SVM gives largest $|f(\mathbf{x})|$ (or highest probability if calibrated).

* `e1071` typically uses OVO for multi-class SVM.

2.13.9   Relationship to Logistic Regression

* SVC optimization problem has similarities to logistic regression, particularly when L_2 penalty is used with logistic regression.

* Loss function for SVC (hinge loss) vs. logistic loss. Hinge loss is zero for points correctly classified beyond the margin.

* In practice, performance can be similar, especially with linear boundaries. SVMs can be more robust to outliers when classes are well-separated due to focus on support vectors.

2.13.10   Exercise: ISLR 9.7

(Corresponds to 'ex9.7.R')

* (a) Create binary variable `high`: 1 if `mpg > median(mpg)`, 0 otherwise.

* (b) SVC (linear kernel): Use `tune()` to find best `cost`. *From 'ex9.7.R'*: Best cost is often 1.

* (c) SVM with polynomial and radial kernels: Use `tune()` for `cost`, `degree` (poly), `gamma` (radial). *From 'ex9.7.R'*:

    · Polynomial: E.g., cost=10, degree=2.

    · Radial: E.g., cost=10, gamma=0.01.

* (d) Plot results and compare predictions (e.g., training confusion matrices).

    · Radial kernel often gives good separation on training data, but generalization to test data is key.

    · Linear SVM with optimal cost usually performs reasonably well.

```r
# library(ISLR); library(e1071)
# Auto$high <- as.factor(ifelse(Auto$mpg > median(Auto$mpg), 1,
    0))
# Auto_svm <- Auto[, !(names(Auto) %in% c("mpg", "name"))] #
    Exclude original mpg and name

# set.seed(1)
# tune.radial <- tune(svm, high ~ ., data=Auto_svm, kernel="
    radial",
#                     ranges=list(cost=c(0.1, 1, 10, 100),
    gamma=c(0.01, 0.1, 1, 5)))
# summary(tune.radial)
# best.svm.radial <- tune.radial$best.model
# plot(best.svm.radial, Auto_svm, horsepower ~ weight) #
    Example plot
```

Listing 50: ISLR 9.7 - Tuning SVM (Example for Radial)

2.14 Lecture 14: Unsupervised Learning

2.14.1 Introduction to Unsupervised Learning

* **Key Characteristic**: We only observe predictors $\mathbf{X}$, no response variable y.
* **Goals**:
  · Exploratory / descriptive data analysis.
  · Detect unusual/remarkable patterns or subgroups in observations.
  · Generate hypotheses about the data structure.
  · Can be seen as "Advanced descriptive analysis."

* **Main Techniques Covered**:
  · Principal Component Analysis (PCA) for dimension reduction and visualization.
  · Clustering (K-Means, Hierarchical) for finding subgroups.

2.14.2 Principal Component Analysis (PCA)

* **Motivation** (Slides L13 p.3-4):
  · Visualizing high-dimensional data ($p > 2$) is difficult. A pairs plot for $p = 10$ yields $\binom{10}{2} = 45$ scatterplots.
  · PCA aims to summarize p variables into a smaller set of $M < p$ new variables (principal components) that capture as much of the original variability as possible.

* **What are Principal Components?**
  · Principal components (PCs) are linear combinations of the original predictors:

$$Z_m = \phi_{1m}X_1 + \phi_{2m}X_2 + \cdots + \phi_{pm}X_p = \sum_{j=1}^{p} \phi_{jm}X_j$$

  · The coefficients ϕ_{jm} are called **loadings**. The loading vector for PC m is $\boldsymbol{\phi}_m = (\phi_{1m}, \ldots, \phi_{pm})'$.
  · Loadings define the direction in feature space along which the data varies the most (for Z_1), or varies most subject to being uncorrelated with previous PCs (for $Z_2, Z_3, \ldots$).

* **First Principal Component (Z_1)** (Slide L13 p.5):
  · Z_1 is the normalized linear combination of X_j's that has the largest variance:

$$\text{maximize } \text{Var}(Z_1) = \text{Var}\left(\sum_{j=1}^{p} \phi_{j1}X_j\right)$$

  subject to the normalization constraint $\sum_{j=1}^{p} \phi_{j1}^2 = 1$.
  · The loadings ϕ_{j1} define the first principal component direction.

* **Second Principal Component (Z_2)** (Slide L13 p.6):

· Z_2 is the normalized linear combination of X_j's that has the largest variance, subject to being uncorrelated with Z_1 ($\text{Cov}(Z_1, Z_2) = 0$).

$$\text{maximize Var}(Z_2) \quad \text{subject to} \quad \sum_{j=1}^{p} \phi_{j2}^2 = 1 \quad \text{and} \quad \text{Cov}(Z_1, Z_2) = 0$$

* Subsequent PCs are defined similarly (maximize variance, uncorrelated with all previous PCs). Up to $\min(n-1, p)$ PCs can be constructed.

* **Geometric Interpretation**: PCs represent new axes in the feature space. Z_1 is the direction of greatest spread, Z_2 is the next direction of greatest spread orthogonal to Z_1, etc.

* **Scaling of Variables**: It is crucial to scale variables to have standard deviation one (and often mean zero) before performing PCA if they are measured on different scales. Otherwise, variables with larger variance will dominate the PCs. (Slide L13 p.13 uses `scale=TRUE` with `prcomp()`).

* **Proportion of Variance Explained (PVE)**:

  · The variance explained by PC m is $\text{Var}(Z_m)$. Total variance is $\sum_{j=1}^{p} \text{Var}(X_j)$ (if scaled, this is p).

  · $PVE_m = \frac{\text{Var}(Z_m)}{\sum_{j=1}^{p} \text{Var}(X_j)}$.

  · Cumulative PVE helps decide how many PCs to retain.

* **Scree Plot**: Plot of PVE for each PC (or eigenvalues $\text{Var}(Z_m)$). Look for an "elbow" to decide number of components to keep.

* **R Implementation** (`prcomp()` or `princomp()`):

  · `prcomp()` is generally preferred (uses SVD).

  · `rotation` element contains loadings ϕ_{jm}.

  · `x` element contains the principal component scores z_{im}.

  · `sdev` contains standard deviations of PCs (square for variances).

* R Example (Simulated $p = 3$ data, Slides L13 p.7-10):

```
# Assume 'x' is the n x 3 matrix of simulated data
# pca_sim <- princomp(x) # Or prcomp(x, scale.=TRUE, center.=
    TRUE)
# summary(pca_sim)
# Output from slide L13 p.9:
# Importance of components:
#                           Comp.1    Comp.2       Comp.3
# Standard deviation     1.5940343 0.6333490 0.031345120
# Proportion of Variance 0.8633688 0.1362973 0.000333842
# Cumulative Proportion  0.8633688 0.9996662 1.000000000
# Almost all variance (99.97%) captured by first two PCs.

# plot(pca_sim$scores[,1:2]) # Plot observations in PC1-PC2
    space (Slide L13 p.10)
```

Listing 51: PCA on Simulated 3D Data (Slides L13 p.9-10)

* R Example (`ISLR USArrests` data, Slides L13 p.11-16):

```
1   # library(ISLR)
2   # data(USArrests)
3   # # summary(USArrests) # Variables have different scales/means
4   # # apply(USArrests, 2, mean)
5   # # apply(USArrests, 2, var)
6
7   # pca.usarrests <- prcomp(USArrests, scale=TRUE, center=TRUE)
8   # summary(pca.usarrests)
9   # # Importance of components:
10  # #                          PC1    PC2    PC3     PC4
11  # # Standard deviation     1.5749 0.9949 0.59713 0.41645
12  # # Proportion of Variance 0.6201 0.2474 0.08914 0.04336
13  # # Cumulative Proportion  0.6201 0.8675 0.95664 1.00000
14  # # First two PCs explain ~87% of variance.
15
16  # pca.usarrests$rotation # Loadings (Slide L13 p.15)
17  # # PC1: High negative loadings for Murder, Assault, Rape (
        crime factor), negative for UrbanPop.
18  # #       States with low PC1 scores have high crime rates.
19  # # PC2: High positive loading for UrbanPop, small negative for
        Assault. (urbanization factor)
20  # #       States with high PC2 scores are highly urbanized.
21
22  # # Biplot (Slide L13 p.14 shows state names, Slide L13 p.16
        shows arrows for variables)
23  # biplot(pca.usarrests, scale=0) # scale=0 ensures arrows
        represent loadings
```

Listing 52: PCA on USArrests Data (Slides L13 p.13-16)

2.14.3 Clustering Methods

* Goal: Find homogeneous subgroups (clusters) among observations. Observations within a cluster should be similar; observations in different clusters should be dissimilar.

* "Similarity" is context-dependent, often based on Euclidean distance between observations in feature space.

* Examples (Slide L13 p.17): Classifying plants based on features, identifying consumer subgroups based on shopping patterns.

2.14.4 K-Means Clustering

* **Objective**:

  · Partition n observations into a pre-specified number of K clusters $C_1, \ldots, C_K$.

  · Clusters must be non-overlapping ($C_k \cap C_{k'} = \emptyset$) and cover all observations ($\cup C_k = \{1, \ldots, n\}$).

  · Minimize within-cluster variation (WCV). A common measure is sum of squared Euclidean distances between observations in a cluster and the cluster centroid, or sum of pairwise squared Euclidean distances within a cluster.

$$\min_{C_1,\ldots,C_K} \sum_{k=1}^{K} W(C_k)$$

75

where $W(C_k) = \frac{1}{|C_k|} \sum_{i,i' \in C_k} \sum_{j=1}^{p} (x_{ij} - x_{i'j})^2$. This is equivalent to minimizing $2 \sum_{k=1}^{K} \sum_{i \in C_k} \sum_{j=1}^{p} (x_{ij} - \bar{x}_{kj})^2$, where $\bar{x}_{kj}$ is the mean of feature j for cluster k (centroid). (ISLR Ch10 Exercise 10.1).

* **Algorithm** (Greedy, iterative, Slide L13 p.20):

  1. Randomly assign each observation to one of K clusters.

  2. Iterate until cluster assignments stop changing:

     (a) For each cluster k, compute its centroid $\boldsymbol{\mu}_k$ (vector of feature means for observations in C_k).

     (b) Reassign each observation to the cluster whose centroid is closest (e.g., by Euclidean distance).

* **Sensitivity to Initialization**: Final clustering can depend on initial random assignment. Run algorithm multiple times with different random starts (argument `nstart` in R's `kmeans()`) and choose the solution with smallest total WCV.

* **Choosing K**: No single best method.

  · Elbow method: Plot total WCV (or within-cluster sum of squares, `tot.withinss`) vs. K. Look for an "elbow" point where adding more clusters gives diminishing returns.

  · Silhouette analysis.

  · Domain knowledge or practical considerations.

* R Example (Simulated 2D data, 3 clusters, Slides L13 p.21-27):

```
# Assume 'x_cluster_sim' is an n x 2 matrix of simulated data
    with 3 true clusters
# K_val <- 3

# Manual iteration steps (conceptual from slides L13 p.22-25)
# 1. Random initial assignment:
#     cluster_colors <- c("black", "red", "blue")
#     initial_assignment <- sample(1:K_val, size=nrow(x_cluster_
    sim), replace=TRUE)
#     plot(x_cluster_sim, col=cluster_colors[initial_assignment
    ])
#     centroids <- matrix(0, K_val, ncol(x_cluster_sim))
#     for(k_idx in 1:K_val) {
#        centroids[k_idx,] <- colMeans(x_cluster_sim[initial_
    assignment==k_idx,])
#        points(centroids[k_idx,1], centroids[k_idx,2], pch=17,
    col=cluster_colors[k_idx], cex=2)
#     }
# 2. Reassign based on distance to centroids, recompute
    centroids ... (iterate)

# Using kmeans() function (Slide L13 p.26-27)
# set.seed(123) # For reproducibility of kmeans if nstart=1
# km.out <- kmeans(x_cluster_sim, centers=K_val, nstart=20) #
    nstart=20 runs 20 random initializations
# names(km.out) # "cluster", "centers", "totss", "withinss", "
    tot.withinss", "betweenss", "size"
#
```

```
21  # plot(x_cluster_sim, col=cluster_colors[km.out$cluster], main
        ="K-Means Clustering Result (K=3)")
22  # points(km.out$centers, col=cluster_colors, pch=17, cex=2, lwd
        =2)
```

Listing 53: K-Means Clustering in R (Slides L13 p.22, 26-27)

2.14.5 Hierarchical Clustering

* Does not require pre-specifying K.

* Produces a tree-based representation (dendrogram) of the observations.

* **Agglomerative (Bottom-Up) Approach** (Algorithm on Slide L13 p.32):

    1. Start with each of n observations as its own cluster.

    2. Compute all $n(n-1)/2$ pairwise dissimilarities (e.g., Euclidean distances).

    3. For $i = n, n-1, \ldots, 2$:

        (a) Fuse the two closest (most similar) clusters. The distance at which they fuse becomes the height in the dendrogram.

        (b) Compute new pairwise inter-cluster dissimilarities among the $i - 1$ remaining clusters.

* **Linkage Methods** (How to define dissimilarity between clusters, Slide L13 p.33):

    · *Complete*: Maximal inter-cluster dissimilarity (largest distance between an obs in cluster A and an obs in cluster B).

    · *Single*: Minimal inter-cluster dissimilarity (smallest distance). Can result in "chaining."

    · *Average*: Mean inter-cluster dissimilarity.

    · *Centroid*: Dissimilarity between centroids of two clusters. Can lead to inversions in dendrogram (later fusions at lower height).

    · *Ward's method*: Minimizes increase in total within-cluster variance upon merging.

* **Dendrogram** (Slide L13 p.30, USArrests example):

    · Leaves are individual observations.

    · Fusions are represented by horizontal lines; height indicates dissimilarity at fusion.

    · Cut dendrogram at a certain height to obtain K clusters. (Slide L13 p.31)

* Choice of dissimilarity measure (Euclidean, correlation-based, etc.) and linkage method can significantly affect results.

* R Example (`USArrests` data, Slide L13 p.29-30):

```
1  # data(USArrests)
2  # # Hierarchical clustering with complete linkage and Euclidean
        distance
3  # hc.complete <- hclust(dist(USArrests), method="complete")
4  # plot(hc.complete, main="Complete Linkage", xlab="", sub="",
        cex=0.9)
```

77

```
 5  #
 6  # # Cutting the tree to get specific number of clusters
 7  # cutree(hc.complete, k=3) # Get 3 clusters
 8  #
 9  # # Scaling data is often recommended before calculating
       distances
10  # hc.complete.scaled <- hclust(dist(scale(USArrests)), method="
       complete")
11  # plot(hc.complete.scaled, main="Complete Linkage (Scaled Data)
       ")
```

Listing 54: Hierarchical Clustering in R (Slides L13 p.30)

2.14.6 Practical Issues in Clustering

* **Scaling Variables**: Important if variables are on different scales, as distance measures will be dominated by variables with larger magnitudes/variances. Standardize to mean 0, std dev 1.
* **Choice of K (for K-Means) or Cut Height (for Hierarchical)**: Often subjective or guided by external criteria/domain knowledge.
* **Interpreting Clusters**: Examine feature means/distributions within each cluster.
* Clustering is exploratory; results should be interpreted with caution and domain expertise.

2.14.7 Exercises: ISLR 12.8, 12.9 (

* **ISLR 12.8 (old 10.8, PCA on USArrests)** (Corresponds to 'ch10-lab1.R', 'ch10ex8.R'):
    · (a) Calculate PVE for each PC.
    · (b) Relate PVE to sum of squared PC scores and sum of squared scaled original data. (PVE for PC m is $\frac{\sum_{i=1}^{n} z_{im}^2}{\sum_{j=1}^{p} \sum_{i=1}^{n} \tilde{x}_{ij}^2}$, where z_{im} are scores for PC m, and $\tilde{x}_{ij}$ are scaled original data.)
* **ISLR 12.9 (old 10.9, Hierarchical Clustering on USArrests)** (Corresponds to 'ch10-lab2.R', 'ch10ex9.R'):
    · (a) Hierarchical clustering with complete linkage, Euclidean distance. Plot dendrogram.
    · (b) Cut tree to get 3 clusters. Identify states in each.
    · (c) Repeat with scaled data. Plot.
    · (d) Compare results. Scaling can change cluster assignments significantly. Discuss if scaling is appropriate (generally yes if variables have different units/scales).

3 Previous Exams

3.1 Exam Spring 2021

Overall Datasets: `Smarket` (ISLR), `dataCar(insuranceData)`

Task 1: Analysis of `Smarket` Data

3.1.1 1a: Bootstrap histogram for volatility

Question *Use the bootstrap to create a histogram of the sampling distribution for the volatility (i.e., the standard deviation of the returns) of the S&P stock index. Observations of the returns are in the variable `Today`. Start by setting the seed to 1 (`set.seed(1)`). Use 1000 bootstrap replicates.*

Base R Solution

```r
library(ISLR)
library(boot)
data(Smarket)

set.seed(1)
calculate_sd_Smarket <- function(data_vector, index) {
  return(sd(data_vector[index]))
}
boot_volatility_Smarket <- boot(data = Smarket$Today, statistic =
    calculate_sd_Smarket, R = 1000)
hist(boot_volatility_Smarket$t, main = "Bootstrap Dist. of
    Volatility (Smarket$Today)",
    xlab = "Estimated Volatility (Std Dev)", col = "lightblue",
        breaks = "Scott")
```

Listing 55: 1a: Bootstrap Volatility - Base R (Exam 2021)

Tidyverse Solution (Data Prep + Base R boot)

```r
library(ISLR); library(boot); library(dplyr); library(ggplot2)
data(Smarket)
today_returns_Smarket <- Smarket %>% pull(Today)
set.seed(1)
# calculate_sd_Smarket function is the same
boot_volatility_tidy_Smarket <- boot(data = today_returns_Smarket,
    statistic = calculate_sd_Smarket, R = 1000)
ggplot(data.frame(volatility = boot_volatility_tidy_Smarket$t), aes(
    x = volatility)) +
  geom_histogram(aes(y=..density..), fill = "lightblue", color = "
      black", bins=30) +
  geom_density(alpha=.2, fill="#FF6666") +
  labs(title="Bootstrap Distribution of Volatility (Smarket$Today)",
      x="Estimated Volatility", y="Density")
```

Listing 56: 1a: Bootstrap Volatility - Tidyverse (Exam 2021)

3.1.2 1b: 95% CI for volatility (normal assumption)

Question *Compute a 95% confidence interval based on the bootstrap, assuming that the volatility is normally distributed.*

Base R Solution (using `boot`)

```r
# Assuming boot_volatility_Smarket from 1a
ci_normal_Smarket <- boot.ci(boot_volatility_Smarket, type = "norm")
print(ci_normal_Smarket)
# Interpretation: Based on the normal approximation, we are 95%
    confident the true volatility is between [lower] and [upper].
```

Listing 57: 1b: Normal Bootstrap CI - Base R boot (Exam 2021)

3.1.3   1c: 95% CI for volatility (percentile method)

Question   *Compute a 95% confidence interval based on the bootstrap, not making the normality assumption but use percentiles from the bootstrapped volatilities.*

Base R Solution (using `boot`)

```
# Assuming boot_volatility_Smarket from 1a
ci_percentile_Smarket <- boot.ci(boot_volatility_Smarket, type = "
    perc")
print(ci_percentile_Smarket)
# Interpretation: Using percentiles, we are 95% confident the true
    volatility is between [lower] and [upper].
```

Listing 58: 1c: Percentile Bootstrap CI - Base R boot (Exam 2021)

3.1.4   1d: Regression of Squared Returns

Question   *Squared returns, `Today^2`, can be used as a proxy for volatility. Squared returns are often autocorrelated, indicating that volatility can be predicted. Run a regression of the square of `Today` on the square of `Lag1` and interpret the estimate of the regression parameter for the square of `Lag1`.*

Base R Solution

```
model_sq_returns_Smarket <- lm(I(Today^2) ~ I(Lag1^2), data =
    Smarket)
summary_model_sq_Smarket <- summary(model_sq_returns_Smarket)
print(summary_model_sq_Smarket)
# Interpretation of I(Lag1^2) coefficient:
# A one-unit increase in Lag1^2 (yesterday's squared return) is
    associated
# with an estimated [coefficient_value, e.g., 0.165] change in Today
    ^2 (today's squared return).
# This relationship is statistically significant (p-value = [p_val])
    .
```

Listing 59: 1d: Regression of Squared Returns - Base R (Exam 2021)

3.1.5   1e: Bootstrap SE of regression coefficient

Question   *Compute the bootstrapped standard error of the regression coefficient in front of the square of `Lag1` in the model in d) and compare it with the standard error from the regression output.*

Base R Solution

```
coef_sq_lag1_Smarket_bs <- function(data, index) {
  subset_data <- data[index, ]
  fit <- lm(I(Today^2) ~ I(Lag1^2), data = subset_data)
  return(coef(fit)[2])
}
set.seed(1) % \textit{(Or seed from marking scheme for this specific
    subtask if different, e.g., if a variant was used.)} %
```

```
7  boot_coef_Smarket <- boot(data = Smarket, statistic = coef_sq_lag1_
       Smarket_bs, R = 1000)
8  bootstrap_se_coef_Smarket <- sd(boot_coef_Smarket$t)
9  cat("Bootstrap SE for I(Lag1^2) coef (Smarket):", bootstrap_se_coef_
       Smarket, "\n")
10
11 lm_se_coef_Smarket <- summary(model_sq_returns_Smarket)$coefficients
       ["I(Lag1^2)", "Std. Error"]
12 cat("LM Output SE for I(Lag1^2) coef (Smarket):", lm_se_coef_Smarket
       , "\n")
13 # Comparison: The bootstrap SE ([value]) is [larger/smaller/similar]
        than the lm output SE ([value]).
14 # Marking scheme: Bootstrap SE (0.0488) was almost twice lm SE
       (0.0279).
```
Listing 60: 1e: Bootstrap SE for Coefficient - Base R (Exam 2021)

Task 2: Analysis of `dataCar` Data - Predicting `claimcst0`

3.1.6 2a: Data Filtering and Variable Removal

Question *Start by first selecting the observations where there is an occurrence of claim, `clm!=0`. Also, remove the variable `X_OBSTAT_`. Work with this new data set for the rest of Task 2.*

Base R Solution
```
1  library(insuranceData)
2  data(dataCar)
3  xdata_task2_car <- dataCar[dataCar$clm != 0, ]
4  xdata_task2_car <- xdata_task2_car[, names(xdata_task2_car) != "X_
       OBSTAT_"]
5  # print(paste("Dimensions after filtering:", paste(dim(xdata_task2_
       car), collapse="x")))
```
Listing 61: 2a: Filter and Remove (dataCar) - Base R (Exam 2021)

Tidyverse Solution
```
1  library(insuranceData); library(dplyr)
2  data(dataCar)
3  xdata_task2_car_tidy <- dataCar %>%
4    filter(clm != 0) %>%
5    select(-any_of("X_OBSTAT_"))
6  # glimpse(xdata_task2_car_tidy)
```
Listing 62: 2a: Filter and Remove (dataCar) - Tidyverse (Exam 2021)

3.1.7 2b: Descriptive statistics and variable formatting

Question *Use descriptive statistics, the description of the data in the help-function and common sense to remove variables that you think will not be helpful. Motivate your choices thoroughly. Also, make sure that all variables are on the right format.*

Base R / Tidyverse Solution (Conceptual - based on marking scheme)

```
1  # Assuming xdata_task2_car from 2a
2  # Marking scheme solution: Remove clm (col 3 in their xdata after X_
       OBSTAT_ removal)
3  # Original columns: veh_value, exposure, clm, numclaims, claimcst0,
       veh_body, veh_age, gender, area, agecat
4  # After X_OBSTAT_ removal and clm filtering, clm column is now
       constant 1.
5  # If X_OBSTAT_ was, say, col 11 as in marking scheme:
6  # data(dataCar); xdata_task2_car <- dataCar[dataCar$clm !=0,]; xdata
       _task2_car <- xdata_task2_car[,-11]
7  # Original column indices matter for marking scheme solution.
8  # Let's assume xdata_task2_car is as per previous step (X_OBSTAT_
       removed).
9  # 'clm' column is now all 1s.
10
11 # Removing 'clm' as it's constant after filtering:
12 xdata_task2_car_processed <- xdata_task2_car[, names(xdata_task2_car
       ) != "clm"]
13
14 # Recoding 'veh_age' and 'agecat' to factors as they represent
       categories
15 xdata_task2_car_processed$veh_age <- as.factor(xdata_task2_car_
       processed$veh_age)
16 xdata_task2_car_processed$agecat <- as.factor(xdata_task2_car_
       processed$agecat)
17 # 'numclaims' is kept as numeric. 'exposure' also kept.
18
19 # Motivation for removals/changes:
20 # - 'clm': After filtering for clm!=0, this variable is constant and
       uninformative for predicting claimcst0 among those who claimed.
21 # - 'veh_age', 'agecat': These are described as age categories (e.g
       ., "youngest"), making factor representation appropriate.
22 # - 'X_OBSTAT_': Likely an ID, not predictive.
23 # - Other variables (e.g., 'exposure', 'numclaims') are kept as
       potential predictors for claim amount.
24 # print(sapply(xdata_task2_car_processed, class))
```

Listing 63: 2b: Data Cleaning Formatting (dataCar) - Base R (Exam 2021)

3.1.8 2c: Cross-validation for linear models

Question *Use a cross validation approach to evaluate predictions for `claimcst0` using: 1) Full LM, 2) All simple LMs, 3) Intercept-only LM. Motivate CV method. Seed 123 (Marking guide used 543, this exam used 123 in task description, so using 123).*

Base R Solution (5-fold CV)

```
1  # Assuming xdata_task2_car_processed from 2b
2  set.seed(123)
3  K_cv <- 5
4  n_car_proc <- nrow(xdata_task2_car_processed)
5  folds_cv_car <- sample(cut(seq(1, n_car_proc), breaks=K_cv, labels=
       FALSE))
6
7  response_car <- "claimcst0"
8  predictors_all_car <- names(xdata_task2_car_processed)[names(xdata_
       task2_car_processed) != response_car]
9
```

```
10  # Matrix to store MSEs
11  mse_results_car <- matrix(NA, nrow=K_cv, ncol=2 + length(predictors_
        all_car))
12  colnames(mse_results_car) <- c("FullLM", "InterceptOnly", predictors
        _all_car)
13
14  for (k_idx_cv in 1:K_cv) {
15    test_idx_cv <- which(folds_cv_car == k_idx_cv)
16    train_df_cv <- xdata_task2_car_processed[-test_idx_cv, ]
17    test_df_cv  <- xdata_task2_car_processed[test_idx_cv, ]
18
19    # Full LM
20    lm_full_cv_car <- lm(paste(response_car, "~ ."), data=train_df_cv)
21    pred_full_cv_car <- predict(lm_full_cv_car, newdata=test_df_cv)
22    mse_results_car[k_idx_cv, "FullLM"] <- mean((test_df_cv[[response_
        car]] - pred_full_cv_car)^2, na.rm=TRUE)
23
24    # Intercept-only LM
25    lm_int_cv_car <- lm(paste(response_car, "~ 1"), data=train_df_cv)
26    pred_int_cv_car <- predict(lm_int_cv_car, newdata=test_df_cv)
27    mse_results_car[k_idx_cv, "InterceptOnly"] <- mean((test_df_cv[[
        response_car]] - pred_int_cv_car)^2, na.rm=TRUE)
28
29    # Simple LMs
30    for (pred_name_car in predictors_all_car) {
31      # Ensure factor levels are consistent or handle potential errors
            if a level is missing in a fold
32      # This might require more careful data handling in a real
            scenario for factors
33      if (is.factor(train_df_cv[[pred_name_car]]) && length(levels(
          droplevels(train_df_cv[[pred_name_car]]))) < 2) next
34
35      formula_simple_cv_car <- as.formula(paste(response_car, "~",
          pred_name_car))
36      tryCatch({ # In case of errors with factors in small folds
37          lm_simple_cv_car <- lm(formula_simple_cv_car, data=train_df_
              cv)
38          pred_simple_cv_car <- predict(lm_simple_cv_car, newdata=test
              _df_cv)
39          mse_results_car[k_idx_cv, pred_name_car] <- mean((test_df_cv
              [[response_car]] - pred_simple_cv_car)^2, na.rm=TRUE)
40      }, error = function(e) { print(paste("Error with", pred_name_car
          , "in fold", k_idx_cv))})
41    }
42  }
43  cv_avg_mses_car <- colMeans(mse_results_car, na.rm=TRUE)
44  print("Average CV Test MSEs for different models:")
45  print(sort(cv_avg_mses_car))
46  # Motivation for K=5: Balances bias (not too small training folds)
        and variance (multiple estimates)
47  # of the test error estimate, and is computationally feasible for
        fitting many models.
```

Listing 64: 2c: CV for Linear Models (dataCar) - Base R (Exam 2021)

3.1.9 2d: Generalized Additive Model (GAM)

Question *Use a GAM with smoothing splines for the numerical variables. Include categorical variables in the model as well. Use* ***s()****-function in* **gam***-library, with*

*argument **df=4**, without motivation. Plot the result and comment on the relationship between predictors and dependent variable.*

Base R Solution (using gam)

```
1  library(gam)
2  # Assuming xdata_task2_car_processed from 2b
3  num_preds_car_gam <- names(xdata_task2_car_processed)[sapply(xdata_
      task2_car_processed, is.numeric) &
4                                          names(xdata_task2_car_
                                              processed) != "
                                              claimcst0"]
5  cat_preds_car_gam <- names(xdata_task2_car_processed)[sapply(xdata_
      task2_car_processed, is.factor)]
6
7  gam_terms_str_car <- c()
8  for (pred in num_preds_car_gam) { gam_terms_str_car <- c(gam_terms_
      str_car, paste0("s(", pred, ", df=4)")) }
9  for (pred in cat_preds_car_gam) { gam_terms_str_car <- c(gam_terms_
      str_car, pred) }
10 gam_formula_full_car_str <- paste("claimcst0 ~", paste(gam_terms_str
      _car, collapse=" + "))
11 gam_formula_full_car <- as.formula(gam_formula_full_car_str)
12
13 # Fit GAM on the processed dataset (xdata_task2_car_processed) for
      plotting
14 gam_model_car_full <- gam(gam_formula_full_car, data=xdata_task2_car
      _processed)
15 # summary(gam_model_car_full)
16
17 # Plot partial effects (adjust mfrow based on number of terms)
18 # num_plots <- length(num_preds_car_gam) + length(cat_preds_car_gam)
19 # par(mfrow=c(ceiling(num_plots/3), 3))
20 # plot(gam_model_car_full, se=TRUE, col="blue", ask=TRUE)
21 # Comment: For each s() term: is it linear, U-shaped, increasing/
      decreasing?
22 # Example from marking guide: numclaims linear, veh_value non-linear
       for small values, exposure non-linear.
```

Listing 65: 2d: GAM with Smoothing Splines (dataCar) - Base R (Exam 2021)

3.1.10 2e: Specify appropriate GAM and compute testMSE

Question *Use your conclusions in 2d to specify an appropriate GAM-model. Compute the testMSE for this GAM-model.*

Base R Solution (Integrating with CV loop from 2c)

```
1  # Based on marking scheme plots/conclusions: s(veh_value), s(
      exposure), linear numclaims, and all factors.
2  # Assuming 'numclaims' is numeric, others like 'veh_body', 'veh_age
      ', 'gender', 'area', 'agecat' are factors.
3  gam_formula_refined_car_str <- "claimcst0 ~ s(veh_value, df=4) + s(
      exposure, df=4) + numclaims + veh_body + veh_age + gender + area
       + agecat"
4  gam_formula_refined_car <- as.formula(gam_formula_refined_car_str)
5
6  # To compute CV testMSE, add this to the CV loop from 2c:
7  # (Inside the loop for k_idx_cv in 1:K_cv)
8  # Ensure factors have same levels in train_df_cv and test_df_cv, or
      handle carefully.
```

```
9  #    tryCatch({
10 #       gam_refined_cv_car <- gam(gam_formula_refined_car, data=train_
     df_cv)
11 #       pred_gam_refined_cv_car <- predict(gam_refined_cv_car, newdata
     =test_df_cv)
12 #       # Add a new column to mse_matrix_car if not already defined:
13 #       # mse_matrix_car[k_idx_cv, "RefinedGAM"] <- mean((test_df_cv[[
     response_car]] - pred_gam_refined_cv_car)^2, na.rm=TRUE)
14 #   }, error = function(e) { print(paste("Error with GAM in fold", k
     _idx_cv, ":", e$message))})
15 # } %% End of loop %
16 # # After loop:
17 # # cv_avg_mses_car <- colMeans(mse_matrix_car, na.rm=TRUE)
18 # # print(cv_avg_mses_car["RefinedGAM"])
19 # # Marking guide: Full GAM had testMSE ~12,243,092. Full linear reg
      ~12,277,335. GAM slightly better.
```

Listing 66: 2e: Refined GAM and Test MSE (dataCar) - Base R (Exam 2021)

Task 3: Analysis of `dataCar` Data (Original) - Predicting `clm`
*Note: For this task, use the original **dataCar**, not the version from Task 2.*

3.1.11 3a: Data prep, descriptive stats for `clm`

Question *Start by removing variables X_OBSTAT_, numclaims and claimcst0. Compute and interpret appropriate cross-tables of clm and the categorical variables. Also compute and interpret well-chosen descriptive statistics (such as central tendencies and measures of variation) for the numerical variables for each of the two categories of clm. Are some of the categories of the categorical variables more interesting than others? Based on this, do you see any promising predictors for clm? Mention a drawback with such pairwise comparisons.*

Base R Solution
```
1  data(dataCar) # Reload original
2  xdata_task3_orig_car <- dataCar[, !(names(dataCar) %in% c("X_OBSTAT_
     ", "numclaims", "claimcst0"))]
3  xdata_task3_orig_car$clm <- as.factor(xdata_task3_orig_car$clm) #
     Target variable
4
5  # Ensure other categoricals are factors
6  xdata_task3_orig_car$veh_body <- as.factor(xdata_task3_orig_car$veh_
     body)
7  xdata_task3_orig_car$veh_age <- as.factor(xdata_task3_orig_car$veh_
     age)
8  xdata_task3_orig_car$gender <- as.factor(xdata_task3_orig_car$gender
     )
9  xdata_task3_orig_car$area <- as.factor(xdata_task3_orig_car$area)
10 xdata_task3_orig_car$agecat <- as.factor(xdata_task3_orig_car$agecat
     )
11
12 # cat_predictors_task3 <- names(xdata_task3_orig_car)[sapply(xdata_
     task3_orig_car, is.factor) & names(xdata_task3_orig_car)!="clm"]
13 # num_predictors_task3 <- names(xdata_task3_orig_car)[sapply(xdata_
     task3_orig_car, is.numeric)]
14
```

```
15  # for (pred in cat_predictors_task3) {
16  #   cat("\nCross-table for clm vs", pred, ":\n")
17  #   tbl <- table(Claim=xdata_task3_orig_car$clm, Predictor=xdata_
    task3_orig_car[[pred]])
18  #   print(tbl)
19  #   print(round(prop.table(tbl, margin = 2)100,1)) # Column
    percentages (P(Claim | Predictor Level))
20  # }
21  # # Interpretation: Look for levels of categorical predictors with
    notably higher/lower claim rates.
22  # # E.g., "BUS and MCARA vehicle types show higher claim rates (19%
    and 11% vs. overall 6.8%)."
23
24  # for (pred_num in num_predictors_task3) {
25  #   cat("\nSummary of", pred_num, "by clm status:\n")
26  #   print(by(xdata_task3_orig_car[[pred_num]], xdata_task3_orig_car$
    clm, summary))
27  #   # boxplot(as.formula(paste(pred_num, "~ clm")), data=xdata_task3
    _orig_car, main=paste(pred_num, "by Claim"))
28  # }
29  # # Interpretation: "Policies with claims have, on average, higher
    veh_value and exposure."
30  # # Promising predictors: veh_body, agecat, veh_value, exposure (
    from marking guide).
31  # # Drawback of pairwise: Does not account for confounding or
    interaction effects between predictors.
```

Listing 67: 3a: Descriptives for clm (dataCar Original) - Base R (Exam 2021)

3.1.12 3b: Logistic regression for `clm`

Question *Based on your hypotheses from 3a, use logistic regression based on all or a subset of the predictors to predict* `clm`*. Interpret the estimated coefficients.*

Base R Solution
```
1   # Assuming xdata_task3_orig_car from 3a.
2   # Predictors from marking guide: veh_body, agecat, veh_value,
    exposure
3   logit_model_clm_car_task3 <- glm(clm ~ veh_body + agecat + veh_value
    + exposure,
4                                    data=xdata_task3_orig_car, family=
                                        binomial)
5   summary(logit_model_clm_car_task3)
6   # Interpretation (odds ratios):
7   # coefs <- coef(logit_model_clm_car_task3)
8   # odds_ratios <- exp(coefs)
9   # print(odds_ratios)
10  # Example: For agecat2, odds of claim are exp(coef_agecat2) (e.g.,
    0.81) times odds for agecat1 (baseline),
11  # holding other variables constant.
12  # For veh_value (continuous), a one-unit increase (10000 dollars)
    multiplies odds of claim by exp(0.06) ~ 1.06.
```

Listing 68: 3b: Logistic Regression for clm (dataCar Original) - Base R (Exam 2021)

3.1.13 3c: Validation set for logistic regression, thresholding

Question *Modify model if you believe it to be sensible. Use the validation set approach with a 50/50 training/test split to evaluate the predictions from the logistic*

regression. Choose a threshold for when an estimated probability should lead to a prediction of a claim (imagine cost FN ¿ FP). Compute a confusion matrix and argue for your choice of threshold. Set the seed to 1234 (Marking guide seed was different, here 1234).

Base R Solution

```r
# Using xdata_task3_orig_car from 3a
set.seed(1234)
n_task3 <- nrow(xdata_task3_orig_car)
train_idx_task3 <- sample(1:n_task3, floor(n_task3/2))
train_data_task3 <- xdata_task3_orig_car[train_idx_task3,]
test_data_task3  <- xdata_task3_orig_car[-train_idx_task3,]

logit_train_task3 <- glm(clm ~ veh_body + agecat + veh_value +
    exposure,
                         data=train_data_task3, family=binomial)
probs_test_task3 <- predict(logit_train_task3, newdata=test_data_
    task3, type="response")

# Threshold choice (e.g., 0.07 as in marking scheme to catch more
    actual claims)
# Motivation: Overall claim rate is low (~6.8%). A standard 0.5
    threshold will predict
# very few claims. To reduce False Negatives (costly if missing an
    actual claim),
# lower the threshold. This increases True Positives but also False
    Positives.
chosen_threshold_task3 <- 0.07
pred_class_task3_logistic <- ifelse(probs_test_task3 > chosen_
    threshold_task3,
                                    levels(train_data_task3$clm)[2],
                                        # "1" or "Yes"
                                    levels(train_data_task3$clm)[1])
                                        # "0" or "No"
pred_class_task3_logistic <- factor(pred_class_task3_logistic,
    levels=levels(train_data_task3$clm))

conf_matrix_task3_logistic <- table(Actual = test_data_task3$clm,
    Predicted = pred_class_task3_logistic)
print(conf_matrix_task3_logistic)
# print(prop.table(conf_matrix_task3_logistic, margin=1)) # Row
    percentages (Sensitivity, Specificity)
# Interpretation: "With a threshold of 0.07, we correctly identify [
    Sensitivity]% of actual claims,
# while incorrectly classifying [1-Specificity]% of non-claims as
    claims. This trade-off was
# chosen to prioritize detecting claims."
```

Listing 69: 3c: Validation Threshold for Logistic (dataCar Original) - Base R (Exam 2021)

3.1.14 3d: Boosted trees for `clm`

Question *Use boosted trees to predict* `clm` *and evaluate the predictions with the same approach as in 3c.*

Base R Solution (using gbm)

```r
library(gbm)
# Using train_data_task3, test_data_task3, chosen_threshold_task3
    from 3c
```

```
3  # GBM requires numeric 0/1 response for bernoulli
4  train_data_task3$clm_numeric <- as.numeric(train_data_task3$clm) - 1
5  test_data_task3$clm_numeric  <- as.numeric(test_data_task3$clm) - 1
6
7  formula_gbm_task3 <- clm_numeric ~ veh_body + agecat + veh_value +
       exposure
8
9  set.seed(1234) # For reproducibility
10 boost_model_task3 <- gbm(formula_gbm_task3,
11                           data=train_data_task3,
12                           distribution="bernoulli",
13                           n.trees=100, # Marking scheme suggests 100
                                  trees
14                           interaction.depth=4, # Example, tune this
15                           shrinkage=0.01)      # Example, tune this
16
17 probs_boost_task3_test <- predict(boost_model_task3, newdata=test_
       data_task3,
18                                  n.trees=100, type="response")
19 pred_class_boost_task3 <- ifelse(probs_boost_task3_test > chosen_
       threshold_task3, 1, 0)
20
21 conf_matrix_boost_task3 <- table(Actual = test_data_task3$clm_
       numeric, Predicted = pred_class_boost_task3)
22 print(conf_matrix_boost_task3)
23 # print(prop.table(conf_matrix_boost_task3, margin=1))
24 # Compare with logistic regression from 3c. Marking scheme: GBM
       improved TP rate (0.67 vs 0.59).
```

Listing 70: 3d: Boosted Trees for clm (dataCar Original) - Base R (Exam 2021)

3.2 Exam Spring 2022

Overall Datasets: `OJ` (ISLR), `Computers` (Ecdat)
Task 1: Analysis of `OJ` Data - Predicting `Purchase`
(Assume `LoyalCH` unavailable except for 1f)

3.2.1 1a: Data preparation and splitting

Question *Remove the variable `LoyalCH`. If necessary, recode categorical variables as factors and remove variables that cannot be used in the analysis. Base your reasoning to do this on the help function and by looking at the data. Split the data in a 50/50 split to create a training and a test dataset. Use the seed 8655, when you split the data, `set.seed(8655)`.*

Base R Solution (basert på løsningsforslag)
```
1 library(ISLR)
2 data(OJ)
3
4 # Remove LoyalCH and StoreID (StoreID is same as STORE after factor
      conversion)
5 oj_data_prep <- OJ[, !(names(OJ) %in% c("LoyalCH", "StoreID"))]
6
7 # Identify variables to convert to factor (<= 5 unique values in
      solution)
8 # card_func <- function(x) length(table(x))
```

```
9  # len_oj <- sapply(oj_data_prep, card_func)
10 # print(len_oj)
11 # Purchase:2, SpecialCH:2, SpecialMM:2, Store7:2, STORE:5 should
       become factors
12 # WeekofPurchase:52, PriceCH:10, PriceMM:8, DiscCH:12, DiscMM:12,
       etc. are numeric
13
14 for (col_name in names(oj_data_prep)) {
15   if (length(unique(oj_data_prep[[col_name]])) <= 5) {
16     oj_data_prep[[col_name]] <- as.factor(oj_data_prep[[col_name]])
17   }
18 }
19 # Verify: sapply(oj_data_prep, class)
20
21 # Split data
22 set.seed(8655) # As specified in this exam version
23 n_oj <- nrow(oj_data_prep)
24 train_idx_oj <- sample(1:n_oj, floor(n_oj/2))
25 train_oj <- oj_data_prep[train_idx_oj, ]
26 test_oj <- oj_data_prep[-train_idx_oj, ]
27
28 # print(paste("Train OJ dimensions:", paste(dim(train_oj), collapse
       ="x")))
29 # print(paste("Test OJ dimensions:", paste(dim(test_oj), collapse="x
       ")))
30 # Motivation: LoyalCH removed as per instruction. StoreID removed as
        STORE factor captures same info.
31 # Variables with few unique values and no inherent order converted
       to factors.
```

Listing 71: 1a: Data Prep Split (OJ) - Base R (Exam 2022)

Tidyverse Solution
```
1  library(ISLR); library(dplyr); library(purrr)
2  data(OJ)
3
4  oj_data_prep_tidy <- OJ %>%
5    select(-LoyalCH, -StoreID) %>%
6    mutate(across(where(~length(unique(.)) <= 5), as.factor))
7
8  set.seed(8655)
9  train_idx_oj_tidy <- sample(1:nrow(oj_data_prep_tidy), floor(nrow(oj
       _data_prep_tidy)/2))
10 train_oj_tidy <- oj_data_prep_tidy[train_idx_oj_tidy, ]
11 test_oj_tidy <- oj_data_prep_tidy[-train_idx_oj_tidy, ]
12 # glimpse(train_oj_tidy)
```

Listing 72: 1a: Data Prep Split (OJ) - Tidyverse (Exam 2022)

3.2.2   1b: Descriptive statistics for promising predictors

Question   *Use descriptive statistics to detect promising predictors for* `Purchase`.
Present the most interesting results as tables and graphs and comment on them.

Base R / Tidyverse Solution (Conceptual - Løsningsforslag fokuserer på plott og 'by')
```
1  # Assuming train_oj from 1a
2  # Visual exploration (Base R)
```

```
3   # par(mfrow=c(1,2)) # Example for specific plots
4   # plot(Purchase ~ STORE, data=train_oj, main="Purchase by Store")
5   # boxplot(PriceDiff ~ Purchase, data=train_oj, main="PriceDiff by
        Purchase")
6   # graphics.off()
7
8   # Numerical summaries (Base R)
9   # cat("Mean PctDiscCH by Purchase:\n")
10  # print(by(train_oj$PctDiscCH, train_oj$Purchase, mean, na.rm=TRUE))
11  # cat("Mean PctDiscMM by Purchase:\n")
12  # print(by(train_oj$PctDiscMM, train_oj$Purchase, mean, na.rm=TRUE))
13  # cat("Table of SpecialMM by Purchase:\n")
14  # print(table(train_oj$SpecialMM, train_oj$Purchase))
15
16  # Tidyverse for summaries
17  # library(dplyr)
18  # train_oj %>% group_by(Purchase) %>%
19  #   summarise(across(c(PriceDiff, PctDiscCH, PctDiscMM),
20  #                    list(mean=mean, median=median), na.rm=TRUE))
21  # train_oj %>% count(STORE, Purchase) %>% group_by(STORE) %>% mutate
        (prop = n/sum(n))
22
23  # Comments (basert p    l sningsforslag):
24  # - STORE: Variation in brand purchased by store.
25  # - PriceDiff: Higher (MM more expensive) when CH bought.
26  # - PctDiscCH/PctDiscMM: Discounts seem to influence choice.
27  # - SpecialMM: Appears differently distributed between CH/MM
        purchases.
```

Listing 73: 1b: Descriptive Stats for Purchase (OJ) - Conceptual (Exam 2022)

3.2.3   1c: Logistic regression

Question   *Fit a logistic regression of* `Purchase` *on all the variables you found promising in 1b, evaluate the predictions using test accuracy and a confusion matrix. If you, in the process, detect that you would modify your model, you should do so and explain why. Use the threshold 0.5 when you go from estimated probability to a prediction.*

Base R Solution (Løsningsforslag bruker `PriceDiff + STORE`)

```
1   # Assuming train_oj, test_oj from 1a
2   # L sningsforslag ender opp med PriceDiff og STORE etter    ha
        vurdert PctDisc-variablene som ikke-signifikante.
3   # Vi starter med de 'lovende' og fjerner evt. ikke-signifikante.
4   # Initial promising: STORE, PriceDiff, PctDiscCH, PctDiscMM,
        SpecialMM
5   # logit_oj_initial <- glm(Purchase ~ STORE + PriceDiff + PctDiscCH +
         PctDiscMM + SpecialMM,
6   #                         data=train_oj, family=binomial)
7   # summary(logit_oj_initial)
8   # % If PctDiscCH, PctDiscMM, SpecialMM are not significant, remove
        them. %
9
10  # Final model from solution:
11  logit_oj_final <- glm(Purchase ~ PriceDiff + STORE, data=train_oj,
        family=binomial)
12  # print(summary(logit_oj_final))
13
```

```
14  # Predictions on test set
15  probs_logit_oj_test <- predict(logit_oj_final, newdata=test_oj, type
        ="response")
16  pred_logit_oj_test <- ifelse(probs_logit_oj_test > 0.5, "MM", "CH")
        # OJ levels: CH, MM
17  pred_logit_oj_test <- factor(pred_logit_oj_test, levels = levels(
        test_oj$Purchase))
18
19  # Confusion matrix and accuracy
20  conf_matrix_logit_oj <- table(Actual=test_oj$Purchase, Predicted=
        pred_logit_oj_test)
21  print(conf_matrix_logit_oj)
22  accuracy_logit_oj <- sum(diag(conf_matrix_logit_oj)) / sum(conf_
        matrix_logit_oj)
23  cat("Test Accuracy (Logistic Regression):", accuracy_logit_oj, "\n")
24  # prop_table_logit_oj <- prop.table(conf_matrix_logit_oj, margin=1)
25  # print(round(prop_table_logit_oj,3))
26  # Interpretation: The model predicts [value]% of CH purchases
        correctly and [value]% of MM purchases.
```

Listing 74: 1c: Logistic Regression (OJ) - Base R (Exam 2022)

3.2.4   1d: Classification tree

Question   *Fit a classification tree, using all variables (except the one(s) you removed in 1a), and plot it. Give an example of how to interpret a branch of the tree. Also, use the same evaluation measures as in 1c and evaluate the predictions. Use the threshold 0.5.*

Base R Solution (using `tree`**)**

```
1   library(tree)
2   # Assuming train_oj, test_oj from 1a (with all original predictors
        kept, except LoyalCH, StoreID)
3   tree_oj <- tree(Purchase ~ ., data=train_oj)
4   # summary(tree_oj)
5   # plot(tree_oj); text(tree_oj, pretty=0)
6
7   # Interpretation example (basert p   l sningsforslagets trefigur):
8   # "If STORE is 'ae' (0 or 4 in factor levels) AND PriceDiff < -0.35
        AND WeekofPurchase < 258.5,
9   # then the prediction is MM." (Values vil variere basert p   seed).
10
11  # Predictions on test set
12  # predict() for tree with type="class" gives direct class
        predictions.
13  # For 0.5 threshold on probabilities, get probabilities first:
14  probs_tree_oj_test <- predict(tree_oj, newdata=test_oj, type="vector
        ") # Gives matrix of probs
15  # Assuming Purchase is factor with levels CH, MM. probs_tree_oj_test
        [,2] is P(MM)
16  pred_tree_oj_test <- ifelse(probs_tree_oj_test[,2] > 0.5, "MM", "CH"
        )
17  pred_tree_oj_test <- factor(pred_tree_oj_test, levels = levels(test_
        oj$Purchase))
18
19  conf_matrix_tree_oj <- table(Actual=test_oj$Purchase, Predicted=pred
        _tree_oj_test)
20  print(conf_matrix_tree_oj)
```

```
21  accuracy_tree_oj <- sum(diag(conf_matrix_tree_oj)) / sum(conf_matrix
        _tree_oj)
22  cat("Test Accuracy (Classification Tree):", accuracy_tree_oj, "\n")
23  # prop_table_tree_oj <- prop.table(conf_matrix_tree_oj, margin=1)
24  # print(round(prop_table_tree_oj,3))
```

Listing 75: 1d: Classification Tree (OJ) - Base R (Exam 2022)

3.2.5   1e: Cross-validation for optimal threshold for tree

Question   *In 1c and 1d you have used the threshold 0.5. Use cross-validation to find the threshold that maximizes the accuracy for the tree. Evaluate predictions from the tree using the best threshold. Are the predictions improved? Use the seed 4598 when you split the data for CV (this implies an inner CV split on the original training data).*

Base R Solution (Validation set for threshold tuning on original training data)

```
1   # Using original train_oj and test_oj from 1a.
2   # To find optimal threshold, we need a validation set, split from
        original train_oj.
3   set.seed(4598) # Seed for splitting train_oj further
4   n_train_oj <- nrow(train_oj)
5   val_idx_oj <- sample(1:n_train_oj, floor(n_train_oj/2))
6   train2_oj <- train_oj[val_idx_oj, ]
7   valid_oj <- train_oj[-val_idx_oj, ]
8
9   # Fit tree on train2_oj
10  tree_for_thresh <- tree(Purchase ~ ., data=train2_oj)
11  probs_valid_oj <- predict(tree_for_thresh, newdata=valid_oj, type="
        vector")[,2] # P(MM)
12
13  thresholds <- seq(0.1, 0.9, by=0.05)
14  accuracies_thresh <- numeric(length(thresholds))
15
16  for (i in 1:length(thresholds)) {
17    th <- thresholds[i]
18    pred_class_thresh <- ifelse(probs_valid_oj > th, "MM", "CH")
19    pred_class_thresh <- factor(pred_class_thresh, levels = levels(
        valid_oj$Purchase))
20    accuracies_thresh[i] <- mean(pred_class_thresh == valid_oj$
        Purchase)
21  }
22
23  # plot(thresholds, accuracies_thresh, type="b", xlab="Threshold",
        ylab="Validation Accuracy")
24  best_threshold_idx <- which.max(accuracies_thresh)
25  optimal_threshold_oj <- thresholds[best_threshold_idx]
26  cat("Optimal threshold found:", optimal_threshold_oj, "\n")
27
28  # Evaluate original tree_oj (fit on full train_oj) on test_oj using
        this optimal_threshold_oj
29  probs_tree_fulltrain_test <- predict(tree_oj, newdata=test_oj, type=
        "vector")[,2] # P(MM)
30  pred_tree_optimal_thresh <- ifelse(probs_tree_fulltrain_test >
        optimal_threshold_oj, "MM", "CH")
31  pred_tree_optimal_thresh <- factor(pred_tree_optimal_thresh, levels=
        levels(test_oj$Purchase))
```

```
32
33  conf_matrix_tree_optimal <- table(Actual=test_oj$Purchase, Predicted
        =pred_tree_optimal_thresh)
34  print(conf_matrix_tree_optimal)
35  accuracy_tree_optimal <- sum(diag(conf_matrix_tree_optimal)) / sum(
        conf_matrix_tree_optimal)
36  cat("Test Accuracy (Tree with Optimal Threshold):", accuracy_tree_
        optimal, "\n")
37  # Compare accuracy_tree_optimal with accuracy_tree_oj (from 1d using
         0.5 threshold)
38  # Marking guide's plot actually uses the same tree (tree1) for
        validation, which is fine.
39  # Their plot shows threshold around 0.4-0.6 being optimal. Accuracy
        doesn't improve much.
```

Listing 76: 1e: CV for Optimal Threshold (OJ Tree) - Base R (Exam 2022)

3.2.6 1f: Analysis with `LoyalCH`

Question   *In a) you removed the variable `LoyalCH`. Imagine now that this variable is actually available at the time of prediction. By augmenting your analysis with this variable both for the logistic regression and the classification tree, conclude if this variable contribute to predicting `Purchase`.*

Solution Outline   *Repeat steps 1a-1d, but KEEP LoyalCH in the dataset. Compare test accuracies and model summaries.*

```
1   # 1. Data Prep: Start with original OJ, remove only StoreID. Convert
        factors. Split.
2   #     oj_with_loyalch <- OJ[, names(OJ) != "StoreID"]
3   #     ... convert factors, split into train_oj_L, test_oj_L ...
4
5   # 2. Logistic Regression with LoyalCH:
6   #     logit_oj_L <- glm(Purchase ~ PriceDiff + STORE + LoyalCH, data=
        train_oj_L, family=binomial)
7   #     % (Or include all promising predictors + LoyalCH) %
8   #     summary(logit_oj_L) % Check significance of LoyalCH %
9   #     % Calculate test accuracy and confusion matrix on test_oj_L %
10  #     % Compare accuracy with result from 1c. %
11
12  # 3. Classification Tree with LoyalCH:
13  #     tree_oj_L <- tree(Purchase ~ ., data=train_oj_L)
14  #     summary(tree_oj_L) % See if LoyalCH is used in splits %
15  #     plot(tree_oj_L); text(tree_oj_L, pretty=0)
16  #     % Calculate test accuracy and confusion matrix on test_oj_L %
17  #     % Compare accuracy with result from 1d. %
18
19  # Conclusion: If LoyalCH is significant in logistic regression and/
        or used prominently
20  # in the tree, and if test accuracies improve notably, then LoyalCH
        contributes.
21  # (LoyalCH is expected to be a very strong predictor).
```

Listing 77: 1f: Re-analysis with LoyalCH - Conceptual (Exam 2022)

Task 2: Analysis of `Computers` Data - Predicting `price`

3.2.7 2a: Data preparation and splitting

Question *Just as in 1a; if necessary, recode categorical variables as factors and remove variables that cannot be used in the analysis. Base your reasoning to do this on the help function and by looking at the data. Split the data in a 50/50 split to create a training and a test dataset. Use the seed 4598, when you split the data,* `set.seed(4598)`*.*

Base R Solution (basert på løsningsforslag)

```r
library(Ecdat)
data(Computers)

# Recode categoricals (<= 5 unique values, from marking scheme)
# Inspect: sapply(Computers, function(x) length(unique(x)))
# screen (3), cd (2), multi (2), premium (2) should become factors
cols_to_factor_comp <- c("screen", "cd", "multi", "premium")
for (col in cols_to_factor_comp) {
  Computers[[col]] <- as.factor(Computers[[col]])
}
# sapply(Computers, class) # Verify

# Split data
set.seed(4598)
n_comp <- nrow(Computers)
train_idx_comp <- sample(1:n_comp, floor(n_comp/2))
train_comp <- Computers[train_idx_comp, ]
test_comp <- Computers[-train_idx_comp, ]
# print(paste("Train Computers dimensions:", paste(dim(train_comp),
    collapse="x")))
```

Listing 78: 2a: Data Prep Split (Computers) - Base R (Exam 2022)

Tidyverse Solution

```r
library(Ecdat); library(dplyr)
data(Computers)
Computers_prep_tidy <- Computers %>%
  mutate(across(all_of(c("screen", "cd", "multi", "premium")), as.
      factor))

set.seed(4598)
train_idx_comp_tidy <- sample(1:nrow(Computers_prep_tidy), floor(
    nrow(Computers_prep_tidy)/2))
train_comp_tidy <- Computers_prep_tidy[train_idx_comp_tidy, ]
test_comp_tidy <- Computers_prep_tidy[-train_idx_comp_tidy, ]
# glimpse(train_comp_tidy)
```

Listing 79: 2a: Data Prep Split (Computers) - Tidyverse (Exam 2022)

3.2.8 2b: Descriptive statistics for promising predictors

Question *Use descriptive statistics to find promising predictors. Comment on your observations.*

Base R / Tidyverse Solution (Conceptual - Løsningsforslag bruker `plot(price .,data=train)`**)**

```r
# Assuming train_comp from 2a
# Visual exploration (Base R)
```

```
3  # par(mfrow=c(3,3)) # Adjust layout as needed
4  # plot(price ~ ., data=train_comp) # As per marking guide
5  # graphics.off()
6  # Comment: "Speed, hd, ram show positive trends with price. Screen17
       has higher price than screen15.
7  # cd=yes, multi=yes, premium=yes associated with higher price. Ads
       and trend also show relationships."
8
9  # Tidyverse for specific plots
10 # library(ggplot2)
11 # ggplot(train_comp, aes(x=speed, y=price)) + geom_point() + geom_
       smooth()
12 # ggplot(train_comp, aes(x=ram, y=price)) + geom_point() + geom_
       smooth()
13 # ggplot(train_comp, aes(x=screen, y=price)) + geom_boxplot()
```

Listing 80: 2b: Descriptive Stats for Price (Computers) - Conceptual (Exam 2022)

3.2.9 2c: Linear regression for `price`

Question *Run a linear regression on all explanatory variables. Interpret some of the coefficients. Evaluate the prediction performance on an appropriate measure.*

Base R Solution
```
1  # Assuming train_comp, test_comp from 2a
2  lm_comp_full <- lm(price ~ ., data=train_comp)
3  summary_lm_comp <- summary(lm_comp_full)
4  print(summary_lm_comp)
5  # Interpretation: e.g., "A 1MHz increase in speed is associated with
       an expected price
6  # increase of [coef_speed] dollars, holding other factors constant (
       p=[p_val])."
7  # "Having a CD-ROM (cdyes) is associated with an expected price
       increase of
8  # [coef_cdyes] compared to no CD-ROM, holding others constant (p=[p_
       val])."
9
10 # Evaluate performance: Test MSE
11 pred_lm_comp_test <- predict(lm_comp_full, newdata=test_comp)
12 mse_lm_comp_test <- mean((test_comp$price - pred_lm_comp_test)^2, na
       .rm=TRUE)
13 cat("Test MSE (Full LM for Computers):", mse_lm_comp_test, "\n")
14 # Marking guide result: MSE1 = 72200
```

Listing 81: 2c: Linear Regression (Computers) - Base R (Exam 2022)

3.2.10 2d: Linear regression of `log(price)`

Question *Investigate if the predictions of `price` would be better if we use a linear regression of `log(price)` on the other variables.*

Base R Solution
```
1  # Create log(price) variable, ensuring it's available in train and
       test
2  # train_comp_log <- train_comp
3  # test_comp_log <- test_comp
4  # train_comp_log$log_price <- log(train_comp_log$price)
```

```
5  # test_comp_log$log_price <- log(test_comp_log$price) # For actual
       values if needed, but predict() on log scale
6
7  # Fit model on training data using log_price, excluding original
       price
8  # lm_comp_log <- lm(log(price) ~ . - price, data=train_comp) #
       Simpler way
9  lm_comp_log <- lm(log(price) ~ speed + hd + ram + screen + cd +
       multi + premium + ads + trend,
10                   data=train_comp)
11 summary(lm_comp_log)
12
13 # Predict on test set (predictions will be on log scale)
14 pred_log_price_test <- predict(lm_comp_log, newdata=test_comp)
15 # Transform predictions back to original price scale
16 pred_price_from_log_test <- exp(pred_log_price_test)
17
18 # Calculate Test MSE on original price scale
19 mse_loglm_comp_test <- mean((test_comp$price - pred_price_from_log_
       test)^2, na.rm=TRUE)
20 cat("Test MSE (Log-Linear LM for Computers):", mse_loglm_comp_test,
       "\n")
21 # Marking guide result: MSE2 = 72610.7. In this case, log transform
       didn't improve MSE.
```

Listing 82: 2d: Log-Linear Regression (Computers) - Base R (Exam 2022)

3.2.11   2e: GAM for `price`

Question   *In 2d you fitted a model where* `price` *is a particular nonlinear function of the other variables. You should now investigate another non-linear models. First, fit a GAM-model, plot the result and evaluate the predictions. Is there a reason to not allow for some variables to have a non-linear relationship with* `price`?

Base R Solution (using gam)

```
1  library(gam)
2  # Assuming train_comp, test_comp from 2a
3  # Formula for GAM: s() for continuous, direct for factors
4  # From marking guide: s(speed)+s(hd)+s(ram)+screen+cd+multi+premium+
       s(ads)+s(trend)
5  gam_comp1 <- gam(price ~ s(speed, df=4) + s(hd, df=4) + s(ram, df=4)
       + screen + cd + multi +
6                     premium + s(ads, df=4) + s(trend, df=4), data=train
                       _comp)
7  # summary(gam_comp1)
8
9  # Plot partial effects
10 # par(mfrow=c(3,3)) # To fit all plots, adjust as needed
11 # plot(gam_comp1, se=TRUE, col="blue", ask=TRUE)
12 # Reason to allow non-linear: Some relationships (e.g. price vs hd
       or speed) might
13 # exhibit diminishing returns or other non-linear patterns. Plots
       will confirm.
14
15 # Evaluate Test MSE
16 pred_gam_comp_test <- predict(gam_comp1, newdata=test_comp)
17 mse_gam_comp_test <- mean((test_comp$price - pred_gam_comp_test)^2,
       na.rm=TRUE)
18 cat("Test MSE (GAM for Computers):", mse_gam_comp_test, "\n")
```

```
19  # Marking guide result: MSE3 = 60947.91. GAM significantly improved
       predictions.
```

Listing 83: 2e: GAM for Price (Computers) - Base R (Exam 2022)

3.2.12   2f: Explanation of backfitting for GAMs

Question   *Give a brief explanation of how backfitting works and in what situations it is possible to fit a generalized additive model with ordinary least square regression.*

Explanation   ***Backfitting Algorithm***: *For a GAM model* $y_i = \beta_0 + \sum_{j=1}^{p} f_j(x_{ij}) + \epsilon_i$:

1. *Initialize:* $\hat{\beta}_0 = \bar{y}$, *and* $\hat{f}_j(x_{ij}) = 0$ *for all* j *(or some initial guess, e.g., from linear model).*

2. *Cycle through predictors* $j = 1, \ldots, p, 1, \ldots, p, \ldots$ *until convergence:*

   * *For current predictor* j, *compute partial residuals:* $r_{ij} = y_i - \hat{\beta}_0 - \sum_{k \neq j} \hat{f}_k(x_{ik})$.

   * *Update* $\hat{f}_j$ *by fitting a smoother (e.g., spline, loess) to the partial residuals* r_{ij} *against* x_{ij}. *That is,* $\hat{f}_j \leftarrow smoother(r_{.j} \sim x_{.j})$.

   * *Often,* $\hat{f}_j$ *are centered to ensure identifiability.*

3. *Convergence is reached when the functions* $\hat{f}_j$ *do not change much between iterations.*

GAM with OLS: *A GAM can be fitted with OLS if each function* $f_j(x_j)$ *can be represented by a linear combination of a set of basis functions. For example:*

   * ***Polynomial Regression***: $f_j(x_j) = \beta_{j1}x_j + \beta_{j2}x_j^2 + \ldots$. *The "basis functions" are* $x_j, x_j^2, \ldots$.

   * ***Regression Splines***: $f_j(x_j)$ *is represented by a sum of spline basis functions (e.g., truncated power basis, B-spline basis). The model* $y = \beta_0 + \sum_j \sum_l \beta_{jl} b_{jl}(x_j) + \epsilon$ *is then linear in the* β_{jl} *coefficients and can be fit with OLS.*

If non-parametric smoothers like smoothing splines (where λ *is chosen by GCV) or local regression are used within backfitting, the overall procedure is not a single OLS fit, but each step of backfitting might involve an OLS-like weighted least squares or penalized least squares.*

3.2.13   2g: Bagged trees for `price`

Question   *Use bagged trees to predict **price**. Evaluate the predictions. Compute variable importance measures and comment on the results. Explain the logic behind the measure for variable importance that you are using. (Hint: If the computation takes a long time you can reduce the number of trees fitted using the argument **ntree**.)*

Base R Solution (using `randomForest`)

```
1  library(randomForest)
2  # Assuming train_comp, test_comp from 2a
3  # For bagging, mtry = number of predictors.
4  # price ~ . means all other columns are predictors.
5  # ncol(train_comp) - 1 (if price is a column, and all others are
       predictors)
```

```
6  # If train_comp has 10 predictors plus 'price' and 'logp' (from 2d),
       p=10.
7  # Marking guide uses mtry=9, perhaps 'logp' was still in data.
       Assuming p=10 here.
8  num_predictors_comp <- ncol(train_comp) - 1 # Assuming price is the
       only response
9  if ("logp" %in% names(train_comp)) num_predictors_comp <- num_
       predictors_comp -1
10
11 set.seed(4598) # For reproducibility
12 bag_comp <- randomForest(price ~ . - logp, data=train_comp, #
       Exclude logp if it exists
13                            mtry=num_predictors_comp,
14                            ntree=100, # Reduced from default 500 for
                                 speed as hinted
15                            importance=TRUE)
16 # print(bag_comp)
17
18 # Evaluate predictions
19 pred_bag_comp_test <- predict(bag_comp, newdata=test_comp)
20 mse_bag_comp_test <- mean((test_comp$price - pred_bag_comp_test)^2,
       na.rm=TRUE)
21 cat("Test MSE (Bagging for Computers):", mse_bag_comp_test, "\n")
22 # Marking guide result: MSE4 = 28272.51 (significant improvement)
23
24 # Variable Importance
25 importance_bag_comp <- importance(bag_comp)
26 print(importance_bag_comp)
27 varImpPlot(bag_comp)
28 # Logic for IncNodePurity (default for regression):
29 # "IncNodePurity (or %IncMSE if type=1 is used) measures the total
       decrease in node impurity
30 # (RSS for regression trees) from splitting on that variable,
       averaged over all trees.
31 # A higher value indicates greater importance."
32 # Comment: "Ram, trend, speed, hd appear as most important variables
       ."
```

Listing 84: 2g: Bagging for Price (Computers) - Base R (Exam 2022)

3.2.14   2h: Explanation of bagging

Question   *Give a brief explanation of bagging.*

Explanation   *Bagging (Bootstrap Aggregating): Bagging is an ensemble learning technique designed to improve the stability and accuracy of machine learning algorithms, particularly those with high variance like decision trees. The process involves:*

1. ***Bootstrap Sampling****: Create B new training datasets by sampling with replacement from the original training dataset. Each bootstrap sample has the same size as the original.*

2. ***Model Fitting****: Fit a separate model (e.g., a decision tree) independently to each of the B bootstrap samples. These trees are typically grown deep and not pruned.*

3. ***Aggregation****:*

   ∗ For **regression**, the predictions from the B models are averaged to get the final prediction.

   ∗ For **classification**, a majority vote is taken among the B model predictions.

The main benefit of bagging is variance reduction. By averaging many (potentially noisy) models fit to slightly different versions of the data, the overall variance of the ensemble prediction is reduced, often leading to improved predictive performance.

3.3 Exam Spring 2023

Overall Dataset: `Churn.csv` (Provided on Canvas) *Refer to variable definitions table at the end of this exam section.*

Task 1: Predict Average Bill (`bill_avg`)

3.3.1 1a: Data preparation and splitting

Question   *If necessary, recode categorical variables as factors and remove variables that cannot be used in the analysis. Also, remove variables if you mean that they are unreasonable to include. One such reason could be that a variable is not available at the time the prediction is made but there might be other reasons too. Motivate your choices in words. Split the data in a 50/50 split to create a training and a test dataset. Use the seed 86554354, when you split the data, set.seed(86554354).*

Base R Solution (basert på løsningsforslag)

```
1  # Churn <- read.csv("Churn.csv") % \textit{(Assuming data is read)}
     %
2
3  # Remove id
4  xdata_churn <- Churn[, -1] % \textit{Assuming 'id' is the first
     column} %
5
6  # Recode specified variables to factors
7  factor_cols_churn <- c("is_tv_subscriber", "is_movie_package_
     subscriber", "churn")
8  for (col in factor_cols_churn) {
9    xdata_churn[[col]] <- as.factor(xdata_churn[[col]])
10 }
11 # print(sapply(xdata_churn, class))
12
13 # Split data
14 set.seed(86554354) # Seed from exam question (v1 from marking
     instructions might differ)
15 n_churn <- nrow(xdata_churn)
16 train_idx_churn <- sample(1:n_churn, floor(n_churn/2))
17 train_churn <- xdata_churn[train_idx_churn, ]
18 test_churn <- xdata_churn[-train_idx_churn, ]
19
20 # print(paste("Train Churn dimensions:", paste(dim(train_churn),
     collapse="x")))
21 # print(paste("Test Churn dimensions:", paste(dim(test_churn),
     collapse="x")))
22 # Motivation: 'id' is an identifier and not predictive. 'is_tv_
     subscriber',
```

```
23  # 'is_movie_package_subscriber', and 'churn' are categorical (0/1
        originally)
24  # and explicitly converted to factors for proper handling by
        modeling functions.
25  # The question of whether 'churn' is available when predicting 'bill
        _avg' is raised
26  # in the marking scheme; it's assumed available for this task.
```

Listing 85: 1a: Data Prep Split (Churn) - Base R (Exam 2023)

Tidyverse Solution
```
1   # library(dplyr)
2   # Churn <- read.csv("Churn.csv")
3   #
4   # xdata_churn_tidy <- Churn %>%
5   #   select(-id) %>%
6   #   mutate(across(all_of(c("is_tv_subscriber", "is_movie_package_
        subscriber", "churn")), as.factor))
7   #
8   # set.seed(86554354)
9   # train_idx_churn_tidy <- sample(1:nrow(xdata_churn_tidy), floor(
        nrow(xdata_churn_tidy)/2))
10  # train_churn_tidy <- xdata_churn_tidy[train_idx_churn_tidy, ]
11  # test_churn_tidy <- xdata_churn_tidy[-train_idx_churn_tidy, ]
12  # glimpse(train_churn_tidy)
```

Listing 86: 1a: Data Prep Split (Churn) - Tidyverse (Exam 2023)

3.3.2   1b: Descriptive methods for `bill_avg`

Question   *Use descriptive methods to find useful predictors for* `bill_avg`. *Write in words which R functions you used, present the most interesting results as tables and graphs and comment on them.*

Base R / Tidyverse Solution (Conceptual - Marking guide uses
   `plot(bill_avg ., data=train))`
```
1   # Assuming train_churn from 1a
2   # Visual exploration (Base R as per marking guide)
3   # par(mfrow=c(3,3)) # To fit all plots (9 predictors excluding id
        and bill_avg itself)
4   # plot(bill_avg ~ ., data=train_churn)
5   # graphics.off()
6   # Comments (based on marking guide comment):
7   # "bill_avg seems to vary with all predictors, possibly with the
        exception of
8   # download_avg and upload_avg, where patterns are difficult to see
        visually."
9   # For factors like 'is_tv_subscriber', 'churn', boxplots are
        generated by plot(y~.).
10  # For numeric like 'subscription_age', scatterplots are generated.
11
12  # Tidyverse for specific plots (examples)
13  # library(ggplot2)
14  # ggplot(train_churn, aes(x=is_tv_subscriber, y=bill_avg)) + geom_
        boxplot()
15  # ggplot(train_churn, aes(x=subscription_age, y=bill_avg)) + geom_
        point() + geom_smooth()
```

Listing 87: 1b: Descriptive Stats for $bill_avg(Churn) - Conceptual(Exam 2023)$

3.3.3 1c: Best OLS for `bill_avg`

Question *Produce the best possible predictions of **bill_avg** using standard linear regression (OLS) and evaluate them. Motivate your choices of variables, how you evaluate the predictions and the evaluation measure used.*

Base R Solution (Comparing full model vs. model without download/upload_avg)

```r
# Assuming train_churn, test_churn from 1a
# Model 1: Full model with all predictors
ols_full_churn <- lm(bill_avg ~ ., data=train_churn)
# summary(ols_full_churn)
pred_ols_full_churn <- predict(ols_full_churn, newdata=test_churn)
mse_ols_full_churn <- mean((test_churn$bill_avg - pred_ols_full_
    churn)^2, na.rm=TRUE)
cat("Test MSE (Full OLS for bill_avg):", mse_ols_full_churn, "\n")
# Marking guide value: 117.8603

# Model 2: Reduced model (based on 1b, removing download_avg, upload
    _avg)
ols_reduced_churn <- lm(bill_avg ~ . - download_avg - upload_avg,
    data=train_churn)
# summary(ols_reduced_churn)
pred_ols_reduced_churn <- predict(ols_reduced_churn, newdata=test_
    churn)
mse_ols_reduced_churn <- mean((test_churn$bill_avg - pred_ols_
    reduced_churn)^2, na.rm=TRUE)
cat("Test MSE (Reduced OLS for bill_avg):", mse_ols_reduced_churn, "
    \n")
# Marking guide value: 146.7376

# Motivation:
# Variables: Compared a full model with one excluding download_avg
    and upload_avg, based on visual
# inspection in 1b suggesting weak relationships.
# Evaluation Measure: Test MSE is used to assess predictive accuracy
     on unseen data.
# A lower Test MSE indicates better out-of-sample prediction.
# Conclusion: The full model performed much better (lower Test MSE)
    than the reduced model.
```

Listing 88: 1c: OLS for $\text{bill}_a vg(Churn) - BaseR(Exam 2023)$

3.3.4 1d: LASSO regression for `bill_avg`

Question *Fit a LASSO regression with all variables. Here you should use the tools you have learned to find appropriate tuning parameters. Are you standardizing the predictors or not? Motivate your choice. Compare the estimated coefficients with the estimated coefficient from an OLS regression on all predictors. (Hint: Using categorical (factor) variables in LASSO, you will have to create dummy variables. **dummy_cols()** from **fastDummies** or **model.matrix()** can be used. **glmnet()** requires matrices as input).*

Base R Solution (using `glmnet`)

```r
library(glmnet)
# Assuming train_churn, test_churn from 1a. Factors are already
    created.
```

```r
 3 | # For glmnet , factors need to be converted to dummy variables.
 4 | # We use the processed data before splitting for model.matrix, then
   |     split X and y.
 5 | # xdata_churn was defined in 1a with factors.
 6 |
 7 | # Create model matrix for predictors (converts factors to dummies)
 8 | x_matrix_churn <- model.matrix(bill_avg ~ . -1, data=xdata_churn) #
   |     -1 removes intercept if glmnet adds it
 9 | y_vector_churn <- xdata_churn$bill_avg
10 |
11 | # Split X and y matrices into train/test using train_idx_churn from
   |     1a
12 | trainX_churn <- x_matrix_churn[train_idx_churn, ]
13 | trainY_churn <- y_vector_churn[train_idx_churn]
14 | testX_churn <- x_matrix_churn[-train_idx_churn, ]
15 | testY_churn <- y_vector_churn[-train_idx_churn] % \textit{(Actually
   |     test_churn$bill_avg)} %
16 |
17 | # Determine lambda by cross-validation (alpha=1 for LASSO)
18 | set.seed(86554354) % \textit{Ensure consistent seed if CV is
   |     stochastic} %
19 | cv_lasso_churn <- cv.glmnet(trainX_churn, trainY_churn, alpha=1,
   |     standardize=TRUE)
20 | # standardize=TRUE is default and recommended for LASSO/Ridge when
   |     variables are on different scales.
21 | # Motivation for standardization: LASSO penalty applies equally to
   |     all coefficients. If predictors
22 | # have different scales, those with larger scales might be unfairly
   |     penalized or vice versa.
23 | # Standardization puts all predictors on a common scale.
24 |
25 | best_lambda_lasso_churn <- cv_lasso_churn$lambda.min
26 | # print(paste("Best lambda (min):", best_lambda_lasso_churn))
27 |
28 | # Fit LASSO with optimal lambda on training data
29 | lasso_model_churn <- glmnet(trainX_churn, trainY_churn, alpha=1,
   |     lambda=best_lambda_lasso_churn, standardize=TRUE)
30 | lasso_coeffs <- coef(lasso_model_churn)
31 | # print("LASSO Coefficients:")
32 | # print(lasso_coeffs)
33 |
34 | # OLS model from 1c (fitted on train_churn, which has factors not
   |     dummies)
35 | # ols_full_churn <- lm(bill_avg ~ ., data=train_churn)
36 | # ols_coeffs <- coef(ols_full_churn)
37 | # print("OLS Coefficients (factors handled by lm):")
38 | # print(ols_coeffs)
39 |
40 | # Comparison:
41 | # Lasso coefficients are shrunk towards zero. Some might be exactly
   |     zero.
42 | # OLS coefficients (for the dummy variables created from factors)
   |     will generally be larger in magnitude.
43 | # Marking scheme showed factor variables were converted to 0/1
   |     numeric for glmnet,
44 | # and coefficients were very similar, none set to zero, indicating
   |     LASSO restriction not very large.
45 | # This implies xdata in marking scheme was xdata_num after
   |     converting factors to numeric.
```

```
46  # If xdata_num <- xdata; for(i in 1:length(fa)) xdata_num[,fa[i]] <-
        as.numeric(xdata_num[,fa[i]])-1
47  # Then use xdata_num for trainX, trainy.
48  # cbind(ols_full_churn$coefficients, lasso_coeffs) % \textit{(Align
        names carefully for direct comparison)}%
```

Listing 89: 1d: LASSO for bill$_a vg(Churn) - BaseR(Exam2023)$

3.3.5 1e: Evaluate LASSO model

Question *Compute predictions for* `bill_avg` *with the model fitted in 1d) and evaluate them with an appropriate measure.*

Base R Solution
```
1  # Assuming lasso_model_churn (fitted on trainX_churn) and testX_
        churn from 1d
2  pred_lasso_churn_test <- predict(lasso_model_churn, newx=testX_churn
        )
3  mse_lasso_churn_test <- mean((test_churn$bill_avg - pred_lasso_churn
        _test)^2, na.rm=TRUE)
4  cat("Test MSE (LASSO for bill_avg):", mse_lasso_churn_test, "\n")
5  # Marking guide value: 117.8655 (very similar to OLS in this case)
```

Listing 90: 1e: Evaluate LASSO (Churn) - Base R (Exam 2023)

3.3.6 1f: Regression tree for `bill_avg`

Question *Fit a regression tree to* `bill_avg`*. Explain the choices you are making. Interpret the tree.*

Base R Solution (using `tree`)
```
1   library(tree)
2   # Assuming train_churn from 1a (with factors)
3   # Choices: Using all predictors. Default tree.control parameters
        initially.
4   tree_bill_avg <- tree(bill_avg ~ ., data=train_churn)
5   # summary(tree_bill_avg)
6   # plot(tree_bill_avg)
7   # text(tree_bill_avg, pretty=0)
8
9   # Interpretation (based on marking scheme tree plot which is complex
        ):
10  # - Follow branches based on predictor conditions.
11  # - Terminal nodes give the predicted average bill_avg for
        observations in that region.
12  # - Example from marking scheme plot:
13  #    "Lowest avg bill (4.589) for: download_avg < 333.2 AND download_
        over_limit < 0.5 AND
14  #     upload_avg < 15.25 AND subscription_age < 0.075 (almost new
        customers)."
15  #    "Highest avg bill (264.6) for: download_avg >= 333.2 AND is_tv_
        subscriber=0 (No)
16  #     AND download_avg >= 861.5."
17  # Tree structure highlights important variables and interaction-like
         effects through sequential splits.
```

Listing 91: 1f: Regression Tree (Churn) - Base R (Exam 2023)

3.3.7 1g: Evaluate regression tree

Question *Predict `bill_avg` with the regression tree in 1f) and evaluate the predictions.*

Base R Solution

```
# Assuming tree_bill_avg from 1f and test_churn from 1a
pred_tree_bill_avg_test <- predict(tree_bill_avg, newdata=test_churn
    )
mse_tree_bill_avg_test <- mean((test_churn$bill_avg - pred_tree_bill
    _avg_test)^2, na.rm=TRUE)
cat("Test MSE (Regression Tree for bill_avg):", mse_tree_bill_avg_
    test, "\n")
# Marking guide value: 120.2793 (comparable to OLS/LASSO, slightly
    worse)
```

Listing 92: 1g: Evaluate Regression Tree (Churn) - Base R (Exam 2023)

3.3.8 1h: Random forest for `bill_avg`

Question *Fit a random forest to `bill_avg`. Plot a variable importance measure for the predictors, interpret it, and, briefly, explain the measure. (If computations are too slow, reduce `ntree` or use smaller training data.)*

Base R Solution (using `randomForest`)

```
library(randomForest)
# Assuming train_churn from 1a
set.seed(86554354) % \textit{(Consistent seed)} %
# Marking guide used ntree=50 for speed
rf_bill_avg <- randomForest(bill_avg ~ ., data=train_churn,
                            ntree=50, importance=TRUE, na.action=na.
                               roughfix)
# importance=TRUE is needed for varImpPlot. na.action handles
    potential NAs if any.

# Variable Importance
# print(importance(rf_bill_avg))
varImpPlot(rf_bill_avg, main="Variable Importance for bill_avg (
    Random Forest)")
# Explanation of IncNodePurity (default for regression):
# "IncNodePurity measures the total decrease in node impurity (RSS
    for regression trees)
# from splitting on that variable, averaged over all trees in the
    forest.
# A higher value indicates that the variable is more important for
    partitioning the data
# and improving the homogeneity of nodes regarding bill_avg."
# Interpretation: "Based on the plot, [predictor1] and [predictor2]
    appear most important.
# download_avg and upload_avg, which had unclear patterns in initial
     plots, now show high
# importance, contrasting with the single tree and OLS findings (
    from marking guide)."
```

Listing 93: 1h: Random Forest (Churn) - Base R (Exam 2023)

3.3.9 1i: Evaluate random forest

Question *Use the model in 1h) to predict `bill_avg` and evaluate the predictions.*

Base R Solution

```r
# Assuming rf_bill_avg from 1h and test_churn from 1a
pred_rf_bill_avg_test <- predict(rf_bill_avg, newdata=test_churn, na
    .action=na.pass)
mse_rf_bill_avg_test <- mean((test_churn$bill_avg - pred_rf_bill_avg
    _test)^2, na.rm=TRUE)
cat("Test MSE (Random Forest for bill_avg):", mse_rf_bill_avg_test,
    "\n")
# Marking guide value: 96.54472 (significantly lower than OLS, LASSO
    , single tree)
```

Listing 94: 1i: Evaluate Random Forest (Churn) - Base R (Exam 2023)

3.3.10   1j: Features of customers with high/low `bill_avg`

Question   *Based on your analysis in 1a)-1i, what are the features of customers with high and, respectively, low average bills?*

Solution Outline (Textual)   *Synthesize findings from all models, focusing on consistent patterns and insights from the best performing model (Random Forest).*

Task 2: Predict `churn`

3.3.11   2a: Data formatting, descriptive statistics for `churn`

Question   *Make sure that all variables are on the right format for your analysis. Use tables and graphs and common sense to remove variables that you think will not be helpful. Motivate your choices thoroughly. Use descriptive statistics to find promising predictors for **churn**.*

Solution Outline (Base R / Tidyverse)

```r
# Assuming xdata_churn from 1a (id removed, specified columns are
    factors)
# and train_churn (the training split)
# Data format should be okay from 1a if factors were made correctly.
# Remove unhelpful variables: 'bill_avg' might be considered a
    consequence of churn or post-churn,
# or a predictor. Marking guide implies it's kept. No other obvious
    removals for churn prediction.

# Descriptive stats on train_churn:
# Base R:
# cat_preds_churn <- names(train_churn)[sapply(train_churn, is.
    factor) & names(train_churn)!="churn"]
# num_preds_churn <- names(train_churn)[sapply(train_churn, is.
    numeric)]
# for (pred in c(cat_preds_churn, num_preds_churn)) {
#   if (is.factor(train_churn[[pred]])) {
#     print(paste("Table for churn vs", pred))
#     print(prop.table(table(Churn=train_churn$churn, Predictor=
    train_churn[[pred]]), margin=2))
#   } else {
#     #boxplot(train_churn[[pred]] ~ train_churn$churn, main=paste(
    pred, "by Churn"))
#   }
# }
```

```
19 # Tidyverse (as in marking guide for means):
20 # library(dplyr)
21 # train_churn %>%
22 #   mutate(churn_numeric = as.numeric(churn)-1) %>% # if churn is
         factor "0","1" or "No","Yes"
23 #   group_by(churn) %>%
24 #   summarise(across(where(is.numeric) & !is.factor, list(mean=mean)
         ),
25 #             across(where(is.factor) & !matches("churn"), ~mean(as.
         numeric(.)-1, na.rm=TRUE))) %>%
26 #   t() %>% print()
27 # plot(remaining_contract ~ churn, data=train_churn, main="Remaining
         Contract by Churn")
28
29 # Promising predictors (from marking scheme):
30 # is_movie_package_subscriber (lower for churners), remaining_
         contract (much lower for churners),
31 # download_avg (lower for churners), upload_avg (lower for churners)
         ,
32 # download_over_limit (higher for churners).
```

Listing 95: 2a: Descriptives for churn (Churn) - Conceptual (Exam 2023)

3.3.12   2b: Bootstrap CI for P(churn=1)

Question   *Let Y be a stochastic variable equal to 1 if a customer churn and 0
otherwise and let $p = P(Y = 1)$ be the unconditional probability to churn. The
first 50 (original, not from the training or test data) observations of the variable*
churn *contains observations drawn from the stochastic variable Y. Use the bootstrap
to compute a 95% confidence interval for p. Compare the standard approximation
$\hat{p} \pm 1.96\sqrt{\hat{p}(1-\hat{p})/n}$ where $\hat{p}$ is the sample fraction of churners and n is the number
of observations.*

Base R Solution
```
1  # Using original xdata_churn (after factor conversion in 1a)
2  # churn_numeric_orig <- as.numeric(xdata_churn$churn) - 1 # Convert
      factor to 0/1
3  # first_50_churn <- churn_numeric_orig[1:50]
4  # n_obs_50 <- 50
5
6  # Bootstrap function for proportion
7  prop_churn_stat <- function(data, index) {
8    return(mean(data[index])) # Mean of 0/1 variable is proportion of
        1s
9  }
10 set.seed(86554354) # Consistent seed
11 boot_churn_prop <- boot(data=first_50_churn, statistic=prop_churn_
      stat, R=1000)
12 # print(boot_churn_prop)
13
14 # Bootstrap Percentile CI
15 ci_perc_churn <- boot.ci(boot_churn_prop, type="perc")
16 print("Bootstrap Percentile CI for P(churn=1):")
17 print(ci_perc_churn)
18 # E.g., (0.68, 0.90) from marking guide.
19
20 # Standard Normal Approximation CI
21 p_hat_50 <- mean(first_50_churn)
```

```
22  se_p_hat_50 <- sqrt(p_hat_50  (1 - p_hat_50) / n_obs_50)
23  ci_norm_approx_churn <- c(p_hat_50 - 1.96  se_p_hat_50, p_hat_50 +
        1.96  se_p_hat_50)
24  print("Normal Approximation CI for P(churn=1):")
25  print(ci_norm_approx_churn)
26  # E.g., (0.689, 0.911) from marking guide.
27  # Comparison: "The results are very similar, indicating the normal
        approximation works well for this sample size and proportion."
```

Listing 96: 2b: Bootstrap CI for P(churn) (Churn) - Base R (Exam 2023)

3.3.13 2c: Logistic regression for `churn`

Question *Fit a logistic regression with all variables to* **churn**. *Interpret the coefficients in front of* **is_tv_subscriber** *and* **is_movie_package_subscriber**. *For all coefficients, is the sign as you expected?*

Base R Solution

```
1   # Assuming train_churn from 1a (with factors, including original '
        churn' factor)
2   logit_churn_allvars <- glm(churn ~ ., data=train_churn, family=
        binomial)
3   summary_logit_churn <- summary(logit_churn_allvars)
4   print(summary_logit_churn)
5
6   # Interpretation of coefficients:
7   # coefs_churn <- coef(logit_churn_allvars)
8   # odds_ratios_churn <- exp(coefs_churn)
9   # For is_tv_subscriber1 (assuming '1' means Yes):
10  # OR_tv = exp(coefs_churn["is_tv_subscriber1"]) e.g., exp(-1.75) =
        0.17
11  # "Holding other variables constant, customers with a TV
        subscription have 0.17 times
12  # the odds of churning compared to those without a TV subscription (
        i.e., 83% lower odds)."
13  # For is_movie_package_subscriber1:
14  # OR_movie = exp(coefs_churn["is_movie_package_subscriber1"]) e.g.,
        exp(-0.06) = 0.94
15  # "Holding other variables constant, customers with a movie package
        have 0.94 times
16  # the odds of churning compared to those without (i.e., 6% lower
        odds)."
17  # Expected signs:
18  # Negative for TV/Movie sub (loyalty), subscription_age (more
        invested), remaining_contract (locked in).
19  # Positive for service_failure_count, download_over_limit.
20  # Ambiguous for bill_avg, download_avg, upload_avg (could be good
        service or high cost).
21  # Check significance (p-values). Marking guide: bill_avg, upload_avg
         not significant.
```

Listing 97: 2c: Logistic Regression for churn (Churn) - Base R (Exam 2023)

3.3.14 2d: Evaluate logistic regression, consider removing variables

Question *Use the logistic regression from 2c) to predict* **churn**. *Evaluate the predictions in an appropriate way. Are the predictions improved if you remove some variables from the model?*

Base R Solution

```r
# Assuming logit_churn_allvars from 2c (fitted on train_churn) and
    test_churn
probs_logit_test_churn <- predict(logit_churn_allvars, newdata=test_
    churn, type="response")
# Default threshold 0.5 for initial evaluation
pred_class_logit_churn_05 <- ifelse(probs_logit_test_churn > 0.5,
                                    levels(train_churn$churn)[2],
                                      levels(train_churn$churn)[1])
pred_class_logit_churn_05 <- factor(pred_class_logit_churn_05,
    levels=levels(train_churn$churn))

conf_matrix_logit_05 <- table(Actual=test_churn$churn, Predicted=
    pred_class_logit_churn_05)
print("Confusion Matrix (Full Model, Threshold 0.5):")
print(conf_matrix_logit_05)
# accuracy_logit_05 <- sum(diag(conf_matrix_logit_05)) / sum(conf_
    matrix_logit_05)
# print(paste("Accuracy (Full Model):", accuracy_logit_05))
# print("Proportions table (Full Model):")
# print(prop.table(conf_matrix_logit_05, margin=1))
# Marking scheme results for full model: (row proportions)
#    FALSE       TRUE
# 0 0.81606725 0.18393275
# 1 0.08262943 0.91737057 (Good at predicting non-churn, good at
    predicting churn when it occurs)

# Model with non-significant variables (bill_avg, upload_avg)
    removed:
logit_churn_reduced <- glm(churn ~ . -bill_avg -upload_avg, data=
    train_churn, family=binomial)
# summary(logit_churn_reduced)
probs_logit_reduced_test <- predict(logit_churn_reduced, newdata=
    test_churn, type="response")
pred_class_logit_reduced_05 <- ifelse(probs_logit_reduced_test >
    0.5,
                                        levels(train_churn$churn)[2],
                                          levels(train_churn$churn)
                                          [1])
pred_class_logit_reduced_05 <- factor(pred_class_logit_reduced_05,
    levels=levels(train_churn$churn))
conf_matrix_logit_reduced_05 <- table(Actual=test_churn$churn,
    Predicted=pred_class_logit_reduced_05)
print("Confusion Matrix (Reduced Model, Threshold 0.5):")
print(conf_matrix_logit_reduced_05)
# print("Proportions table (Reduced Model):")
# print(prop.table(conf_matrix_logit_reduced_05, margin=1))
# Marking scheme: "removal of non-significant variables did not
    improve the test confusion matrix much."
# (Compare the prop.tables or specific metrics like sensitivity/
    specificity).
```

Listing 98: 2d: Evaluate Logistic Regression (Churn) - Base R (Exam 2023)

3.3.15 2e: Random forest for `churn`

Question *Use a random forest to predict **churn** and evaluate the predictions. (If computations are too slow, reduce **ntree** or use smaller training data.)*

```
1  library(randomForest)
2  # Assuming train_churn, test_churn from 1a
3  set.seed(86554354) % \textit{(Consistent seed)} %
4  rf_churn <- randomForest(churn ~ ., data=train_churn,
5                           ntree=50, # Reduced for speed as per hint/
                              marking scheme
6                           importance=TRUE, na.action=na.roughfix)
7
8  # Predict classes directly
9  pred_class_rf_churn <- predict(rf_churn, newdata=test_churn, na.
      action=na.pass)
10 conf_matrix_rf_churn <- table(Actual=test_churn$churn, Predicted=
      pred_class_rf_churn)
11 print("Confusion Matrix (Random Forest):")
12 print(conf_matrix_rf_churn)
13 # accuracy_rf_churn <- sum(diag(conf_matrix_rf_churn)) / sum(conf_
      matrix_rf_churn)
14 # print(paste("Accuracy (Random Forest):", accuracy_rf_churn))
15 # print("Proportions table (Random Forest):")
16 # print(prop.table(conf_matrix_rf_churn, margin=1))
17 # Marking scheme results (row proportions):
18 #    FALSE      TRUE
19 # 0 0.94564187 0.05435813
20 # 1 0.06290371 0.93709629
21 # "The random forest improved the predictions considerably" (
      compared to logistic regression).
22 # Higher true positive for churn (0.937 vs 0.917) and higher true
      negative (0.945 vs 0.816).
```

Listing 99: 2e: Random Forest for churn (Churn) - Base R (Exam 2023)

3.3.16 2f: Typical features of customers who churn

Question *Based on your analysis in 2a)-2e), what are the typical features of customers who churn?*

Solution Outline (Textual) *Synthesize findings from descriptive stats, logistic regression (significant coefficients, signs), and random forest (variable importance if examined, though not explicitly asked for plot here).*

3.4 Exam Spring 2024

For variable definitions, see the table at the end of this exam section.

Task 1: Methodological Topics

3.4.1 1a: Explain R-function

Question *Explain what the following R-function is doing.*

```
1  f <- function(x0, x, y, K=3) {
2    d <- abs(x - x0)
3    o <- order(d)[1:K]
4    xl <- x[o]
5    yl <- y[o]
6    xydata <- data.frame(xl=xl, yl=yl)
```

```
7    reg <- lm(yl ~ xl, data=xydata)
8    ypred <- predict(reg, newdata=data.frame(xl=x0, yl=1)) % \textit{(
         yl=1 is placeholder)} %
9    return(ypred)
10 }
```

Listing 100: R-function f for 1a (Exam 2024)

Explanation   *This R function* $f(x0, x, y, K=3)$ *implements a local regression prediction method. 1. It takes a target predictor value* $x0$, *a vector of training predictor values* x, *corresponding training response values* y, *and the number of neighbors K (default 3). 2. It calculates the absolute distances (*d*) between* $x0$ *and all values in* x. *3. It finds the indices (*o*) of the K values in* x *that are closest to* $x0$. *4. It selects these K closest predictor values (*xl*) and their corresponding response values (*yl*). 5. It creates a data frame* $xydata$ *from these K neighbors. 6. It fits a simple linear regression model (*reg*) of* yl *on* xl *using only these K neighbors. 7. Finally, it uses this locally fitted linear model to predict the response (*$ypred$*) for the original target point* $x0$. *The method resembles K-Nearest Neighbors (KNN) regression but instead of averaging the responses of the neighbors, it fits a linear model to them for prediction. The* $yl=1$ *in* $predict()$ *is a placeholder and not used when predicting for a new* xl.

3.4.2   1b: Optimal K for function f using LOOCV

Question   *Consider the following small dataset:* $x = [1, 2, 3, 4, 5, 6, 7, 8, 9, 10]$, $y = [5.26, 9.13, 11.17, 15.64, 25.32, 25.55, 41.39, 48.17, 58.65, 68.24]$. *Use leave-one-out cross-validation (LOOCV) to determine the optimal K in the function* f *for this dataset.*

Base R Solution (basert på løsningsforslag)

```
1  # Dataset
2  x_1b <- c(1, 2, 3, 4, 5, 6, 7, 8, 9, 10)
3  y_1b <- c(5.26, 9.13, 11.17, 15.64, 25.32, 25.55, 41.39, 48.17,
       58.65, 68.24)
4  n_1b <- length(x_1b)
5
6  # Function f from 1a (ensure it's defined)
7  f <- function(x0, x, y, K=3) {
8    d <- abs(x - x0); o <- order(d)[1:K]; xl <- x[o]; yl <- y[o]
9    xydata <- data.frame(xl=xl, yl=yl); reg <- lm(yl ~ xl, data=xydata
       )
10   return(predict(reg, newdata=data.frame(xl=x0))) # Removed yl=1
       placeholder
11 }
12
13 # LOOCV function (as per solution structure)
14 loo_cv_func <- function(K_val, x_full, y_full) {
15   n_loo <- length(x_full)
16   squared_errors <- numeric(n_loo)
17   for (i in 1:n_loo) {
18     x_train <- x_full[-i]
19     y_train <- y_full[-i]
20     x_test_single  <- x_full[i]
21     y_test_single  <- y_full[i]
22
```

```
23      # Ensure K is not greater than number of available training
            points for local lm
24      K_to_use <- min(K_val, length(x_train))
25      if (K_to_use < 2) K_to_use <- 2 # lm needs at least 2 points for
            y~x
26
27      pred_loo <- f(x_test_single, x_train, y_train, K=K_to_use)
28      squared_errors[i] <- (y_test_single - pred_loo)^2
29    }
30    return(mean(squared_errors))
31  }
32
33  # Iterate for different K values (solution implies K from 2 to 10,
        but K=n-1 for LOOCV training set)
34  # K can range from 2 (min for lm) up to n-1 (using all training
        points for local lm)
35  possible_K_vals <- 2:(n_1b-1)
36  loocv_mse_values <- sapply(possible_K_vals, loo_cv_func, x_full=x_1b
        , y_full=y_1b)
37
38  # Plot results and find optimal K
39  plot(possible_K_vals, loocv_mse_values, type="b", xlab="K (Number of
         Neighbors)",
40       ylab="LOOCV MSE", main="LOOCV to find Optimal K")
41  optimal_K_1b <- possible_K_vals[which.min(loocv_mse_values)]
42  points(optimal_K_1b, min(loocv_mse_values), col="red", pch=19, cex
        =1.5)
43  cat("Optimal K found by LOOCV:", optimal_K_1b, "\n")
44  # Exam solution PDF indicates K=5 is optimal.
```

Listing 101: 1b: LOOCV for Optimal K - Base R (Exam 2024)

3.4.3 1c: Plot predictions with optimal K

Question *Plot y against x and add a line with the predictions based on the optimal K.*

Base R Solution
```
1   # Assuming x_1b, y_1b from 1b, and optimal_K_1b = 5
2   optimal_K_plot <- 5
3   predictions_1c <- numeric(n_1b)
4   # To generate a smooth line, we'd typically predict on a grid of x0
        values.
5   # However, to match the solution's likely intent (predicting on
        training points for the line):
6   for (i in 1:n_1b) {
7     # Here, for plotting the fit on training data, f uses the full
          dataset x_1b, y_1b as its "training"
8     # to find neighbors for x_1b[i]. This is for visualizing the
          fitted function.
9     predictions_1c[i] <- f(x_1b[i], x_1b, y_1b, K=optimal_K_plot)
10  }
11
12  plot(x_1b, y_1b, xlab="x", ylab="y", main=paste("Data and Fitted
        Line with K =", optimal_K_plot))
13  order_x_1b <- order(x_1b)
14  lines(x_1b[order_x_1b], predictions_1c[order_x_1b], col="red", lwd
        =2)
15  legend("topleft", legend=paste("Fit with K =", optimal_K_plot), col=
        "red", lty=1, lwd=2)
```

3.4.4 1d: Modify function f for multiple predictors

Question *In the function f, there is only one predictor. One way to allow for more than one predictor is to compute **d** in the f-function in (a) differently. Explain how such a modification can be done; exemplify with a case with two predictors.*

Explanation and Conceptual R Modification

Example with two predictors (X_1, X_2): Let input x0 be a vector `c(x01, x02)`, and x be a matrix where rows are observations and columns are X_1, X_2. The modified d calculation within f would be: *To adapt function f for multiple predictors $(p > 1)$, the distance calculation **d** needs to be changed from `abs(x - x0)` to a multivariate distance. The most common is Euclidean distance.*

Example with two predictors (X_1, X_2)*: Let input **x0** be a vector **c(x01, x02)**, and **x** be a matrix where rows are observations and columns are X_1, X_2. The modified **d** calculation within **f** would be:*

```
# Inside function f_multi(x0_vec, x_matrix, y_vec, K)
# x0_vec is c(x0_pred1, x0_pred2)
# x_matrix has columns for pred1, pred2

# Calculate Euclidean distances
# diff_sq <- sweep(x_matrix, 2, x0_vec, "-")^2 # (x_i1-x01)^2, (x_i2
    -x02)^2 for each row
# d_multi <- sqrt(rowSums(diff_sq))

# The rest of the logic for finding K nearest neighbors (o, xl, yl)
     remains similar,
# but xl will be a matrix of K rows and p columns.
# The lm call becomes: reg <- lm(yl ~ ., data=as.data.frame(xl))
# And newdata for predict needs to be a data frame with column names
     matching xl.
# newdata_pred <- as.data.frame(matrix(x0_vec, nrow=1,
#                                      dimnames=list(NULL, colnames(
    x_matrix))))
```

Listing 103: Conceptual modification for distance in f (2 predictors)

Crucially, if predictors are on different scales, they should be standardized before calculating Euclidean distances to prevent variables with larger scales from dominating the distance metric. The linear model `lm(yl ~ .)` would then be a multiple linear regression on the K selected neighbors' p (standardized, if applicable) predictor values.

3.4.5 1e: Implement backfitting for this method

Question *Another way to allow for several predictors is to use backfitting. Explain how you would implement backfitting for this method (you do not need to do it).*

Explanation

Backfitting Algorithm with function f (example for $p = 2$ ***predictors*** X_1, X_2***):***

1. ***Initialize****:* $\hat{\beta}_0 = mean(Y)$. *Initialize* $\hat{f}_1(x_{i1}) = 0$ *and* $\hat{f}_2(x_{i2}) = 0$ *for all observations* i.

2. ***Iterate*** *until convergence (i.e.,* $\hat{f}_j$ *functions stabilize):*

   (a) ***Update*** $\hat{f}_1(X_1)$*: Compute partial residuals:* $r_{i1} = y_i - \hat{\beta}_0 - \hat{f}_2(x_{i2})$. *For each unique value* x_{u1} *of* X_1 *(or for each* x_{i1}*):* $\hat{f}_1(x_{u1}) \leftarrow$ *Call* $f(x_{u1}, X_1, r_1, K)$. *(Typically, one would re-center:* $\hat{f}_1 \leftarrow \hat{f}_1 - mean(\hat{f}_1)$*).*

   (b) ***Update*** $\hat{f}_2(X_2)$*: Compute partial residuals:* $r_{i2} = y_i - \hat{\beta}_0 - \hat{f}_1(x_{i1})$ *(using updated* $\hat{f}_1$*). For each unique value* x_{u2} *of* X_2*:* $\hat{f}_2(x_{u2}) \leftarrow$ *Call* $f(x_{u2}, X_2, r_2, K)$. *(Re-center:* $\hat{f}_2 \leftarrow \hat{f}_2 - mean(\hat{f}_2)$*).*

The final prediction for a new observation $(\mathbf{x}_{01}, \mathbf{x}_{02})$ *is* $\hat{Y}_0 = \hat{\beta}_0 + \hat{f}_1(\mathbf{x}_{01}) + \hat{f}_2(\mathbf{x}_{02})$. *The function* ***f*** *from 1a acts as the univariate smoother for each predictor on its partial residuals. The optimal* K *for each* f_j *can be chosen using CV, possibly within each backfitting iteration or as a global parameter. Backfitting allows fitting an additive model* $Y = \beta_0 + f_1(X_1) + f_2(X_2) + \cdots + f_p(X_p) + \epsilon$, *where each* f_j *is estimated using a univariate smoother, such as our function* ***f*** *(from 1a, adapted for single predictor use).*

Task 2: Analysis of `airline.csv` Data - Customer Satisfaction

Target variable: `satisfaction`. *Often converted to binary for classification.*

3.4.6   2a: Recode categorical variables and split data

Question   *If necessary, recode categorical variables as factors. Motivate your choices in words. (Implicit: Split data for training/testing as per general instructions and solution PDF structure for this exam).*

Base R Solution (based on solution PDF)

```
1  airline <- read.csv("airline.csv", stringsAsFactors = FALSE) % \
       textit{Read data} %
2  # Convert 'satisfaction' to factor (assuming 'satisfied' vs '
       dissatisfied' for classification)
3  # This step depends on how 'satisfaction' is defined in the original
        CSV.
4  # If it's text "satisfied"/"dissatisfied", it's already good for
       factor.
5  # If numeric, might need binning. Let's assume it becomes a factor.
6  # airline$satisfaction <- as.factor(airline$satisfaction)
7
8  # Recode variables with <= 7 unique values as factors (as per
       solution for this exam)
9  for (i in 1:ncol(airline)) {
10    if (length(unique(airline[,i])) <= 7 && !is.factor(airline[,i])) {
11      airline[,i] <- as.factor(airline[,i])
12    }
13  }
14  # Example: Gender, Customer_Type, Type_of_Travel, Class are likely
       candidates.
15  # Satisfaction scores (0-5) might also be converted if not already
       factors.
16  # Motivation: Ensure categorical predictors are treated as such by
       modeling functions.
17
18  # Split data (as per solution PDF structure for this exam)
19  set.seed(123) # From solution PDF for this task
20  n_airline <- nrow(airline)
21  train_idx_airline <- sample(1:n_airline, floor(n_airline/2))
22  training_airline <- airline[train_idx_airline, ]
23  test_airline <- airline[-train_idx_airline, ]
24  # sapply(training_airline, class)
```

Listing 104: 2a: Recode Categoricals Split (airline) - Base R (Exam 2024)

Tidyverse Solution

```
1  library(dplyr)
2  airline_tidy <- read.csv("airline.csv", stringsAsFactors = FALSE)
       %>%
3    mutate(across(where(~length(unique(.)) <= 7 & !is.factor(.)), as.
       factor))
4    # May need specific mutate for 'satisfaction' if it's not binary
       textual
5    # e.g. mutate(satisfaction = factor(ifelse(satisfaction_score >=
       4, "satisfied", "dissatisfied")))
6
7  set.seed(123)
8  train_idx_airline_tidy <- sample(1:nrow(airline_tidy), floor(nrow(
       airline_tidy)/2))
9  training_airline_tidy <- airline_tidy[train_idx_airline_tidy, ]
10  test_airline_tidy <- airline_tidy[-train_idx_airline_tidy, ]
11  # glimpse(training_airline_tidy)
```

Listing 105: 2a: Recode Categoricals Split (airline) - Tidyverse (Exam 2024)

3.4.7 2b: Motivation for data usage (Inference: "Why dissatisfied?")

Question *In tasks (a)-(g) the methods and models are used to answer the question "Why are some customers dissatisfied?". With this in mind, motivate your choice to evaluate the models on all/training/test data.*

Solution/Motivation (based on exam solution PDF)

*For tasks (a)-(g), which focus on understanding why customers are dissatisfied (an inferential goal rather than pure out-of-sample prediction), one could argue for using the entire dataset to maximize statistical power for inference. However, the provided solution chooses to **split the data into training and test sets from the start, and perform all descriptive analysis and initial model fitting (for inference) on the training data only**. The motivations for this approach are:*

1. ***Comparability**: To allow direct comparison of model performance (e.g., R-squared, variable importance) with models developed later for the explicit prediction task (h-l), which must be evaluated on test data.*

2. ***Preventing Data Leakage**: Using only training data for initial exploration and model specification for the inferential part acts as a safeguard against inadvertently using information from the test set to guide these early modeling choices. This ensures that when the test set is used for the later prediction task, it is truly "unseen".*

3.4.8 2c: Descriptive statistics for

Question *Use descriptive statistics to find variables associated with* `satisfaction`*. First, explain which type of descriptive statistics (types of tables and figures) that you are using. Give examples of R code but do not show all code and output if you are doing the same thing many times. Summarize your results as text, mentioning the relevant numbers.*

Base R Solution (Conceptual based on exam solution)

```
# Assuming training_airline from 2a. Assume 'satisfaction' is a
    binary factor ("satisfied", "dissatisfied").
# Overall satisfaction rate:
# print(prop.table(table(training_airline$satisfaction))) # e.g.,
    55% satisfied

# Function for conditional proportions (as per exam solution)
proptab_satisfaction <- function(data, predictor_col_name) {
  tbl <- table(data[[predictor_col_name]], data$satisfaction)
  return(round(100  prop.table(tbl, margin = 1), 1)) # Proportion
      satisfied/dissatisfied per predictor level
}
# Example: proptab_satisfaction(training_airline, "Gender")
#             satisfaction
# Gender      dissatisfied satisfied
#    Female              35        65
#    Male                56        44
# (Values from exam solution PDF for illustration)

# For numeric variables: Boxplots and grouped summaries
# boxplot(Age ~ satisfaction, data=training_airline, ylab="Age")
# by(training_airline$Flight_Distance, training_airline$satisfaction
    , summary)
```

```
20
21  # Summary of results (textual, based on exam solution PDF):
22  # - Overall 55% satisfied.
23  # - Gender: Females (65% satisfied) more satisfied than Males (44%).
24  # - Customer_Type: Loyal customers more satisfied.
25  # - Type_of_Travel, Class: Business travelers/class more satisfied.
26  # - Specific satisfaction Qs (Inflight_entertainment, Cleanliness,
        etc.): Positively associated with overall satisfaction.
27  # - Age: Satisfied customers tend to be older on average.
28  # - Flight_Distance, Delays: No strong or obvious association with
        overall satisfaction from boxplots/summaries.
```

Listing 106: 2c: Descriptives for Satisfaction (airline) - Base R (Exam 2024)

Tidyverse Solution (Conceptual)

```
1   # library(dplyr); library(ggplot2)
2   # Assuming training_airline_tidy from 2a.
3   # training_airline_tidy %>% count(satisfaction) %>% mutate(prop = n/
        sum(n))
4
5   # Conditional proportions
6   # training_airline_tidy %>%
7   #   count(Gender, satisfaction) %>%
8   #   group_by(Gender) %>%
9   #   mutate(prop = round(100  n / sum(n), 1)) %>% print()
10
11  # Boxplots
12  # ggplot(training_airline_tidy, aes(x=satisfaction, y=Age, fill=
        satisfaction)) + geom_boxplot()
13  # Grouped summaries
14  # training_airline_tidy %>% group_by(satisfaction) %>%
15  #   summarise(mean_age = mean(Age, na.rm=TRUE), median_dist = median
        (Flight_Distance, na.rm=TRUE))
```

Listing 107: 2c: Descriptives for Satisfaction (airline) - Tidyverse (Exam 2024)

3.5   2d: Logistic Regression (Why dissatisfied?)

```
1   logreg_infer <- glm(satisfaction ~ Gender + Customer_Type + Age +
        Type_of_Travel +
2       Class + Flight_Distance + Food_and_drink + Inflight_
            entertainment +
3       Online_support + Baggage_handling + Checkin_service + Arrival_
            Delay_in_Minutes +
4       Inflight_wifi_service + Leg_room_service + On_board_service +
            Gate_location + Seat_comfort,
5       data=training_airline, family=binomial)
6
7   probs_logreg_infer_test <- predict(logreg_infer, newdata=test_
        airline, type="response")
8   preds <- ifelse(probs_logreg_infer_test > 0.5, "satisfied", "
        dissatisfied")
9   conf_matrix <- table(Actual = test_airline$satisfaction, Predicted =
        preds)
10  print(conf_matrix)
```

Listing 108: 2d: Logistic Regression - Inferential (airline)

3.6 2e: Classification Tree (Why dissatisfied?)

```
library(tree)
tree_model <- tree(satisfaction ~ Gender + Customer_Type + Age +
    Type_of_Travel +
    Class + Flight_Distance + Food_and_drink + Inflight_
        entertainment +
    Online_support + Baggage_handling + Checkin_service + Arrival_
        Delay_in_Minutes +
    Inflight_wifi_service + Leg_room_service + On_board_service +
        Gate_location + Seat_comfort,
    data=training_airline)
plot(tree_model)
text(tree_model, pretty=0)
```

Listing 109: 2e: Classification Tree - Inferential

3.7 2f: Random Forest (Why dissatisfied?)

```
library(randomForest)
rf_model <- randomForest(satisfaction ~ Gender + Customer_Type + Age
    + Type_of_Travel +
    Class + Flight_Distance + Food_and_drink + Inflight_
        entertainment +
    Online_support + Baggage_handling + Checkin_service + Arrival_
        Delay_in_Minutes +
    Inflight_wifi_service + Leg_room_service + On_board_service +
        Gate_location + Seat_comfort,
    data=training_airline, importance=TRUE)
varImpPlot(rf_model)
```

Listing 110: 2f: Random Forest - Inferential

3.8 2g: Why are some customers dissatisfied?

Customers who are dissatisfied often report low satisfaction with inflight entertainment, seat comfort, online support, and baggage handling. These variables were consistently important in the models. Loyal customers and business class passengers are generally more satisfied. Demographic variables like age and gender also play a role.

3.9 2h: Assumptions for Prediction Task (Pre-flight variables)

Variables known before flight: Gender, Age, Customer Type, Type of Travel, Class, and Flight Distance. Satisfaction ratings and delay variables are post-flight and not available for prediction before departure.

3.10 2i: Motivation for Train/Test Split in Prediction

Since the goal is to predict dissatisfaction for new customers, it is crucial to evaluate the model on unseen test data to assess generalization performance.

117

3.11 2j: Logistic Regression (Prediction)

```
logreg_pred <- glm(satisfaction ~ Gender + Customer_Type + Age +
    Type_of_Travel + Class + Flight_Distance,
  data=training_airline, family=binomial)
probs_pred <- predict(logreg_pred, newdata=test_airline, type="
    response")
preds <- ifelse(probs_pred > 0.5, "satisfied", "dissatisfied")
conf_matrix <- table(Actual = test_airline$satisfaction, Predicted =
    preds)
print(conf_matrix)
```

Listing 111: 2j: Logistic Regression - Prediction

3.12 2k: Random Forest (Prediction)

```
rf_pred <- randomForest(satisfaction ~ Gender + Customer_Type + Age
    + Type_of_Travel + Class + Flight_Distance,
  data=training_airline, importance=TRUE)
preds_rf <- predict(rf_pred, newdata=test_airline)
conf_matrix <- table(Actual = test_airline$satisfaction, Predicted =
    preds_rf)
print(conf_matrix)
```

Listing 112: 2k: Random Forest - Prediction

3.13 2l: Future Data Collection

To improve predictions, collect data on pre-flight interactions: booking channel, time between booking and flight, seat selection, and support tickets. These can act as proxies for future dissatisfaction and help model potential outcomes more accurately.

4 Tutorials

This section summarizes key tasks and R code from the provided exercise sets, assuming they correspond to the course tutorials.

4.1 Tutorial 1

4.1.1 Task 1: Linear Regression and KNN (One Predictor)

Context *Given a small dataset (x, y pairs).*

1a: Prediction with Linear Regression (Manual/Calculator) *Question: What is the prediction for x = 3 if the linear regression model is used? Solution Approach: Manually calculate $\hat{\alpha}_0, \hat{\alpha}_1$ using formulas for simple linear regression, then predict $\hat{y} = \hat{\alpha}_0 + \hat{\alpha}_1 \cdot 3$.*

```
# x_t1 <- 1:5
# y_t1 <- c(1.88, 4.54, 10.12, 9.14, 11.26)
# n_t1 <- length(x_t1)
# mean_x_t1 <- mean(x_t1)
```

```
5   # mean_y_t1 <- mean(y_t1)
6   #
7   # beta1_hat_t1 <- sum((x_t1 - mean_x_t1)   (y_t1 - mean_y_t1)) / sum
        ((x_t1 - mean_x_t1)^2)
8   # beta0_hat_t1 <- mean_y_t1 - beta1_hat_t1   mean_x_t1
9   #
10  # prediction_x3_t1 <- beta0_hat_t1 + beta1_hat_t1   3
11  # print(prediction_x3_t1) % \textit{Solution gives 7.388} %
```

Listing 113: Manual OLS Calculation (Conceptual - Ex Set 1, T1a)

1b: Fit Linear Regression with `lm()` and Predict   *Question: Use `lm()` to fit the same model and `predict()` for x=3.*

```
1   # x_t1 <- 1:5
2   # y_t1 <- c(1.88, 4.54, 10.12, 9.14, 11.26)
3   # model_t1b <- lm(y_t1 ~ x_t1)
4   # predict_x3_t1b <- predict(model_t1b, newdata = data.frame(x_t1 =
        3))
5   # print(predict_x3_t1b) % \textit{Solution gives 7.388} %
```

Listing 114: Using lm and predict (Ex Set 1, T1b)

1c: KNN Prediction   *Question: Prediction for x=3 if KNN with K=1, K=3, K=5 is used. Solution Approach (Manual/Calculator): For K=1: Find y-value corresponding to x closest to 3. For K=3: Find 3 y-values for x's closest to 3, then average them. For K=5: Average all 5 y-values.*

```
1   # y_t1 <- c(1.88, 4.54, 10.12, 9.14, 11.26)
2   # K=1 (x=3 is closest): 10.12
3   # K=3 (x=2,3,4 are closest): mean(c(4.54, 10.12, 9.14)) % \textit{
        Result: 7.933333} %
4   # K=5 (all points): mean(y_t1) % \textit{Result: 7.388} %
```

Listing 115: KNN Predictions (Conceptual - Ex Set 1, T1c)

1d: Understanding KNN R-function   *Question: Explain each row of the provided KNN R-function.*

```
1   knn_func_t1d <- function(x0, x, y, K=20) { % \textit{Note: Exercise
        sheet has K=20, example uses K=3} %
2     d <- abs(x-x0)                    % \textit{# Calculate absolute
        distances from x0 to all x} %
3     o <- order(d)[1:K]               % \textit{# Get indices of the K
        smallest distances} %
4     ypred <- mean(y[o])              % \textit{# Average y-values of the
         K nearest neighbors} %
5     return(ypred)                    % \textit{# Return the prediction}
        %
6   }
7   # Experiment: K_exp=3; x0_exp=3; x_exp=1:5; y_exp=c
        (1.88,4.54,10.12,9.14,11.26)
8   # d_exp <- abs(x_exp - x0_exp) % Result: 2 1 0 1 2 %
9   # o_exp <- order(d_exp)[1:K_exp] % Result: 3 2 4 (indices for x=3, x
        =2, x=4) %
10  # ypred_exp <- mean(y_exp[o_exp]) % Result: mean(y[c(3,2,4)]) =
        7.933333 %
```

Listing 116: KNN R-function (Ex Set 1, T1d)

1e: Choosing Method Based on Plot *Question: Which method (LR or KNN K=1,2,3) for a given scatter plot? Why? Solution Approach: If plot shows strong linear trend, Linear Regression is preferred for efficiency and interpretability. If non-linear, KNN might be better, with K chosen to balance bias/variance.*

4.1.2 Task 2: Data Splitting and Basic Linear Model

2a: Load, Summarize, Help for `College`

```
library(ISLR)
# summary(College)
# help(College)
```
Listing 117: Loading College Data (Ex Set 1, T2a)

2b: 50/50 Training/Test Split

```
set.seed(123)
n_college <- nrow(College)
train_indicator_college <- sample(1:n_college, size=floor(n_college/
    2))
train_college <- College[train_indicator_college,]
test_college <- College[-train_indicator_college,]
```
Listing 118: Train/Test Split (Ex Set 1, T2b)

2c: Fit Linear Model for `Apps`   `Private + Accept`

```
lm_apps_t2c <- lm(Apps ~ Private + Accept, data=train_college)
# summary(lm_apps_t2c)
```
Listing 119: Linear Model for Apps (Ex Set 1, T2c)

2d: Compute Training MSE

```
pred_train_t2d <- predict(lm_apps_t2c, newdata=train_college)
mse_train_t2d <- mean((train_college$Apps - pred_train_t2d)^2)
# print(mse_train_t2d) % \textit{Solution shows 1437492} %
```
Listing 120: Training MSE (Ex Set 1, T2d)

2e: Model with `Accept` **only, compute Training MSE**

```
lm_apps_t2e <- lm(Apps ~ Accept, data=train_college)
pred_train_t2e <- predict(lm_apps_t2e, newdata=train_college)
mse_train_t2e <- mean((train_college$Apps - pred_train_t2e)^2)
# print(mse_train_t2e) % \textit{Solution shows 1440012, larger than
    2d, as expected} %
```
Listing 121: Reduced Linear Model and Training MSE (Ex Set 1, T2e)

2f: Compute Test MSE for models from 2c and 2e

```
pred_test_t2c <- predict(lm_apps_t2c, newdata=test_college)
mse_test_t2c <- mean((test_college$Apps - pred_test_t2c)^2)
# print(paste("Test MSE (Private + Accept):", mse_test_t2c)) % \
    textit{Sol: 1849668} %

pred_test_t2e <- predict(lm_apps_t2e, newdata=test_college)
mse_test_t2e <- mean((test_college$Apps - pred_test_t2e)^2)
# print(paste("Test MSE (Accept only):", mse_test_t2e)) % \textit{
    Sol: 1854227} %
# Model with Private + Accept has slightly better test MSE here.
```

Listing 122: Test MSE for both models (Ex Set 1, T2f)

2h: KNN for Apps ~ Accept, compare Test MSE

```
# Using knn_func_t1d from earlier, assuming K=3 as per solution's
    use of it
# Need to apply it for each test observation
# x0_knn_t2h <- test_college$Accept
# x_knn_t2h <- train_college$Accept
# y_knn_t2h <- train_college$Apps
#
# pred_knn_t2h <- sapply(x0_knn_t2h, knn_func_t1d, x=x_knn_t2h, y=y_
    knn_t2h, K=3)
# mse_knn_t2h <- mean((test_college$Apps - pred_knn_t2h)^2)
# print(paste("Test MSE (KNN K=3 for Apps ~ Accept):", mse_knn_t2h))
    % \textit{Sol: 4473360} %
# KNN performs much worse than linear regression here.
```

Listing 123: KNN for Apps (Ex Set 1, T2h)

4.1.3   Task 3: Exercise 3.10a-f from ISLR Book

This task involves fitting and interpreting a multiple linear regression model for `Sales ~ Price + Urban + US` *on the* `ISLR::Carseats` *dataset. Steps include fitting, interpreting coefficients, identifying significant predictors, and comparing models.*

```
# library(ISLR)
# data(Carseats)
# (a) Fit model
# model_carseats_310a <- lm(Sales ~ Price + Urban + US, data=
    Carseats)
# (b) Interpret coefficients
# summary(model_carseats_310a)
# (c) Write equation
# (d) Identify significant predictors (Price, US based on p-values)
# (e) Fit model with significant predictors only
# model_carseats_310e <- lm(Sales ~ Price + US, data=Carseats)
# (f) Compare R-squared and Adjusted R-squared (very similar for
    these two models).
# (g) Optional: Split, train, test.
# set.seed(12345)
# n_cs <- nrow(Carseats)
# train_idx_cs <- sample(1:n_cs, n_cs/2)
# train_cs_310 <- Carseats[train_idx_cs,]
# test_cs_310 <- Carseats[-train_idx_cs,]
# lm3_cs <- lm(Sales ~ Price + Urban + US, data=train_cs_310)
# pred3_cs <- predict(lm3_cs, newdata=test_cs_310)
```

```
20  # mse3_cs <- mean((test_cs_310$Sales - pred3_cs)^2) % Sol: 6.662185
        %
21  # lm4_cs <- lm(Sales ~ Price + US, data=train_cs_310)
22  # pred4_cs <- predict(lm4_cs, newdata=test_cs_310)
23  # mse4_cs <- mean((test_cs_310$Sales - pred4_cs)^2) % Sol: 6.580072
        %
24  # Reduced model slightly better on this test split.
```

Listing 124: ISLR 3.10a-f (Conceptual - Ex Set 1, T3)

4.2 Tutorial 2

4.2.1 Task 1: Logistic Regression and KNN for Binary Classification

Context *Small dataset (x=1:7, y binary).*

1a: Fit Logistic Regression, predict for x0=4

```
1  # x_t2 <- 1:7; y_t2 <- c(0,1,0,1,1,1,1)
2  # data_t2 <- data.frame(x=x_t2, y=as.factor(y_t2)) % \textit{Ensure
       y is factor for glm binomial} %
3  # m1_t2 <- glm(y ~ x, data=data_t2, family="binomial")
4  # summary(m1_t2)
5  # Estimated equation: P(Y=1|X=x) = 1 / (1 + exp(-(coef(m1_t2)[1] +
       coef(m1_t2)[2]x)))
6  # For x0=4: predict(m1_t2, newdata=data.frame(x=4), type="response")
        % Sol: 0.8628, predict Y=1 %
```

Listing 125: Logistic Regression Basics (Ex Set 2, T1a)

1b: Plot probabilities from Logistic Regression

```
1  # plot(x_t2, as.numeric(y_t2)-1, ylab="P(Y=1) or Y") # Plot 0/1 Y
       values
2  # probs_m1_t2 <- predict(m1_t2, type="response")
3  # lines(x_t2, probs_m1_t2, col="blue")
```

Listing 126: Plotting Logistic Probabilities (Ex Set 2, T1b)

1c: KNN probability estimation for x0=4 with K=1,3,5

```
1  # knn_prob_t2 <- function(x0, x_train, y_train, K_val) {
2  #   d <- abs(x_train - x0)
3  #   o <- order(d)[1:K_val]
4  #   return(mean(as.numeric(y_train[o])-1)) # Mean of 0/1 y-values
5  # }
6  # knn_prob_t2(4, x_t2, data_t2$y, K=1) % Sol: 1 (y at x=4 is 1) %
7  # knn_prob_t2(4, x_t2, data_t2$y, K=3) % Sol: mean(y at x=3,4,5) =
       mean(c(0,1,1))=0.667 %
8  # knn_prob_t2(4, x_t2, data_t2$y, K=5) % Sol: mean(y at x=2,3,4,5,6)
       =mean(c(1,0,1,1,1))=0.8 %
```

Listing 127: KNN Probabilities (Ex Set 2, T1c)

1d: Plot KNN probabilities for K=1,3,5 and logistic

```
1  # plot(x_t2, as.numeric(data_t2$y)-1, ylab="P(Y=1) or Y", xlab="x")
2  # for (K_val_plot in c(1,3,5)) {
3  #   probs_knn_plot <- sapply(x_t2, knn_prob_t2, x_train=x_t2, y_
       train=data_t2$y, K_val=K_val_plot)
4  #   lines(x_t2, probs_knn_plot, lty=K_val_plot+1, col=K_val_plot+1)
```

```
5   # }
6   # lines(x_t2, probs_m1_t2, col="blue", lty=1) # Logistic
7   # legend("bottomright", legend=c("Logistic", "KNN K=1", "KNN K=3", "
      KNN K=5"),
8   #        col=c("blue", 2, 3, 4), lty=c(1,2,3,4))
```

Listing 128: Plotting All Probabilities (Ex Set 2, T1d)

1e: Confusion matrices for logistic and KNN (K=3)

```
1    # Logistic predictions (threshold 0.5)
2    # pred_log_t2e <- ifelse(probs_m1_t2 > 0.5, 1, 0)
3    # table(Actual=as.numeric(data_t2$y)-1, Predicted=pred_log_t2e)
4    # prop.table(table(Actual=as.numeric(data_t2$y)-1, Predicted=pred_
       log_t2e), margin=1)
5
6    # KNN K=3 predictions
7    # probs_knn3_t2e <- sapply(x_t2, knn_prob_t2, x_train=x_t2, y_train=
       data_t2$y, K_val=3)
8    # pred_knn3_t2e <- ifelse(probs_knn3_t2e > 0.5, 1, 0)
9    # table(Actual=as.numeric(data_t2$y)-1, Predicted=pred_knn3_t2e)
10   # prop.table(table(Actual=as.numeric(data_t2$y)-1, Predicted=pred_
       knn3_t2e), margin=1)
11   # Solution indicates identical predictions for this small dataset.
```

Listing 129: Confusion Matrices (Ex Set 2, T1e)

4.2.2   Task 3: Logistic Regression vs LDA (`Default` data)

3a: Fit Logistic and LDA, 50/50 split

```
1    # library(ISLR); library(MASS)
2    # data(Default)
3    # Default$default_numeric <- ifelse(Default$default == "Yes", 1, 0)
       % \textit{Make numeric for comparison} %
4
5    # set.seed(1) % (Or a seed specified in tutorial if different) %
6    # n_def <- nrow(Default)
7    # train_idx_def <- sample(1:n_def, floor(n_def/2))
8    # train_def <- Default[train_idx_def, ]
9    # test_def  <- Default[-train_idx_def, ]
10
11   # logreg_def <- glm(default ~ balance + income + student, data=train
       _def, family="binomial")
12   # lda_def <- lda(default ~ balance + income + student, data=train_
       def)
13   #
14   # prob_logreg_def_test <- predict(logreg_def, newdata=test_def, type
       ="response")
15   # prob_lda_def_test <- predict(lda_def, newdata=test_def)$posterior
       [, "Yes"] % \textit{Or correct class index} %
16   # head(data.frame(LogRegProb=prob_logreg_def_test, LDAProb=prob_lda_
       def_test))
```

Listing 130: Logistic vs LDA (Ex Set 2, T3a)

3b: Predict test data, confusion matrices. Remove variables?

```
1    # threshold <- 0.5
2    # pred_log_class <- ifelse(prob_logreg_def_test > threshold, "Yes",
       "No")
```

```
3   # conf_log <- table(Actual=test_def$default, Predicted=pred_log_
        class)
4   # print("Logistic Confusion Matrix:"); print(conf_log); print(prop.
        table(conf_log,1))
5   #
6   # pred_lda_class <- ifelse(prob_lda_def_test > threshold, "Yes", "No
        ")
7   # conf_lda <- table(Actual=test_def$default, Predicted=pred_lda_
        class)
8   # print("LDA Confusion Matrix:"); print(conf_lda); print(prop.table(
        conf_lda,1))
9   # Solution note: Both good at Y=0, not Y=1. Lowering threshold might
         help.
10  # Removing variables: Try fitting models without 'student' or '
        income' if they were
11  # not significant or if LDA assumptions are violated by them.
        Compare results.
```

Listing 131: Evaluate Logistic LDA (Ex Set 2, T3b)

3c: Suitability of predictors for LDA. Other methods? *Question: Are student, balance, income all suitable for LDA? Other methods? Solution Notes: - LDA assumes normality of predictors within classes and equal covariance. - 'balance' and 'income' (continuous) might be reasonably normal. 'student' (binary factor) violates normality. - If normality/equal covariance is strongly violated, Logistic Regression is more robust. - QDA could be better if covariance matrices differ between default/non-default groups. - Non-parametric methods like KNN or Trees don't make these distributional assumptions.*

4.2.3   Task 4 : Cross-Validation Methods (`Auto` data)

This task explains data prep for `Auto` (create binary 'y' from 'mpg', 'age' from 'year', remove 'mpg,name,year') and then applies validation set, LOOCV, and k-fold CV to a logistic regression model for 'y'.

```
1   # library(ISLR)
2   # Auto_t4 <- Auto
3   # Auto_t4$y <- as.factor(ifelse(Auto_t4$mpg > median(Auto_t4$mpg), "
        high", "low"))
4   # Auto_t4$age <- 83 - Auto_t4$year
5   # Auto_t4 <- Auto_t4[, !(names(Auto_t4) %in% c("mpg", "name", "year
        "))]
6
7   # (b) Validation set (50/50 split)
8   # set.seed(1) % (Assume a seed) %
9   # n_auto_t4 <- nrow(Auto_t4)
10  # train_idx_t4 <- sample(1:n_auto_t4, floor(n_auto_t4/2))
11  # train_auto_t4 <- Auto_t4[train_idx_t4, ]
12  # test_auto_t4  <- Auto_t4[-train_idx_t4, ]
13  # m1_t4 <- glm(y ~ ., data=train_auto_t4, family="binomial")
14  # prob_t4_val <- predict(m1_t4, newdata=test_auto_t4, type="response
        ")
15  # pred_t4_val <- ifelse(prob_t4_val > 0.5, "high", "low")
16  # # table(test_auto_t4$y, pred_t4_val); mean(pred_t4_val != test_
        auto_t4$y)
17
18  # (d) LOOCV (Manual loop shown in exercise, or use boot::cv.glm)
```

```
19  # library(boot)
20  # m_all_t4 <- glm(y ~ ., data=Auto_t4, family="binomial")
21  # cv_error_loo_t4 <- cv.glm(Auto_t4, m_all_t4, K=nrow(Auto_t4))$
       delta[1] % (Requires cost function for error rate) %
22  # Manual loop from exercise set:
23  # error_rate_LOO_t4 <- mean(pred_LOO_t4_class != Auto_t4$y)
24  # % (Where pred_LOO_t4_class is from the loop. Solution gets 0.089 (
       high=0.923, low=0.117 error prop)) %
25
26  # (e) 8-fold CV (Manual loop shown in exercise, or use boot::cv.glm
       with K=8)
27  # error_rate_k8_t4 <- mean(pred_k8_t4_class != Auto_t4$y)
28  # % (Solution gets 0.094 (high=0.903, low=0.117 error prop)) %
29
30  # (f) Preference for predictors: All vs. some (e.g., age + weight)
31  # Fit LOOCV with reduced model: y ~ age + weight
32  # Compare LOOCV error rates. Solution: Reduced model was worse.
```

Listing 132: CV Methods Example (Ex Set 2, T4)

4.3 Tutorial 3

4.3.1 Task 1: Shrinkage Methods (`College` data)

Predict `Apps` using `Accept, Top10perc, Expend`. Compare OLS, Ridge, Lasso.

1a: Compare OLS, Ridge, Lasso for different λs

```
1   # library(ISLR2); library(glmnet)
2   # xdata_t3 <- subset(College, select=c("Apps", "Accept", "Top10perc
       ", "Expend"))
3   # x_mat_t3 <- as.matrix(xdata_t3[,-1])
4   # y_vec_t3 <- xdata_t3[,1]
5
6   # OLS
7   # ols_t3 <- lm(Apps ~ Accept + Top10perc + Expend, data=xdata_t3)
8   # coef_ols_t3 <- coef(ols_t3)
9
10  # Ridge (example lambda=100)
11  # ridge_t3_100 <- glmnet(x_mat_t3, y_vec_t3, alpha=0, lambda=100)
12  # coef_ridge_t3_100 <- coef(ridge_t3_100)
13  # % (Repeat for other lambdas) %
14
15  # Lasso (example lambda=100)
16  # lasso_t3_100 <- glmnet(x_mat_t3, y_vec_t3, alpha=1, lambda=100)
17  # coef_lasso_t3_100 <- coef(lasso_t3_100)
18  # % (Repeat for other lambdas. Compare how coefficients shrink/go to
       zero) %
```

Listing 133: Shrinkage Comparison (Ex Set 3, T1a)

1b: Optimal λ using `cv.glmnet()` *Explain what `cv.glmnet()` does: Performs k-fold CV (default 10) over a grid of λ values to find λ that minimizes CV error (e.g., MSE).*

```
1   # cv_ridge_t3 <- cv.glmnet(x_mat_t3, y_vec_t3, alpha=0)
2   # lambda_ridge_opt_t3 <- cv_ridge_t3$lambda.min % Sol: 364.8993 %
3   # cv_lasso_t3 <- cv.glmnet(x_mat_t3, y_vec_t3, alpha=1)
4   # lambda_lasso_opt_t3 <- cv_lasso_t3$lambda.min
```

1c: Re-estimate with optimal λ, compare coefficients

```
# ridge_opt_t3 <- glmnet(x_mat_t3, y_vec_t3, alpha=0, lambda=lambda_
    ridge_opt_t3)
# lasso_opt_t3 <- glmnet(x_mat_t3, y_vec_t3, alpha=1, lambda=lambda_
    lasso_opt_t3)
# print(cbind(OLS=coef_ols_t3, Ridge=coef(ridge_opt_t3), Lasso=coef(
    lasso_opt_t3)))
```

Listing 135: Re-estimate and Compare Coefs (Ex Set 3, T1c)

1d: Evaluate OLS, Ridge, Lasso by testMSE (validation set)

```
# n_t3 <- nrow(xdata_t3)
# set.seed(1) % (Assume a seed) %
# ind_t3 <- sample(1:n_t3, size=floor(n_t3/2))
# train_t3 <- xdata_t3[ind_t3,]; test_t3 <- xdata_t3[-ind_t3,]
# Xtrain_t3 <- as.matrix(train_t3[,-1]); ytrain_t3 <- train_t3[,1]
# Xtest_t3  <- as.matrix(test_t3[,-1]);  ytest_t3  <- test_t3[,1]

# OLS_train_t3 <- lm(Apps ~ ., data=train_t3)
# pred_ols_t3 <- predict(OLS_train_t3, newdata=test_t3)
# mse_ols_t3 <- mean((ytest_t3 - pred_ols_t3)^2) % Sol: 1293261 %

# Ridge (using lambda_ridge_opt_t3 from full data CV for simplicity,
     ideally re-CV on train_t3)
# ridge_train_t3 <- glmnet(Xtrain_t3, ytrain_t3, alpha=0, lambda=
    lambda_ridge_opt_t3)
# pred_ridge_t3 <- predict(ridge_train_t3, newx=Xtest_t3)
# mse_ridge_t3 <- mean((ytest_t3 - pred_ridge_t3)^2) % Sol: 1399610
    %
# % (LASSO similarly. Solution: OLS better than Ridge here.) %
```

Listing 136: Test MSE Comparison (Ex Set 3, T1d)

4.3.2   Task 2: Non-linear Models (Small Dataset)

Polynomial regression vs. regression splines.

2a: Plot y against x

2b: Polynomial regression K=7, plot predictions

```
# x_t3_task2 <- 1:8
# y_t3_task2 <- c(8.88, 4.54, 5.12, 1.14, 2.26, 8.43, 10.92, 14.47)
# data_t3_task2 <- data.frame(x=x_t3_task2, y=y_t3_task2)
# poly7_t3 <- lm(y ~ poly(x, 7, raw=TRUE), data=data_t3_task2) % (
    raw=TRUE for x, x^2...) %
# plot(data_t3_task2$x, data_t3_task2$y)
# lines(data_t3_task2$x, predict(poly7_t3), col="red") % (Perfect
    fit as N=P) %
```

Listing 137: Polynomial Regression (Ex Set 3, T2b)

2c: Regression spline (manual basis) with knots at 3, 6. Plot.

```
# h_func <- function(x, xi) pmax((x-xi)^3, 0)
# spline_manual_t3 <- lm(y ~ x + I(x^2) + I(x^3) + h_func(x,3) + h_
    func(x,6), data=data_t3_task2)
# lines(data_t3_task2$x, predict(spline_manual_t3), col="blue")
# % (Expect worse predictions for new x than poly7, especially
    outside [1,8]) %
```

Listing 138: Manual Regression Spline (Ex Set 3, T2c)

2d: Regression spline with `bs()` (knots 3,6). Compare preds.

```
# library(splines)
# spline_bs_t3 <- lm(y ~ bs(x, knots=c(3,6)), data=data_t3_task2)
# lines(data_t3_task2$x, predict(spline_bs_t3), col="green")
# % (Predictions should be identical to 2c if basis is equivalent,
    which it is for cubic) %
```

Listing 139: bs() Spline (Ex Set 3, T2d)

2e: LOOCV for different knot placements

```
# testMSE_spline <- function(knot1, knot2, x_data, y_data) {
#   # ... (LOOCV loop as in earlier examples, fitting lm with bs(x,
    knots=c(knot1,knot2))) ...
#   return(mean_SE_from_LOOCV)
# }
# # testMSE_spline(3, 6, x_t3_task2, y_t3_task2) % Sol: 2044 %
# # testMSE_spline(3.5, 5.5, x_t3_task2, y_t3_task2) % Sol: 377 (
    better) %
```

Listing 140: LOOCV for Knots (Ex Set 3, T2e)

4.3.3   Task 3: Extension of KNN (Nadaraya-Watson)

Implement KNN as weighted average, then Nadaraya-Watson.

3a: Properties of weights w_i   *(Sum to 1, large w_i for small $|x_i - x_0|$).*

3b: Implement KNN as weighted average (`knn2`)

3c: Modify knn2 to Nadaraya-Watson nw), plot for different h.

```
# weight_nw <- function(x0, x_vec, h_val=0.2) dnorm((x_vec-x0)/h_val
    ) / sum(dnorm((x_vec-x0)/h_val))
# nw_func <- function(x0, x_vec, y_vec, h_val=0.2) {
#   w <- weight_nw(x0, x_vec, h_val)
#   return(sum(wy_vec))
# }
# # Simulate x_sim, y_sim from exercise
# # set.seed(12); x_sim <- seq(-50,50,length=100)
# # y_sim <- 1-0.05x_sim-0.02x_sim^2+0.003x_sim^3+50rnorm(100)
# # plot(x_sim, y_sim)
# # pred_nw_h1 <- sapply(x_sim, nw_func, x_vec=x_sim, y_vec=y_sim, h
    _val=1)
# # lines(x_sim, pred_nw_h1, col="red")
# # pred_nw_h5 <- sapply(x_sim, nw_func, x_vec=x_sim, y_vec=y_sim, h
    _val=5)
# # lines(x_sim, pred_nw_h5, col="blue")
```

Listing 141: Nadaraya-Watson (Ex Set 3, T3c)

3d: Meaning of bandwidth h *(Controls "localness" of estimator).*

4.4   Tutorial 4

4.4.1   Task 1 : Regression Trees (`Hitters` data)

Predict `logSalary` using `Hits` and `Years`.

1a: Explain RSS function and `optimise()` for first split.

```
# library(ISLR); data(Hitters)
# Hitters <- Hitters[!is.na(Hitters$Salary),]
# Hitters$logSalary <- log(Hitters$Salary)
# RSS_func_t4 <- function(s, x1_vec, y_vec) {
#    yR1 <- mean(y_vec[x1_vec < s], na.rm=TRUE)
#    yR2 <- mean(y_vec[x1_vec >= s], na.rm=TRUE)
#    rss <- sum((y_vec[x1_vec < s] - yR1)^2, na.rm=TRUE) +
#           sum((y_vec[x1_vec >= s] - yR2)^2, na.rm=TRUE)
#    return(rss)
# }
# opt_hits <- optimise(RSS_func_t4, lower=0, upper=250,
#                      x1_vec=Hitters$Hits, y_vec=Hitters$logSalary)
# % $min=117.5, $obj=160.97 %
# opt_years <- optimise(RSS_func_t4, lower=0, upper=50,
#                      x1_vec=Hitters$Years, y_vec=Hitters$
   logSalary)
# % $min=4.5, $obj=115.06 %
# # Years gives lower RSS, so first split is on Years < 4.5
```
Listing 142: RSS function for Tree Split (Ex Set 4, T1a)

1b: Fit tree with `tree()`, plot. Compare first split.

```
# library(tree)
# tree1_t4 <- tree(logSalary ~ Hits + Years, data=Hitters)
# plot(tree1_t4); text(tree1_t4, pretty=0)
# % First split in tree should be Years < 4.5 (or very close) %
```
Listing 143: Fit and Plot Tree (Ex Set 4, T1b)

1c: Predict for new observation (Hits=125, Years=2.5)   *Follow path in tree1_t4: Years=2.5 (¡4.5) -¿ Left. Hits=125 (¿114 if that's next split) -¿ Right in subtree. Read leaf value. Sol: 5.264*

1d: Fit tree with all predictors

```
# tree2_t4 <- tree(logSalary ~ . - Salary, data=Hitters) # Exclude
   original Salary
# summary(tree1_t4) % RSS is deviance: 69.06 for tree1 %
# summary(tree2_t4) % RSS is deviance: 43.03 for tree2 (lower as
   expected) %
# rss1_t4 <- sum(residuals(tree1_t4)^2)
# rss2_t4 <- sum(residuals(tree2_t4)^2)
```
Listing 144: Tree with All Predictors and RSS (Ex Set 4, T1d)

4.4.2   Task 2: Bagging and Random Forest (`Hitters` data)

Benchmark: `tree2` from T1d.

2a: TestMSE for `tree2` (validation set)

```
1  # set.seed(123)
2  # n_hit <- nrow(Hitters)
3  # ind_hit <- sample(1:n_hit, size=floor(n_hit/2))
4  # train_hit <- Hitters[ind_hit,]; test_hit <- Hitters[-ind_hit,]
5  # tree2_train_t4 <- tree(logSalary ~ . - Salary, data=train_hit)
6  # pred2_test_t4 <- predict(tree2_train_t4, newdata=test_hit)
7  # mse2_test_t4 <- mean((test_hit$logSalary - pred2_test_t4)^2) % Sol
       : 0.428 %
```

Listing 145: Test MSE for tree2 (Ex Set 4, T2a)

2b: Bagging of `tree2`. Variable Importance.

```
1   # library(randomForest)
2   # set.seed(123) % (Ensure consistent seed for RF if comparing) %
3   # # For bagging, mtry = number of predictors (19 for Hitters
        excluding Salary, logSalary)
4   # bagging_t4 <- randomForest(logSalary ~ . - Salary, data=train_hit,
         mtry=19, importance=TRUE)
5   # predbag_t4 <- predict(bagging_t4, newdata=test_hit)
6   # msebag_t4 <- mean((test_hit$logSalary - predbag_t4)^2) % Sol:
        0.3015 %
7   # varImpPlot(bagging_t4)
8   # % x-axis is MeanDecreaseGini (for classification) or IncNodePurity
         (RSS decrease for regression).
9   #   Help says for regression: total decrease in node impurities (
       RSS) from splitting on the variable,
10  #    averaged over all trees. CRuns most important. %
```

Listing 146: Bagging and VarImp (Ex Set 4, T2b)

2c: Random Forest. How is it an extension of bagging?

```
1  # library(randomForest)
2  # set.seed(123)
3  # rf_t4 <- randomForest(logSalary ~ . - Salary, data=train_hit, mtry
      =3, importance=TRUE) % (mtry=3 from solution) %
4  # predrf_t4 <- predict(rf_t4, newdata=test_hit)
5  # mserf_t4 <- mean((test_hit$logSalary - predrf_t4)^2)
6  # % Random Forest is an extension because at each split, only a
      random subset (mtry) of predictors
7  #   is considered, decorrelating trees and potentially improving
      variance reduction. %
```

Listing 147: Random Forest (Ex Set 4, T2c)

4.4.3 Task 3: Boosting (`Hitters` data)

*Predict **logSalary** with boosted tree. Compute testMSE.*

```
1  # library(gbm)
2  # # Using train_hit, test_hit from 2a
3  # set.seed(123)
4  # boost_t4 <- gbm(logSalary ~ . - Salary, data=train_hit,
      distribution="gaussian",
5  #                 n.trees=1000, interaction.depth=4, shrinkage=0.01)
6  # predgbm_t4 <- predict(boost_t4, newdata=test_hit, n.trees=1000)
7  # msegbm_t4 <- mean((test_hit$logSalary - predgbm_t4)^2) % Sol:
      0.2666 %
8  # % Boosting often yields lowest MSE among tree ensembles here. %
```

4.4.4 Task 4: Support Vector Machines (Simulated Data 'xy.csv')

Classes "1" and "-1".

4a: Possible to find maximal marginal classifier from plot? *(No, classes overlap).*

4b: Recreate plot from 'xy.csv'.

4c: Use SVC (linear kernel) to classify. Plot.

```
# xy_data_t4 <- read.csv("xy.csv")
# names(xy_data_t4) <- c("x1", "x2", "y_class") % (Assuming column
    names) %
# xy_data_t4$y_factor <- as.factor(xy_data_t4$y_class)
# library(e1071)
# svmfit_linear_t4 <- svm(y_factor ~ x1 + x2, data=xy_data_t4,
    kernel="linear",
#                         cost=10, scale=FALSE)
# plot(svmfit_linear_t4, xy_data_t4, x2 ~ x1) % (Adjust formula for
    plot if needed) %
```

Listing 149: SVC (Ex Set 4, T4c)

4d: Meaning of tuning parameter C=10. *(Allows for some misclassifications/margin violations, controls trade-off between margin width and violations. Larger C = narrower margin, fewer violations).*

4e: SVM with polynomial kernel. Better separation?

```
# svmfit_poly_t4 <- svm(y_factor ~ x1 + x2, data=xy_data_t4, kernel
    ="polynomial",
#                         cost=10, scale=FALSE) % (Default degree=3) %
# plot(svmfit_poly_t4, xy_data_t4, x2 ~ x1)
# % Polynomial kernel can create non-linear boundary, might separate
     training data better,
#    but risk of overfitting. %
```

Listing 150: SVM Polynomial Kernel (Ex Set 4, T4e)

4.4.5 Task 5: PCA (`advertising.csv`)

PCA on TV, radio, newspaper. Interpret first two PCs from loadings.

```
# advertising_t5 <- read.csv("advertising.csv")[, c("TV", "radio", "
    newspaper")]
# pca_t5 <- prcomp(advertising_t5, scale.=TRUE, center.=TRUE)
# summary(pca_t5)
# print(pca_t5$rotation)
# % PC1: High loading on TV (-0.999), small on others. Represents
    overall ad spend dominated by TV.
#    PC2: High positive on newspaper (0.93), moderate positive on
    radio (0.36), small on TV. Represents
```

```
7  #    print media focus vs. TV. (Signs might be flipped but
          interpretation similar). %
```

<div align="center">Listing 151: PCA (Ex Set 4, T5a)</div>

4.4.6 Task 6: Cluster Analysis (`USArrests` data)

6a: K-means clustering (K=3) on Murder, Assault. Plot.

```
1  # data(USArrests)
2  # x_km_t6 <- scale(USArrests[, c("Murder", "Assault")]) # Scale data
3  # set.seed(1) % (For reproducibility of kmeans) %
4  # km3_t6 <- kmeans(x_km_t6, centers=3, nstart=20)
5  # plot(x_km_t6, col=km3_t6$cluster, pch=19, main="K-Means (K=3) on
          Murder & Assault")
6  # points(km3_t6$centers, col=1:3, pch=4, cex=3, lwd=3)
7  # text(x_km_t6, labels=rownames(USArrests), col=km3_t6$cluster, pos
          =3, cex=0.7)
```

<div align="center">Listing 152: K-Means Clustering (Ex Set 4, T6a)</div>

6b: Hierarchical clustering. Compare.

```
1  # Using scaled data x_km_t6 from above
2  # dist_t6 <- dist(x_km_t6)
3  # hclust_complete_t6 <- hclust(dist_t6, method="complete")
4  # plot(hclust_complete_t6, main="Hierarchical Clustering (Complete
          Linkage)", cex=0.7)
5  # clusters_hclust_k3_t6 <- cutree(hclust_complete_t6, k=3)
6  # table(KMeans=km3_t6$cluster, Hierarchical=clusters_hclust_k3_t6)
7  # % Compare cluster assignments. Differences due to different
          algorithms/criteria. %
```

<div align="center">Listing 153: Hierarchical Clustering (Ex Set 4, T6b)</div>

5 Assignment: Compulsory Assignment BAN404 (Spring 2025x)

Predicting log(CEO Salary) using the `ceosal2` *dataset.*

5.1 Initial Data Loading and Preparation

Loading Packages and Data

```
1  # Ensure wooldridge is installed: install.packages("wooldridge")
2  library(wooldridge)
3  library(dplyr)      # For tidyverse data manipulation
4  library(ggplot2)    # For tidyverse plotting
5  library(glmnet)     # For LASSO
6  library(FNN)        # For knn() function (alternative to manual for
          Task 3)
7  library(gam)        # For GAMs
8  library(boot)       # For cv.glm (though LOOCV can be manual for lm)
9
10 data(ceosal2)
11 # View data structure and summary
12 # str(ceosal2)
13 # summary(ceosal2)
```

```
14  # help(ceosal2) # To understand variables
```
Listing 154: Loading necessary packages and the ceosal2 dataset

Data Splitting (50/50 Train/Test) *As per the assignment's suggested solution, an initial 50/50 split is performed.*

```
1  set.seed(1234) # Seed from assignment
2  n_total <- nrow(ceosal2)
3  ntrain <- floor(n_total/2)
4  train_indices <- sample(1:n_total, size=ntrain)
5
6  # These original splits might be used if subsequent processing is
       done per task
7  # For consistency with assignment solution, a global processed
       dataset is often made first
8  original_train_data <- ceosal2[train_indices,]
9  original_test_data <- ceosal2[-train_indices,]
```
Listing 155: Splitting data into 50/50 training and test sets

5.2 Task 1: Descriptive Statistics and Predictor Investigation

Question *Describe relevant features of the output variable (`lsalary`) with descriptive statistics (tables and graphs). Also use descriptive statistics to investigate promising predictors for `lsalary`. Remove variables that cannot possibly be available when the salary of a new CEO should be predicted.*

Solution

Processing Data for Task 1 *Assumptions based on provided solution: New CEO hired internally. `comten` (company tenure) and `comtensq` are available. `ceoten` (CEO tenure in current company) and `ceotensq` are NOT available for a new CEO. `salary` (untransformed) is obviously not available. `college` is removed due to low variability in the training data example in the solution.*

```
1  # Create the 'lsalary' variable if it's not already present (it is
       in ceosal2)
2  # ceosal2$lsalary is log(salary)
3
4  # Variables to remove
5  vars_to_remove_t1 <- c("salary", "ceoten", "ceotensq", "college")
6  ceosal_task1_filtered <- ceosal2[, !(names(ceosal2) %in% vars_to_
       remove_t1)]
7
8  # Convert 'grad' to factor
9  ceosal_task1_filtered$grad <- as.factor(ceosal_task1_filtered$grad)
10
11  # Now split the PROCESSED data
12  train_t1 <- ceosal_task1_filtered[train_indices,]
13  test_t1  <- ceosal_task1_filtered[-train_indices,]
```
Listing 156: Data processing for Task 1 (Applied to whole dataset first, then split)

Descriptive Statistics for `lsalary` (on `train_t1`)

```r
# Base R
summary(train_t1$lsalary)
hist(train_t1$lsalary, main="Histogram of log(Salary) - Training
    Data", xlab="log(Salary)", col="lightblue", breaks=12)
boxplot(train_t1$lsalary, main="Boxplot of log(Salary) - Training
    Data", ylab="log(Salary)")

# Tidyverse
# library(dplyr); library(ggplot2)
# train_t1 %>% select(lsalary) %>% summary()
# ggplot(train_t1, aes(x=lsalary)) +
#   geom_histogram(aes(y=..density..), fill="lightblue", color="
    black", bins=12) +
#   geom_density(alpha=0.2, fill="red") +
#   ggtitle("Distribution of log(Salary) - Training Data")
```

Listing 157: Descriptive statistics for lsalary (Task 1)

Interpretation: Comment on the distribution of `lsalary` - is it symmetric, skewed, any outliers? The log transformation often makes salary data more symmetric.

Investigating Promising Predictors for `lsalary` (on `train_t1`)

```r
# Base R (as in solution)
# par(mfrow=c(3,3)) # Adjust layout based on number of predictors
    remaining
# plot(lsalary ~ ., data=train_t1)
# graphics.off()

# Tidyverse (example for a few predictors)
# ggplot(train_t1, aes(x=lsales, y=lsalary)) + geom_point() + geom_
    smooth(method="lm") + ggtitle("lsalary vs lsales")
# ggplot(train_t1, aes(x=grad, y=lsalary)) + geom_boxplot() +
    ggtitle("lsalary vs grad")
```

Listing 158: Plotting lsalary against other predictors (Task 1)

Interpretation of plots: - `lsales` shows a clear positive linear relationship. - `sales` also positive, might capture non-linearity if `lsales` also included. - `mktval`, `lmktval` show positive trends. - For categorical `grad`: Check if means/medians differ significantly.

```r
# Base R
# by(train_t1$lsalary, train_t1$grad, summary)
# t_test_grad_0 <- t.test(train_t1$lsalary[train_t1$grad == 0]) # If
    enough observations
# t_test_grad_1 <- t.test(train_t1$lsalary[train_t1$grad == 1])
# print(t_test_grad_0$conf.int)
# print(t_test_grad_1$conf.int)
# if (abs(t_test_grad_0$estimate - t_test_grad_1$estimate) is small
    and CIs overlap substantially,
# then 'grad' might not be a strong predictor alone).

# Tidyverse
# train_t1 %>% group_by(grad) %>% summarise(mean_lsalary = mean(
    lsalary), sd_lsalary = sd(lsalary))
```

Listing 159: Statistical test for 'grad' (Task 1)

Conclusion from Task 1 based on solution: `lsales` is the most promising. Others warrant investigation in modeling. `grad` seems weak.

5.3 Task 2: Linear Regression and LASSO

Question *Use linear regression (on all or a subset of the predictors) and LASSO to predict* `lsalary`. *Evaluate the predictions on test data using an appropriate error measure (MSE).*

Solution *We use* `train_t1` *and* `test_t1` *from the end of Task 1 preparation.*

OLS with All Predictors

```
ols_model_all <- lm(lsalary ~ ., data=train_t1)
# summary(ols_model_all) # For coefficient details

ols_all_preds_test <- predict(ols_model_all, newdata=test_t1)
ols_all_mse_test <- mean((test_t1$lsalary - ols_all_preds_test)^2)
cat("Test MSE (OLS All Predictors):", ols_all_mse_test, "\n")
# Solution had Test MSE = 0.3044908
```

Listing 160: OLS with all predictors (Task 2)

OLS with `lsales` Only

```
ols_model_lsales <- lm(lsalary ~ lsales, data=train_t1)
# summary(ols_model_lsales)

ols_lsales_preds_test <- predict(ols_model_lsales, newdata=test_t1)
ols_lsales_mse_test <- mean((test_t1$lsalary - ols_lsales_preds_test
    )^2)
cat("Test MSE (OLS lsales Only):", ols_lsales_mse_test, "\n")
# Solution had Test MSE = 0.309084
```

Listing 161: OLS with lsales only (Task 2)

Comment: The test MSEs are very similar, suggesting `lsales` *captures most of the predictive power. The solution mentions LOOCV for a more robust comparison, where the* `lsales`*-only model was better.*

LASSO Regression

```
library(glmnet)
# Prepare matrices for glmnet (factors need to be dummified)
# 'grad' is the only factor in train_t1/test_t1
x_train_lasso_t2 <- model.matrix(lsalary ~ . , data=train_t1)[,-1] #
    -1 to remove intercept column
y_train_lasso_t2 <- train_t1$lsalary
x_test_lasso_t2  <- model.matrix(lsalary ~ . , data=test_t1)[,-1]

set.seed(1234) # For cv.glmnet
cv_lasso_t2 <- cv.glmnet(x_train_lasso_t2, y_train_lasso_t2, alpha
    =1, standardize=TRUE)
best_lambda_t2 <- cv_lasso_t2$lambda.min
# print(paste("Best lambda for LASSO:", best_lambda_t2))

lasso_model_t2 <- glmnet(x_train_lasso_t2, y_train_lasso_t2, alpha
    =1,
                         lambda=best_lambda_t2, standardize=TRUE)
# print("LASSO Coefficients:")
# print(coef(lasso_model_t2))
# Solution notes its LASSO (on its training split) picked lsales and
    lmktval.
```

```
19   lasso_preds_test_t2 <- predict(lasso_model_t2, newx=x_test_lasso_t2)
20   lasso_mse_test_t2 <- mean((test_t1$lsalary - lasso_preds_test_t2)^2)
21   cat("Test MSE (LASSO):", lasso_mse_test_t2, "\n")
22   # Solution LOOCV MSE for LASSO was 0.3777, worse than lsales-only
         OLS.
23   # Compare your test MSE here with the OLS models.
```

Listing 162: LASSO regression (Task 2)

5.4 Task 3: KNN Regression

Question *Predict `lsalary` by KNN, using only `lsales` as a predictor. Determine K using leave-one-out cross-validation. Also, evaluate the predictions and compare with the results in 2.*

Solution *The assignment solution uses LOOCV on the full processed dataset (`ceosal_final_f` to determine K. For strictness, K should be determined using only the training portion (`train_t1`) or an inner CV loop if the final evaluation is on `test_t1`. We follow the provided solution's approach for K determination for comparability with its results.*

Determine Optimal K using LOOCV
```
1    # Using ceosal_final_filtered from Task 1 for LOOCV to find K
2    # (This uses data that will also be in the "test" part of the
         original 50/50 split)
3    knn_reg_func_t3 <- function(x0, x_train_vec, y_train_vec, K_val) {
4      d <- abs(x_train_vec - x0)
5      o <- order(d)
6      # Ensure K_val is not larger than the number of available
          neighbors
7      K_actual <- min(K_val, length(o))
8      return(mean(y_train_vec[o[1:K_actual]]))
9    }
10
11   lsales_all_t3 <- ceosal_task1_filtered$lsales % \textit{Using the
         dataset after initial processing} %
12   lsalary_all_t3 <- ceosal_task1_filtered$lsalary
13   n_all_t3 <- nrow(ceosal_task1_filtered)
14   k_max_t3 <- 40 # As per solution
15   loocv_mse_k_t3 <- numeric(k_max_t3 - 1)
16
17   for (k_val_t3 in 2:k_max_t3) {
18     se_fold_t3 <- numeric(n_all_t3)
19     for (i in 1:n_all_t3) {
20       xtrain_knn_t3 <- lsales_all_t3[-i]
21       ytrain_knn_t3 <- lsalary_all_t3[-i]
22       xtest_knn_t3  <- lsales_all_t3[i]
23       ytest_knn_t3  <- lsalary_all_t3[i]
24
25       pred_knn_fold_t3 <- knn_reg_func_t3(xtest_knn_t3, xtrain_knn_t3,
             ytrain_knn_t3, K_val=k_val_t3)
26       se_fold_t3[i] <- (ytest_knn_t3 - pred_knn_fold_t3)^2
27     }
28     loocv_mse_k_t3[k_val_t3-1] <- mean(se_fold_t3)
29   }
30
```

```
31  # plot(2:k_max_t3, loocv_mse_k_t3, type="b", xlab="K", ylab="LOOCV
        MSE", main="LOOCV for KNN K selection")
32  optimal_K_t3 <- (2:k_max_t3)[which.min(loocv_mse_k_t3)]
33  # cat("Optimal K from LOOCV on full processed data:", optimal_K_t3,
        "\n")
34  # Solution found K=24 (index 23 in their 2:kmax loop).
35  min_loocv_mse_knn_t3 <- min(loocv_mse_k_t3)
36  # cat("LOOCV MSE for KNN with K=", optimal_K_t3, ":", min_loocv_mse_
        knn_t3, "\n")
37  # Solution result for K=24: 0.2840139.
```
Listing 163: KNN: Determine K using LOOCV (Task 3)

Comparison with Task 2 Results *LOOCV MSE for KNN (K=[optimal_K_t3]): [min_loocv_mse_knn_t3] From solution's LOOCV for Task 2: - OLS (lsales only) LOOCV MSE: 0.2685442 - LASSO LOOCV MSE: 0.3777567 - OLS (all predictors) LOOCV MSE: 0.4253121 Conclusion: KNN with only `lsales` (optimal K) performs better than LASSO and full OLS in terms of LOOCV MSE, but slightly worse than the simple OLS model using only `lsales`.*

5.5 Task 4: KNN with Multiple Predictors and GAM

Question (Part 1 - KNN with Multiple Predictors) *Explain how you would use KNN if you had more than one predictor. Two alternative ways to do this have been mentioned in the course. If you would like to you can also write R-code to do this but this is not required to pass the assignment.*

Solution (Textual) *(Same textual explanation as in Hypothetical Exam 2025 for: 1. Multivariate Distance Metric for KNN, 2. GAM Framework with KNN as Univariate Smoother via Backfitting).*

Question (Part 2 - GAM) *Argue for which, if any, variables might have a nonlinear relationship with `lsalary` and find a good model within the generalized additive models (GAM) framework. Evaluate the predictions.*

Solution (using gam) *From Task 1, plots suggested `sales` and `mktval` might have non-linear relationships. The solution investigates these using LOOCV.*

```
1   library(gam)
2   # Using ceosal_final_train, ceosal_final_test (or ceosal_task1_
        filtered for LOOCV as in solution)
3
4   # Visual inspection on training data (as in solution)
5   # gam_inspect_t4 <- gam(lsalary ~ s(sales, df=4) + s(mktval, df=4) +
        lsales + lmktval +
6   #                                   age + grad + comten + profits +
        comtensq + profmarg,
7   #                       data=train_t1) % \textit{(Using train_t1,
        consistent with solution's Task 1 plots)} %
8   # par(mfrow=c(1,2))
9   # plot(gam_inspect_t4, se=TRUE, terms=c("s(sales, df = 4)", "s(
        mktval, df = 4)"))
10  # graphics.off()
```

```
11  # Solution comment based on this: "contribution of non-linear mktval
        looks questionable".
12
13  # LOOCV for GAM with s(sales) and other predictors linearly.
14  # (Using the full processed dataset ceosal_task1_filtered for LOOCV
        as per solution)
15  n_gam_t4 <- nrow(ceosal_task1_filtered)
16  se_gam_s_sales_t4 <- numeric(n_gam_t4)
17
18  # Formula based on solution: only s(sales) is non-linear, others
        linear.
19  # (The solution eventually only uses s(sales) in its final GAM LOOCV
        comparison)
20  formula_gam_s_sales <- lsalary ~ s(sales, df=4)
21  # To compare with other models, it's better to include other linear
        terms if they were in OLS.
22  # For this example, let's test the simple s(sales) model as in the
        solution's final GAM part.
23
24  # for(i in 1:n_gam_t4) {
25  #    train_gam_fold_t4 <- ceosal_task1_filtered[-i,]
26  #    test_gam_fold_t4  <- ceosal_task1_filtered[i,]
27  #
28  #    # Ensure factor levels are handled if factors were included in
        formula_gam_s_sales
29  #    # For lsalary ~ s(sales, df=4), this is not an issue.
30  #    gam_fit_fold_t4 <- gam(formula_gam_s_sales, data=train_gam_fold_
        t4)
31  #    pred_gam_fold_t4 <- predict(gam_fit_fold_t4, newdata=test_gam_
        fold_t4)
32  #    se_gam_s_sales_t4[i] <- (test_gam_fold_t4$lsalary - pred_gam_
        fold_t4)^2
33  # }
34  # mse_gam_s_sales_loocv_t4 <- mean(se_gam_s_sales_t4)
35  # cat("LOOCV MSE (GAM with s(sales, df=4) only):", mse_gam_s_sales_
        loocv_t4, "\n")
36  # Solution result for GAM with s(sales) only: 0.2918272.
37
38  # Comparison:
39  # KNN (K=24, lsales only) LOOCV MSE: 0.2840139
40  # OLS (lsales only) LOOCV MSE: 0.2685442
41  # Conclusion: "GAM with s(sales) is slightly worse than KNN with
        lsales, and also worse than
42  # simple OLS with lsales. The log-transformation in lsales seems to
        capture the main
43  # non-linearity effectively for this dataset."
```

Listing 164: GAM for lsalary (Task 4)